arXiv:2610.17925v1 [cs.LG] 13 Oct 2026

Stein Variational Distributional Reinforcement Learning: Score-Guided Particles for Multivariate Return Distributions

Manuel Haußmann **Melih Kandemir**
University of Southern Denmark

Abstract

Decisions that weigh several reward signals against each other require the joint distribution of the return vector. Categorical critics scale exponentially in the reward dimension, and quantiles have no canonical multivariate analogue. Multivariate distributional critics therefore represent each return distribution by M equal-weight particles trained with a discrepancy loss. Beyond their atoms, the particles provide no sampler and cannot resolve events rarer than $1/M$. We propose Stein variational distributional RL (SVDRL). Its critic outputs Stein particles, which Stein variational gradient descent transports along a learned score field of the distributional Bellman target. We learn this field by bootstrapped denoising score matching across noise levels and tie it across the Bellman recursion by a consistency loss, which we prove is a non-expansive temporal-difference scheme in score space. The same field certifies the particles. Its reverse-diffusion sampler satisfies a supremal Wasserstein policy-evaluation bound that holds for bounded, possibly multimodal returns. The Stein particles inherit this bound up to a gap that we measure without ground truth, and the sampler answers tail-risk queries below the particle resolution at any sample size. Empirically, SVDRL is more accurate on rare events than particle critics with the same number of particles, and on par on MuJoCo continuous control.

1 INTRODUCTION

Distributional reinforcement learning models the distribution of the return rather than only its expectation (Bellemare et al., 2017). With multivariate rewards, risk- and correlation-sensitive decisions depend on the joint distribution of the return components (Zhang et al., 2021; Hayes et al., 2022). The univariate representations do not carry over: categorical grids (Bellemare et al., 2017) need $\mathcal{O}(M^d)$ atoms, and quantiles (Dabney et al., 2018a,b) have no canonical analogue for $d > 1$. Existing multivariate methods therefore represent the return distribution by M particles trained with a discrepancy loss, e.g., MMD (Zhang et al., 2021), Sinkhorn divergences (Sun et al., 2024), or sliced divergences (Debes and Tuytelaars, 2026b).

The theory of these methods covers the Bellman operator, which contracts in the respective supremal divergence, but not the learned particles. We show that this gap is structural. Equal-weight particles carry no density and no sampler, and M of them cannot represent an event of probability below $1/M$, whatever loss trains them (Proposition 1 and Figure 1).

We instead learn the *return score field*, the smoothed score of the return distribution at all noise levels. Smoothing commutes with the Bellman update up to a rescaling of the noise level (Lemma 1), so the field satisfies its own Bellman equation, and we learn it along a noise schedule (Song and Ermon, 2019) by bootstrapped denoising score matching (Lemma 4) and by temporal differences in score space (Theorem 2). A single noise level does not suffice, as it cannot certify mode masses (Remark 2). We sample from the field by reverse diffusion. For fast action selection we additionally keep M particles per state–action pair, the Stein particles, which we transport along the field by Stein variational gradient descent (SVGD) (Liu and Wang, 2016) and amortize into a critic.

The training risk controls the sampled return distribution in total variation and w_1 for bounded multimodal returns, where log-concavity and related assumptions fail (Theorem 3). This yields a policy-evaluation bound in $\bar{w}_1$ (Theorem 4), certified mode masses (Corollaries 6 and 8), and a CVaR estimate whose error does not depend on M (Theorem 6).

No single critic is accurate on every return law, and we show which error limits each family. Equal-weight particles suffer from a *resolution floor*: they round

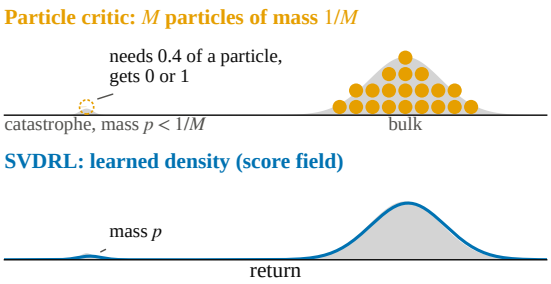

Figure 1: *A rare event below the particle resolution.* Schematic. A catastrophe of mass $p < 1/M$ needs a fraction of a particle, so particle critics round its probability to a multiple of $1/M$ (Proposition 1). SVDRL learns a density, which carries any mass; on the rare-event chain it estimates $p = 0.005$ far better than particle critics with $M = 128$ (Section 5 and Figure 5).

event probabilities and mixture weights to multiples of $1/M$ (Proposition 1). A score field suffers from *smoothing bias* and is hard to learn where the return law is atomic or its modes lie far apart relative to their widths (Proposition 2 and Lemma 2). SVDRL keeps both objects, and a *hybrid rule* reports, per state–action pair and without ground truth, the Stein particles on separated or atomic clusters and the sampler draws elsewhere (Definition 1 and Table 1).

SVDRL is, to our knowledge, the first certified multivariate score-based critic for possibly multimodal returns, and its Stein particles inherit the certificate up to a gap measured without ground truth (Theorem 4). At equal particle count, SVDRL is more accurate than the best particle critic in every rare-event comparison (Section 5). On all eight synthetic MDPs the hybrid wins within-mode shape against the competitors as specified, although at matched compute they win the unimodal ones. The competitors' own benchmarks place the boundary at atomic laws, and on MuJoCo SAC-SVDRL is on par in return and has the most accurate critic distribution on three of four. Its stratified Bellman target set and running-mean moment matching improve every particle critic we tested.

2 BACKGROUND AND NOTATION

Markov Decision Processes with Multivariate Rewards. We consider a Markov Decision Process (MDP) defined by $(\mathcal{X}, \mathcal{A}, P, \mathcal{R}, \gamma)$, where $\mathcal{X}$ and $\mathcal{A}$ denote state and action spaces, $\gamma \in (0, 1)$ is a discount factor, and $P : \mathcal{X} \times \mathcal{A} \to \mathcal{P}(\mathcal{X})$ is a transition kernel. Unlike standard RL, where the reward is a scalar, we define the reward as a potentially random vector $R(x, a) \sim \mathcal{R}(\cdot|x, a) \in \mathcal{P}(\mathbb{R}^d)$. We define the number of reward components by $d \geq 1$ and recover the classical univariate distributional RL setting for $d = 1$.

Given a policy $\pi : \mathcal{X} \to \mathcal{P}(\mathcal{A})$, the *random return* is

$$Z^\pi(x, a) \triangleq \sum_{t \geq 0} \gamma^t R_t \in \mathbb{R}^d, \quad R_t \sim \mathcal{R}(\cdot|X_t, A_t)$$

with $(X_0, A_0) = (x, a)$, $X_{t+1} \sim P(\cdot|X_t, A_t)$, and $A_t \sim \pi(\cdot|X_t)$ for $t \geq 1$. $V^\pi(x, a) \triangleq \mathbb{E}[Z^\pi(x, a)] \in \mathbb{R}^d$ is its expected *value function*. We aim to infer the *return distribution function* $\eta^\pi(x, a) \triangleq \mathrm{Law}(Z^\pi(x, a)) \in \mathcal{P}(\mathbb{R}^d)$.

Assumption 1 (Bounded returns)**.** The rewards are bounded per dimension, $|R_i| \leq r_{\max,i}$ almost surely for $i \leq d$, and $\gamma \in [0, 1)$. Write $z_{\max} \triangleq r_{\max}/(1 - \gamma) \in \mathbb{R}^d_+$, let $B_R \triangleq \prod_{i \leq d}[-z_{\max,i}, z_{\max,i}]$ be the return box, $R_Z \triangleq \|z_{\max}\|_2$ its circumradius, and $D \triangleq 2R_Z$. Then $|r_i + \gamma z_i| \leq z_{\max,i}$ whenever $z \in B_R$, so $\mathcal{T}^\pi$ maps laws supported in B_R to laws supported in B_R, and $\mathrm{supp}\, \eta^\pi(x, a) \subseteq B_R \subseteq B(0, R_Z)$ for all (x, a).

Distributional Bellman Operator. For $r \in \mathbb{R}^d$ and $\gamma \in (0, 1)$, the *affine bootstrap map* $b_{r,\gamma} : \mathbb{R}^d \to \mathbb{R}^d$ is defined as $b_{r,\gamma}(z) \triangleq r + \gamma z$. The *push-forward* of a distribution $\nu \in \mathcal{P}(\mathbb{R}^d)$ under $b_{r,\gamma}$ is $(b_{r,\gamma})_\# \nu \triangleq \mathrm{Law}(r + \gamma Z)$ for $Z \sim \nu$. The *distributional Bellman operator (DBO)* $\mathcal{T}^\pi : \mathcal{P}(\mathbb{R}^d)^{\mathcal{X} \times \mathcal{A}} \to \mathcal{P}(\mathbb{R}^d)^{\mathcal{X} \times \mathcal{A}}$ is defined as

$$(\mathcal{T}^\pi \eta)(x, a) \triangleq \mathbb{E}_\pi \left[(b_{R,\gamma})_\# \eta(X', A') | X = x, A = a \right], \tag{1}$$

where $R \sim \mathcal{R}(\cdot|x, a)$, $X' \sim P(\cdot|x, a)$, and $A' \sim \pi(\cdot|X')$. The return distribution η^π is the unique fixed point of this operator, i.e., $\mathcal{T}^\pi \eta^\pi = \eta^\pi$.

DBO's convergence is naturally analyzed under the *Wasserstein metric* w_p and its supremal version $\bar{w}_p(\eta, \eta') \triangleq \sup_{(x,a)} w_p(\eta(x, a), \eta'(x, a))$.

Theorem 1 (Contraction of DBO (Bellemare et al., 2017; Zhang et al., 2021))**.** *For any $p \geq 1$ and $\eta, \eta' \in \mathcal{P}_p(\mathbb{R}^d)^{\mathcal{X} \times \mathcal{A}}$, $\bar{w}_p(\mathcal{T}^\pi \eta, \mathcal{T}^\pi \eta') \leq \gamma \bar{w}_p(\eta, \eta')$.*

So $\eta_{k+1} = \mathcal{T}^\pi \eta_k$ converges to η^π in $\bar{w}_p$ at rate γ^k (Bellemare et al., 2023). $\mathcal{T}^\pi$ also contracts in supremal ℓ_p, in supremal MMD (Gretton et al., 2012) for suitable kernels, and in sup-sliced divergences (Nguyen-Tang et al., 2021; Wiltzer et al., 2024; Debes and Tuytelaars, 2026b), but its convergence in $\bar{w}_p$ motivates measuring approximation errors in $\bar{w}_p$.

Smoothed Scores, SVGD, and DSM. We define the smoothing operator H_σ as the convolution with a normal distribution, i.e., $\mathsf{H}_\sigma \mu \triangleq \mu * \mathcal{N}(0, \sigma^2 \mathbf{I})$. The score and smoothed score of μ are given as $s_\mu \triangleq \nabla \log \mu$, and $s_{\mu,\sigma} \triangleq \nabla \log(\mathsf{H}_\sigma \mu)$. Score matching is a minimum Stein discrepancy estimator (Barp

et al., 2019), and denoising score matching (DSM), the regression of $s_\psi(y+\sigma\varepsilon,\sigma)$ on $-\varepsilon/\sigma$ for $y\sim\mu$, $\varepsilon\sim\mathcal{N}(0,\mathbf{I})$, is score matching on the smoothed law $\mathsf{H}_\sigma\mu$ (Vincent, 2011). SVGD moves particles along $\hat{\phi}^*(z)=\frac{1}{M}\sum_j[k(z_j,z)\,s(z_j)+\nabla_{z_j}k(z_j,z)]$, the steepest descent direction of $\mathrm{KL}(\cdot\,\|\,p)$ in the unit ball of the RKHS of k for score $s=\nabla\log p$ (Liu and Wang, 2016).

Representations. We focus on particle representations, i.e., equal-weight Dirac mixtures of the form $\mathcal{F}_M\triangleq\{\frac{1}{M}\sum_{i=1}^M\delta_{z_i}:z_i\in\mathbb{R}^d\}$. Table 1 in Section 4.4 compares the critic families by their main error.

3 STEIN VARIATIONAL DISTRIBUTIONAL RL

We model the return distribution function through its smoothed score over all noise levels, the *return score field*, and use the Stein particles only to represent it. The smoothed score exists for every return law, and smoothing commutes with the affine bootstrap map up to a rescaling of the noise level (Lemma 1), which yields a Bellman equation in score space (Corollary 1).

The Return Score Field. The return score field of η is the map $(z,x,a,\sigma)\mapsto s_{\eta(x,a),\sigma}(z)$, which we learn with a noise-conditioned network $s_\psi(z,x,a,\sigma)$.

Lemma 1 (Bellman consistency)**.** *For any $\mu\in\mathcal{P}(\mathbb{R}^d)$, $r\in\mathbb{R}^d$, $\gamma\in(0,1)$, $\sigma>0$, we have*

$$\begin{aligned}&\textit{(i)}\ \ (b_{r,\gamma})_{\#}(\mathsf{H}_{\sigma/\gamma}\mu)=\mathsf{H}_\sigma\left((b_{r,\gamma})_{\#}\mu\right),\quad\textit{and}\\&\textit{(ii)}\ \ s_{(b_{r,\gamma})_{\#}\mu,\sigma}(y)=\frac{1}{\gamma}s_{\mu,\sigma/\gamma}\left(\frac{y-r}{\gamma}\right).\end{aligned}\tag{2}$$

Corollary 1 (Score Bellman equation; general branch mixture)**.** *Assume Assumption 1. Let $J=(R,X',A')\sim q(\cdot\mid x,a)$ index Bellman branches and $p_{j,\sigma}$ be the density of $(b_{r_j,\gamma})_{\#}\mathsf{H}_{\sigma/\gamma}\,\eta(x'_j,a'_j)$. Then, for all y and $\sigma>0$,*

$$\begin{aligned}s_{\mathcal{T}^\pi\eta(x,a),\sigma}(y)&=\mathbb{E}_{j\sim q_\sigma(\cdot|y)}\left[\tfrac{1}{\gamma}\,s_{\eta(x'_j,a'_j),\sigma/\gamma}\Big(\tfrac{y-r_j}{\gamma}\Big)\right]\\q_\sigma(j\mid y)&\propto q(j|x,a)\,p_{j,\sigma}(y).\end{aligned}\tag{3}$$

With a single deterministic branch, $q_\sigma(\cdot\mid y)$ is a point mass and (3) reduces to (2)(ii), our consistency target (5). Stochastic rewards are the case $J=R$.

We learn s_ψ by DSM on the bootstrap targets $y=r+\gamma z'$, with $a'\sim\pi(\cdot|x')$ and z' a particle at (x',a'), and tie the noise levels together by the score Bellman equation (3), with one sampled branch and one shared noise draw:

$$L_{\mathrm{DSM}}(\psi)=\mathbb{E}_{y,\sigma,\varepsilon}\left[\lambda(\sigma)\big\|s_\psi(y+\sigma\varepsilon,x,a,\sigma)+\varepsilon/\sigma\big\|^2\right],\tag{4}$$

Algorithm 1 One tabular SVDRL update at (x,a)

Require: particles $z_{1:M}$; score network s_ψ; hyperparameters ρ, σ_{start}, α

1: $y_{1:M}\leftarrow$ stratified Bellman targets $r_j+\gamma z'_j$
2: DSM and consistency steps (4)–(5); replace $\mathrm{round}(\rho M)$ particles by targets
3: $\tilde{z}\leftarrow$ SVGD along $s_\psi(\cdot,\sigma_\ell)$ from z, for $\sigma_\ell\le\sigma_{\mathrm{start}}$
4: $z\leftarrow z+\alpha(\tilde{z}-z)$, then map $z_{1:M}$ onto the moments of $y_{1:M}$

$$\begin{aligned}L_{\mathrm{cons}}(\psi)=\mathbb{E}\Big[\lambda(\sigma)\big\|&s_\psi(r+\gamma z+\sigma\varepsilon,x,a,\sigma)\\&-\tfrac{1}{\gamma}\,\mathrm{sg}\Big[s_{\bar{\psi}}(z+\tfrac{\sigma}{\gamma}\varepsilon,x',a',\tfrac{\sigma}{\gamma})\Big]\big\|^2\Big],\end{aligned}\tag{5}$$

with $\bar{\psi}$ an exponential moving average (EMA) of ψ. The pairing of σ with σ/γ is (2), and it makes score TD non-expansive (Theorem 2). Both losses stay unbiased under one-sample bootstrapping, the property (U) of Debes and Tuytelaars (2026b) (Lemmas 4 and 5). Tabular runs use the EDM preconditioning (Karras et al., 2022) in coordinates that standardize the targets at (x,a), with $\lambda(\sigma)=\sigma^2$ (Section B).

The Resample-Move Update. In the tabular experiments SVDRL keeps M particles per state–action pair and updates them by a resample-move step of sequential Monte Carlo (Gilks and Berzuini, 2001) in four steps (Algorithm 1). The *(i) stratified Bellman target set* draws the M successors by systematic resampling, so a branch of probability p receives $\lfloor Mp\rfloor$ or $\lceil Mp\rceil$ targets, and the reward noise is stratified within each successor block (Section B). *(ii) Resampling* replaces a fraction ρ of the particles by targets; $\rho=1$ carries the branch masses exactly up to $1/M$. The *(iii) Stein move* runs SVGD along s_ψ down the levels $\sigma_\ell\le\sigma_{\mathrm{start}}$ with relaxation α, and *(iv) moment matching* maps the cloud affinely onto the target set's moments. Resampling carries the mass, the Stein move the shape.

The Sampler and the Hybrid Rule. The sampler draws from the field by reverse diffusion of the EMA networks (Definition 2); one set of $N\gg M$ draws answers CVaR of any scalarization, event probabilities and conditional queries (Corollaries 10 and 11). Lemma 2 says the field is learnable when the return law has a density whose components are not far apart relative to their widths; we turn it into a rule that the agent evaluates on its own particles, without ground truth, with a threshold fixed in advance (Section B).

Definition 1 (Hybrid rule)**.** Let r be the ratio of the smallest gap between single-linkage clusters of the particle cloud at (x,a) to the largest within-cluster standard deviation, after merging particles closer than the field's smallest noise level. The hybrid rule reports the particles if $r>6$, or if a cluster of at least two

particles is narrower than the field's smallest noise level; otherwise it reports the sampler draws.

SAC-SVDRL. In continuous control SVDRL is the critic of SAC (Haarnoja et al., 2018): an amortized critic outputs M particles regressed onto the transported cloud, i.e., amortized SVGD (Feng et al., 2017), and a moment head sets the score network's standardization and moment-matching target (Section B).

4 CONVERGENCE ANALYSIS

All results concern policy evaluation for a fixed policy π under Assumption 1. We certify the reverse-diffusion sampler, for bounded and possibly multimodal returns, with no isoperimetric, dissipativity or mode-separation assumption; the Stein particles enter through one measurable term. Proofs are in Section A, and Section 5 checks Theorems 2 and 3 on trained networks.

Definition 2 (Noise grid, multi-level risk, reverse-diffusion sampler)**.** Let $\sigma_0 > \sigma_1 > \cdots > \sigma_{N_g} > 0$ be the noise grid on which the score network is trained, with variance steps $h_n \triangleq \sigma_n^2 - \sigma_{n+1}^2$, and let $\mathcal{I}(\nu) \triangleq \mathbb{E}_\nu\left[\|\nabla \log \nu\|^2\right]$ denote the Fisher information of a law with a smooth positive density. For probability measure μ with bounded support and score model $s_\psi(\cdot,\sigma)$ that is bounded on bounded sets, define the DSM excess risk at level σ (Vincent, 2011) and the *multi-level risk* by

$$\mathcal{E}_\sigma(\psi;\mu) \triangleq \mathbb{E}_{\mathsf{H}_\sigma \mu}\left[\|s_\psi(\cdot,\sigma) - s_{\mu,\sigma}\|^2\right],$$
$$\mathcal{E}_{\mathrm{ml}}(\psi;\mu) \triangleq \sum_{n=0}^{N_g-1} h_n\, \mathcal{E}_{\sigma_n}(\psi;\mu).$$

The reverse-diffusion sampler of s_ψ starts from $\hat{X}_0 \sim \mathcal{N}(z_c, \sigma_0^2 \mathbf{I})$, for a center $z_c \in B_R$ fixed before sampling (the standardizing mean at (x,a) in Algorithm 2), and iterates, for $n = 0,\dots,N_g-1$ and i.i.d. $\xi_n \sim \mathcal{N}(0,\mathbf{I})$, $\hat{X}_{n+1} = \hat{X}_n + h_n\, s_\psi(\hat{X}_n,\sigma_n) + \sqrt{h_n}\,\xi_n$. We write $\hat{Q}_\psi$ for the law of $\hat{X}_{N_g}$ and $Q_\psi \triangleq (\mathrm{proj}_{B_R})_{\#}\hat{Q}_\psi$ for the law of its projection onto the return box (Algorithm 2). At a state–action pair we take $\mu = \mu(x,a)$ and $s_\psi(\cdot,\sigma) = s_\psi(\cdot,x,a,\sigma)$.

4.1 Score-Space Temporal Differences

Theorem 2 (Score-space TD, informal)**.** *Assume a geometric grid with $\kappa_g = \gamma^{1/m}$ for an integer m and the Gaussian prior score $-(z-c)/\sigma^2$ above σ_0, with c the center of B_R. Write $\tilde{e}_n$ and $\tilde{r}_n$ for the score error and the consistency residual at level n, each normalized by σ_n and uniform over (x,a).* (i) *The score Bellman operator is non-expansive in the multi-level risk: the risk of its image is at most that of its input on the noise levels shifted up by m.* (ii) *At a fixed point of training, $\tilde{e}_n \le \sum_{i=0}^{\lfloor n/m \rfloor} \tilde{r}_{n-im} + D/(2\sigma_0)$ for every level n.*

By *(ii)* the consistency residuals, which need no ground truth, bound the error per level; exact sweeps converge in finitely many steps, with the noise schedule as discount factor (Theorem 5 and Corollaries 3 and 4).

4.2 A Bound on the Sampled Return Law

Theorem 3 (The multi-level risk controls the sampled return law)**.** *Let μ be supported in B_R and $s_\psi(\cdot,\sigma)$ be bounded on $\mathbb{R}^d$ for every grid level. Then*

$$\mathrm{KL}\left(\mathsf{H}_{\sigma_{N_g}}\mu \,\|\, \hat{Q}_\psi\right) \le \frac{m_2(\mu;z_c)}{2\sigma_0^2} + \mathcal{E}_{\mathrm{ml}}(\psi;\mu) + \varepsilon_{\mathrm{disc}}, \tag{6}$$
$$\textit{where } \varepsilon_{\mathrm{disc}} \triangleq \sum_{n=0}^{N_g-1} h_n\big(\mathcal{I}(\mathsf{H}_{\sigma_{n+1}}\mu) - \mathcal{I}(\mathsf{H}_{\sigma_n}\mu)\big) \ge 0.$$

Here $m_2(\mu;z_c) \triangleq \mathbb{E}_{Y\sim\mu}\left[\|Y - z_c\|^2\right]$ is the second moment of μ about the sampler's center, which is at most $\sup_{y\in B_R}\|y - z_c\|^2$, i.e., at most R_Z^2 for $z_c = 0$. It is measurable from the bootstrap targets, which are draws of μ. For a geometric grid, $\sigma_{n+1} = \kappa_g\sigma_n$ with $\kappa_g \in (0,1)$, $\varepsilon_{\mathrm{disc}} \le d\,\kappa_g^{-2}(1-\kappa_g^2)\big(1 + 2\log(\sigma_0/\sigma_{N_g})\big)$. Consequently, with $\varepsilon_{\mathrm{TV}}$ the square root of half the right side of (6),

$$\mathrm{TV}\big(\hat{Q}_\psi, \mathsf{H}_{\sigma_{N_g}}\mu\big) \le \varepsilon_{\mathrm{TV}} \qquad \textit{and}$$
$$w_1(Q_\psi,\mu) \le D\,\varepsilon_{\mathrm{TV}} + \sigma_{N_g}\sqrt{d}. \tag{7}$$

A clipping argument extends it to score networks of linear growth such as ours (Remark 3).

4.3 Policy Evaluation

Theorem 4 (Policy evaluation from the sampler and from Stein particles)**.** *For $k = 0,\dots,K-1$ let $\nu_k \in \mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$ be the bootstrap source of iteration k, let ψ_{k+1} be a score model toward the Bellman targets $\mu_k(x,a) \triangleq \mathcal{T}^\pi\nu_k(x,a)$ with $s_{\psi_{k+1}}(\cdot,\sigma)$ bounded on $\mathbb{R}^d$ at every grid level, and let $Q_{k+1}(x,a)$ be its sampler law. Let $\varepsilon_{\mathrm{TV},k}(x,a)$ be the bound of Theorem 3 for $\mu_k(x,a)$ and ψ_{k+1}, and $\varepsilon_{\mathrm{gen}} \triangleq D\sup_{k,(x,a)}\varepsilon_{\mathrm{TV},k}(x,a) + \sigma_{N_g}\sqrt{d}$.* (i) Sampler bootstrapping. *If $\nu_k = Q_k$ for every k, with $Q_0 \in \mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$ arbitrary, then $\bar{w}_1(Q_K,\eta^\pi) \le \gamma^K\,\bar{w}_1(Q_0,\eta^\pi) + \varepsilon_{\mathrm{gen}}/(1-\gamma)$.* (ii) Stein particles. *If $\nu_k = \hat{\eta}_k$ are the Stein particles of iteration k, then, with the Stein-particle gap*

$\varepsilon_{\mathrm{St}} \triangleq \sup_{k,(x,a)} w_1\big(\hat{\eta}_{k+1}(x,a), Q_{k+1}(x,a)\big)$,

$$\bar{w}_1(\hat{\eta}_K, \eta^\pi) \le \gamma^K \bar{w}_1(\hat{\eta}_0, \eta^\pi) + \frac{\varepsilon_{\mathrm{gen}} + \varepsilon_{\mathrm{St}}}{1-\gamma},$$
$$\bar{w}_1(Q_K, \eta^\pi) \le \varepsilon_{\mathrm{gen}} + \gamma\, \bar{w}_1(\hat{\eta}_{K-1}, \eta^\pi).$$

Every term is proved or measurable without ground truth, the multi-level risk through the consistency residuals at a fixed point of training (Corollary 5), and the analogues in $\bar{w}_2$ hold (Corollary 7). Sampler bootstrapping needs moment matching, since draws without it inflate the variance of the targets through the recursion, by $\gamma^2\sigma^2_{N_g}/(1-\gamma^2)$ per coordinate on a deterministic self-loop (Remark 5). The Stein particles enter *(ii)* only through the measured $\varepsilon_{\mathrm{St}}$: the exact-score SVGD drift is not globally dissipative for any bounded even C^1 kernel with bounded gradient (Proposition 4).

4.4 Three Sources of Error

A critic's error has three sources: the $1/M$ resolution floor of a finite particle set, the sampler's smoothing bias, and the regularity of the field it must learn.

Proposition 1 (What particles cannot answer, informal)**.** *Let $p \in (0, \frac{1}{2M})$ and let $\lambda \triangleq z_{\max}/\|z_{\max}\|_2$.* (i) *For every equal-weight M-atom law Q and every event A with target mass p, either $Q(A) = 0$ or $Q(A) > 2p$.* (ii) *There is a two-point return law η in B_R with a catastrophe of mass p at depth D along λ such that* every *equal-weight M-atom law Q, read by plug-in, errs by at least $\frac{D}{2}(1-2pM)$ in $\mathrm{CVaR}_\alpha(\lambda^\top Z)$ at $\alpha = p$ or at $\alpha = \frac{1}{M}$, and takes the wrong chance-constrained decision for every threshold δ in an interval of length at least p.* (iii) *With probability at least $1-\beta$ over N draws, the reverse-diffusion sampler estimates the probability of every half-space event, simultaneously for all scalarizations, with an error that does not depend on M: training risk $\varepsilon_{\mathrm{TV}}$, plus a smoothing term that is at most $\frac{1}{2}e^{-w^2/(2\sigma^2_{N_g})}$ when the slab of half-width w around the event boundary carries no target mass, plus $\mathcal{O}(\sqrt{(d\log N + \log(1/\beta))/N})$.*

(i) and *(ii)* hold for every particle critic, the Stein particles included, and *(i)* also rounds mixture weights to multiples of $1/M$. At rare-event scale the certificate is qualitative, since $\varepsilon_{\mathrm{TV}}$ can exceed p (Section 5); the claim is what controls the error: $1/M$ for atoms, the training risk and N for the sampler.

The sampler instead has a smoothing bias, first order on point masses and second order on laws with densities.

Proposition 2 (Smoothing bias)**.** *Let $\eta \in \mathcal{P}_2(\mathbb{R}^d)$ and $\sigma > 0$.* (i) $w_2(\mathsf{H}_\sigma\eta, \eta) \le \sigma\sqrt{d}$. (ii) *If $\eta = \sum_k w_k\delta_{y_k}$ with $\|y_j - y_k\| \ge \Delta$ for all $j \ne k$, then $w_2^2(\mathsf{H}_\sigma\eta, \eta) \ge \sigma^2 d\,\big(1 - \Pr(\chi^2_{d+2} \ge \Delta^2/(4\sigma^2))\big)$, where χ^2_{d+2} is a chi-squared variable with $d+2$ degrees of freedom; for a single atom, $w_2(\mathsf{H}_\sigma\delta_y, \delta_y) = \sigma\sqrt{d}$.* (iii) *If $\eta = \sum_k w_k\mathcal{N}(\mu_k, S_k)$ with $S_k \succeq s^2\mathbf{I}$ for every k, then $w_2(\mathsf{H}_\sigma\eta, \eta) \le \sqrt{d}\,\sigma^2/(2s)$.*

Table 1: *Main errors of the critic families.* The error that limits each family (Propositions 1 and 2 and Lemma 2). d>1: joint return law; <$1/M$: events rarer than one particle's mass. Particle critics with fresh or time-averaged targets differ only in sampling variance. Only SVDRL has both without a grid.

Critic	d>1	<$1/M$	Main error
Categorical grid	m^d bins	on grid	bin width
Quantile, IQN	–	IQN, one λ	d=1 only
Particle critics	✓	–	$1/M$ floor
Scalar generative	–	✓	d=1 only
SVDRL	✓	✓	smoothing, regularity

Part *(iii)* decides where the field can be learned.

Lemma 2 (Regularity of the smoothed score)**.** *Let $\eta \in \mathcal{P}(\mathbb{R}^d)$, $\sigma > 0$, $s_\sigma \triangleq s_{\eta,\sigma}$ and $\mathrm{Lip}(s_\sigma) \triangleq \sup_z \|\nabla s_\sigma(z)\|_{\mathrm{op}}$.* (i) *If $\operatorname{supp}\eta$ has diameter at most D, then $-\sigma^{-2}\mathbf{I} \preceq \nabla s_\sigma(z) \preceq \sigma^{-4}(D^2/4 - \sigma^2)\mathbf{I}$ for every z.* (ii) *If $\eta = \sum_{k=1}^K w_k\delta_{y_k}$ with $K \ge 2$ atoms at pairwise distances at least δ, attained by some pair, then $\liminf_{\sigma\to 0}\sigma^4\,\mathrm{Lip}(s_\sigma) \ge \delta^2/4$.* (iii) *If $\eta = \sum_k w_k\mathcal{N}(\mu_k, S)$ with $S \succeq s^2\mathbf{I}$ and the means in a set of diameter Δ, then $\mathrm{Lip}(s_\sigma) \le \max\{s^{-2}, \Delta^2/(4s^4)\}$ for every $\sigma > 0$.*

So the field of a law with densities is equally regular on every level, while the Lipschitz constant of an atomic law's field grows like σ^{-4} at the lowest noise levels, where a network of bounded regularity cannot fit it uniformly and Euler steps of the sampler become unstable (Section A). Particles need no regularity of the field. This is the split we observe in Section 5, and the hybrid rule of Definition 1 is its agent-side form.

5 EXPERIMENTS

Protocol. In the tabular experiments a seed's error is its mean over the last 20 % of training, which averages out the fluctuation of single evaluations. Configurations are selected on tuning seeds and reported on different seeds that the selection did not use, in the rare-event study additionally on a second set of fresh seeds. We report the interquartile mean (IQM) or median over seeds with 95 % bootstrap intervals; on the synthetic MDPs and the chain a comparison is a win or a loss when the Welch t statistic is at most -2 or at least 2, and in the rare-event study it compares IQM errors (Section C). The particle competitors are MMD critics

(Nguyen-Tang et al., 2021; Zhang et al., 2021) with RBF, MD3QN and energy kernels, sliced Cramér and sliced w_2 critics, the persistent-particle objectives of Debes and Tuytelaars (2026b) and Sinkhorn critics (Sun et al., 2024), each with a tuned learning rate and each also run with our stratified target set (Section B.1). The MMD critic with the energy kernel, energy MMD, is our own variant; its loss is the energy distance, in which the multivariate distributional Bellman operator contracts (Wiltzer et al., 2024).

Rare Events. In the rare-event chain (Section C) the risky root action leads, $D \in \{0, 2\}$ steps later, to a catastrophe of mass $p \in \{0.002, 0.005, 0.01\}$ in a two-dimensional return, and the queries on the joint law are CVaR down to $\alpha = 0.001$, event probabilities, chance constraints, conditional tails and entropic risk. Six particle critics answer from their particles or from a kernel density estimate on them, a quantile critic is trained on the oracle scalarization, and SVDRL bootstraps from moment-matched sampler draws (Theorem 4(ii), Remark 5) and answers with $N = 10^5$ draws.

At equal particle count, $M = 128$, SVDRL has a lower IQM error than the best particle critic in 60 of 60 rare-event and 12 of 12 bulk comparisons, and on fresh seeds its error is lower by a median factor of 5.0 on rare events and 3.9 on the bulk (Figure 2). At comparable compute, against particle critics with $M = 512$, it is still more accurate in 40 of 60 rare-event and 10 of 12 bulk comparisons. At $p = 0.005$ and $D = 2$ it is more accurate than every particle critic on every query of Table 5, from their atoms or from a kernel density estimate, and the quantile critic on the oracle scalarization matches it only at $\mathrm{CVaR}_{0.01}$ of that scalarization. Its $|\log \hat{P}/P|$ there is 0.16, against 4.8 to 8.5 for the particle critics. Particle critics need about $10/p$ particles to locate the catastrophe: with $M = 2048$, sliced w_2 reaches a $\mathrm{CVaR}_{0.001}$ error of 0.30 against SVDRL's 6.9. Even then SVDRL is more accurate on the mass ($|\log \hat{P}/P|$ 0.16 against 0.26), on the chance-constrained decisions (accuracy 1.00 against 0.96) and on the bulk ($\mathrm{CVaR}_{0.1}$ error 0.94 against 1.5). With sampler bootstrapping the catastrophe survives the recursion even at $M = 32$ for $p \geq 0.005$, where every particle critic loses it (Section D.1). This requires the rare mode to have a density.

Resource gathering from MO-Gymnasium (Felten et al., 2023) has atomic returns, with death as the rare event. At the standard attack probability $p = 0.1$ and at $p = 0.02$, SVDRL's particles ($M = 128$) are the best critic: their $|\log \hat{P}/P|$ for death is 0.028 and 0.15 against at best 0.19 and 0.45, their chance-constrained choice is wrong in 0 and 13 % of the cases against 14 and 41 %, and their CVaR error at $p = 0.1$ is 0.0035 against 0.030 (medians over ten seeds, Section D.4). The stratified target set holds each atom at its mass up to $1/M$, and at $p = 0.005 < 1/M$ full replacement rounds the mass without bias: the time-averaged estimate is 0.0053 for $P = 0.005$. The sampler fails on these atomic laws, as Lemma 2 predicts, and the hybrid rule matches the particles on the probability of death and on the chance-constrained choice.

Table 2: *Eight synthetic MDPs.* Within-mode w_2^2 in units of 10^{-3}, averaged over the last 20 % of training. SVDRL reports the hybrid of Definition 1. Ratios to the best competitor (Table 8); own: as specified; all: also on our target set; $M = 4096$: compute-matched, unimodal MDPs only. Bold: SVDRL is better by Welch's test. Mode masses in Table 7.

		Within-mode shape, w_2^2 (↓)			
			SVDRL / best competitor		
MDP	d	SVDRL	own	all	M=4096
MDP1	1	5.7	**0.23**	**0.27**	10
MDP2	1	81	**0.83**	1.6	–
MDP3	2	389	**0.20**	**0.80**	–
MDP4	5	62	**0.020**	**0.097**	–
MDP5	2	6.1	**0.17**	**0.20**	1.2
MDP6	2	50	**0.054**	**0.12**	–
MDP7	3	1.5	**0.19**	**0.27**	4.1
MDP8	2	38	**0.19**	**0.23**	2.3
Wins/ties/losses			8/0/0	7/0/1	0/0/4
its particles alone			6/1/1	4/1/3	0/0/4

Eight Synthetic MDPs. We split the error into the total variation of the mode masses (Table 7) and the within-mode w_2^2 (Table 2), because the full w_2^2 charges a mass error ϵ across a gap Δ as $\epsilon\Delta^2$ and is dominated by rounding when modes are far apart (Section D.2). On within-mode shape the hybrid scores 8/0/0 wins/ties/losses against the competitors as specified and 7/0/1 against every row, competitors given our target set included; a single configuration for all eight MDPs, without per-MDP selection, gives the same counts, and the separation test alone ($r > 6$) makes the same choice of object on every MDP (Table 6). Its particles alone score 6/1/1 and 4/1/3 in the same two framings. On mass every method sits at the $1/M$ limit (0/2/2), as Proposition 1(i) predicts: SVDRL allocates within one particle of the weights and rounds stochastically, so its mass is unbiased over time (bias 0.0013 on MDP2, against a floor of 0.0031 at which persistent clouds get stuck).

On the four unimodal MDPs the hybrid reports $N = 4096$ sampler draws against $M = 64d$ particles, and at matched compute these wins vanish: competitors with $M = 4096$ beat the sampler on all four (0/0/4).

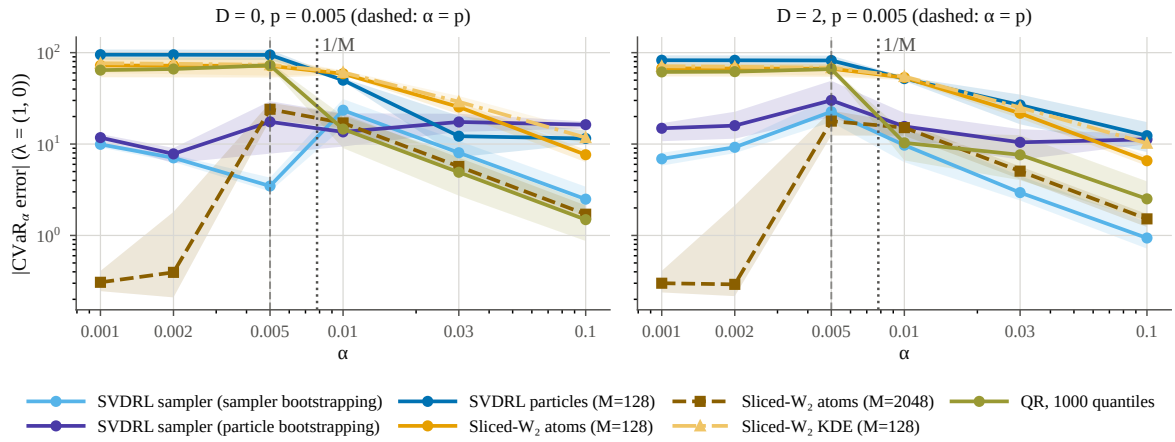

Figure 2: *Rare events.* Rare-event chain, catastrophe of mass $p = 0.005$ at depth $D = 0$ (left) and $D = 2$ (right): CVaR_α error of Z_1. Below $1/M$, particle critics with $M = 128$ miss the catastrophe and SVDRL's sampler errs an order of magnitude less; sliced w_2 needs $M = 2048$ to locate it. IQM over ten seeds, 95 % bootstrap intervals.

The ablations and the checks of the theory support every component (Sections D.2 and D.6). On MDP1, MDP5 and MDP8 the Stein move beats Langevin dynamics along the same score, a Gaussian jitter and no move by factors of 1.19 to 1.39, and without repulsion the cloud collapses on MDP5 and MDP8, with 28.6 and 204 times the error (Table 13). Removing the consistency loss raises the error significantly on four MDPs and never lowers it, and a single-level network misallocates the mode masses (0.32 against 0.056 on MDP2; Remark 2). On trained networks Theorem 2(ii) holds in 87 of 87 seeds and the bound of Theorem 3 in 80 of 80 cases, but loosely: $\varepsilon_{\mathrm{TV}}$ exceeds the measured total variation 4.4 to 34 times. The ground-truth-free gap $\varepsilon_{\mathrm{St}}$ of Theorem 4 equals the particles' true error within 1 to 12 % on the unimodal MDPs.

Chain with Reward Noise. The chain of Debes and Tuytelaars (2026b) has ten reward dimensions, to which we add Gaussian reward noise of scale σ_r. On the atomic chain, $\sigma_r = 0$, their persistent particles win, as Lemma 2 predicts. With reward noise, $\sigma_r \geq 0.05$, and against their critics at their published settings with $M = 128$ instead of 500 particles, SVDRL's moment-matched sampler (4096 draws) wins every shape metric and the mean (Figure 3), e.g., with 0.15 times the best competitor's energy distance and 0.56 times its mean error at $\sigma_r = 0.2$. With more particles they win shape at matched compute: with $M = 4096$ their shape errors are 1.6 to 3.0 times lower than the sampler's. Our stratified target set lowers the energy distance and mean error of every one of their variants, and running-mean moment matching, a shift of their particles to a running mean of the bootstrap targets, lets them win the mean too, so both of our components transfer to their critics (Section D.3).

Continuous Control. On four MuJoCo environments (Towers et al., 2024), SAC with the SVDRL critic runs against SAC and against IQN, MD3QN and energy MMD critics on the same trunk (Table 3). Its return is on par with SAC's, with overlapping 95 % intervals on all four environments, although its IQM on HalfCheetah is lower. Its critic distribution is the most accurate on three of four: its centered w_2, which removes the twin-critic bias of the means, is the lowest on HalfCheetah, Walker2d and Hopper, with the interval entirely below IQN's, and second on Ant, behind energy MMD. With the default configuration of HalfCheetah it is 0.70 on Walker2d and 0.57 on Ant (Table 18), the same ranks.

6 RELATED WORK

Multivariate Distributional RL. MD3QN (Zhang et al., 2021) trains particles with an MMD loss, SinkhornDRL (Sun et al., 2024) with entropic optimal transport, and sliced DRL (Debes and Tuytelaars, 2026b) with sliced divergences, whose max-sliced version also contracts under matrix discounting. Their analyses establish the contraction of the Bellman operator in the respective divergence, and Wiltzer et al. (2024) prove convergence of categorical dynamic programming and TD learning on a fixed support; our guarantees concern a learned score-based representation of possibly multimodal returns (Theorems 3 and 4), also under matrix discounting (Corollary 9). Yoo et al. (2026) extend quantile regression to vector returns through the Knothe–Rosenblatt map.

Generative Critics. Diffusion and flow critics such as VDRL (Hu et al., 2025), DBC (Ding et al., 2026), Value Flows (Dong et al., 2026) and FlowIQN (Groom

Table 3: *Continuous control.* MuJoCo after 10^6 steps, IQM over five to ten seeds with its 95 % bootstrap interval; bold: inside the interval of the best IQM. Return: each run's mean evaluation return over the last 20 % of training. Critic shape: centered w_2 between the critic and Monte Carlo returns of the frozen policy, over their standard deviation. SVDRL's configurations in Section D.5.

	Return (↑)					Critic shape (↓)			
	SVDRL	SAC	IQN	MD3QN	E-MMD	SVDRL	IQN	MD3QN	E-MMD
HalfCheetah	8720[7565, 9889]	**10 014**[9446, 10 512]	**9219**[8338, 9879]	**8986**[8449, 9877]	**10 098**[8885, 11 186]	**0.43**[0.28, 0.56]	1.05[0.82, 1.41]	0.69[0.50, 0.92]	**0.52**[0.39, 0.71]
Walker2d	**5514**[5025, 5780]	**5372**[4939, 5744]	5083[4466, 5496]	**5661**[5129, 5907]	**5609**[5361, 5805]	**0.92**[0.61, 1.31]	3.99[2.32, 6.13]	**1.19**[0.90, 2.75]	2.08[0.82, 7.51]
Hopper	**2187**[2047, 2294]	**2187**[2083, 2319]	2019[1841, 2221]	1822[1743, 1889]	**2180**[2047, 2255]	**2.25**[0.87, 4.56]	10.36[6.69, 20.64]	6.10[1.69, 14.74]	21.26[10.69, 44.55]
Ant	2429[2059, 2654]	2426[2214, 2759]	2411[2095, 3029]	1233[895, 2343]	**2928**[2628, 3598]	0.59[0.57, 0.63]	0.73[0.42, 1.19]	0.72[0.56, 2.95]	**0.49**[0.45, 0.56]

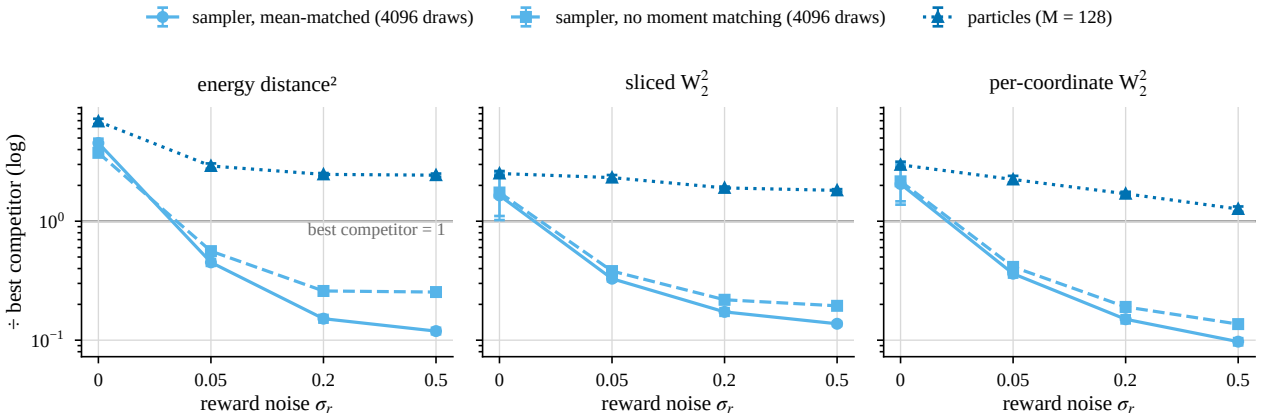

Figure 3: *Chain with reward noise.* SVDRL's error over the best competitor's ($M = 128$) per metric against the reward noise σ_r, for its sampler and particles; below one, SVDRL is better. The sampler loses on the atomic chain and wins for $\sigma_r \geq 0.05$, until competitors get $M = 4096$ particles. Means over ten seeds, 95 % intervals.

et al., 2026) model scalar returns, and DBC and Flow-IQN rely on one-dimensional structure, inverse distribution functions or sorted couplings. Bellman Diffusion (Li et al., 2024) learns unnormalized density fields of scalar returns on which the Bellman operator acts linearly, and Corollary 1 is the normalized counterpart in score space, which is nonlinear through the posterior branch weights. The Bellman GAN (Freirich et al., 2019) samples vector-valued returns and distributional Sobolev training (Debes and Tuytelaars, 2026a) returns with their action gradients, without a density or a score. None of these works bounds the error of a learned multivariate representation.

Particle Transport. SVGD (Liu and Wang, 2016) is a kernelized Wasserstein gradient flow of the KL divergence (Liu, 2017; Chewi et al., 2020), and its convergence guarantees rest on functional inequalities of the target, with constants that grow exponentially with the separation of modes (Proposition 3(i)). In RL, SVGD has so far trained policies (Liu et al., 2017; Messaoud et al., 2024), not return distributions.

7 LIMITATIONS AND CONCLUSION

The main errors of Table 1 also mark where SVDRL is not the best critic. The sampler fails on atomic and widely separated laws, such as the atomic chain, resource gathering and the far modes of MDP4, as Lemma 2 predicts, and the hybrid rule reports the particles there. A light rare mode far from the bulk comes out too wide (Figure 5), and a static fit of the score network to draws of the true law reproduces this, with 13 to 15 times the true width (Section E). At equal N the sampler loses to fitted particle clouds (Section 5). In SAC-SVDRL a linear-variance floor made one Ant seed diverge, which a log-variance head removes on HalfCheetah and a standardized moment head limits on multi-objective MuJoCo.

SVDRL learns a Bellman-consistent score field that yields a certified sampler without a resolution floor and Stein particles. It answers rare-event queries that particles cannot represent, and two of its components improve every particle critic we tested. At equal particle count it is more accurate than the best particle critic in 60 of 60 rare-event comparisons, and in continuous control SAC-SVDRL learns the most accurate critic distribution on three of four MuJoCo environments.

AI Use Statement

In this work, we used generative AI tools for the following tasks. It was used to optimize the original implementation of the methods by the authors to improve performance and experimental throughput. All generated code was reviewed by the authors, compared numerically against the original implementation and an extensive set of unit tests. We further used these tools during the writing process as follows. Drafting of sections, typesetting the proofs by the authors into their final form, and grammar checks. The final form of the proofs was checked manually by the authors to ensure they remained consistent. Whenever a tool suggested the addition of a reference in any of the sections, it was flagged and manually checked for correctness by the authors. To ensure that results could not inadvertently be changed by the tool during typesetting, all are generated automatically from the logged experimental data. Any error in the proofs as well as incorrect attribution of theoretical or empirical results is therefore due to the authors and we take full responsibility for it and everything written in the paper.

References

Anderson, B. D. O. (1982). Reverse-time diffusion equation models. *Stochastic Processes and their Applications*.

Ba, J., Erdogdu, M. A., Ghassemi, M., Sun, S., Suzuki, T., Wu, D., and Zhang, T. (2022). Understanding the variance collapse of SVGD in high dimensions. In *International Conference on Learning Representations*.

Barp, A., Briol, F.-X., Duncan, A., Girolami, M., and Mackey, L. (2019). Minimum Stein discrepancy estimators. In *Advances in Neural Information Processing Systems*.

Bellemare, M. G., Dabney, W., and Munos, R. (2017). A distributional perspective on reinforcement learning. In *International Conference on Machine Learning*.

Bellemare, M. G., Dabney, W., and Rowland, M. (2023). *Distributional Reinforcement Learning*. The MIT Press.

Benton, J., De Bortoli, V., Doucet, A., and Deligiannidis, G. (2024). Nearly d-linear convergence bounds for diffusion models via stochastic localization. In *International Conference on Learning Representations*.

Boucheron, S., Lugosi, G., and Massart, P. (2013). *Concentration Inequalities*. Oxford University Press.

Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Katariya, Y., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., and Zhang, Q. (2018). JAX: Composable transformations of Python+NumPy programs. http://github.com/jax-ml/jax.

Chen, S., Chewi, S., Li, J., Li, Y., Salim, A., and Zhang, A. R. (2023). Sampling is as easy as learning the score: Theory for diffusion models with minimal data assumptions. In *International Conference on Learning Representations*.

Chewi, S., Le Gouic, T., Lu, C., Maunu, T., and Rigollet, P. (2020). SVGD as a kernelized Wasserstein gradient flow of the chi-squared divergence. In *Advances in Neural Information Processing Systems*.

Dabney, W., Ostrovski, G., Silver, D., and Munos, R. (2018a). Implicit quantile networks for distributional reinforcement learning. In *International Conference on Machine Learning*.

Dabney, W., Rowland, M., Bellemare, M., and Munos, R. (2018b). Distributional reinforcement learning with quantile regression. *AAAI Conference on Artificial Intelligence*.

Debes, B. and Tuytelaars, T. (2026a). Distributional value gradients for stochastic environments. In *International Conference on Learning Representations*.

Debes, B. and Tuytelaars, T. (2026b). Multivariate distributional reinforcement learning using sliced divergences. In *International Conference on Machine Learning*.

Devroye, L., Györfi, L., and Lugosi, G. (1996). *A Probabilistic Theory of Pattern Recognition*. Springer.

Ding, S., Zhou, Y., Hu, K., Pan, M., Zhong, S., Fu, Y., Wang, J., and Shi, Y. (2026). Distributional reinforcement learning with diffusion bridge critics. *arXiv preprint arXiv:2602.05783*.

Dong, P., Zheng, C., Finn, C., Sadigh, D., and Eysenbach, B. (2026). Value flows. In *International Conference on Learning Representations*.

Duncan, A., Nüsken, N., and Szpruch, L. (2023). On the geometry of Stein variational gradient descent. *Journal of Machine Learning Research*.

Felten, F., Alegre, L. N., Nowé, A., Bazzan, A. L. C., Talbi, E. G., Danoy, G., and Silva, B. C. da. (2023). A toolkit for reliable benchmarking and research in multi-objective reinforcement learning. In *Advances in Neural Information Processing Systems*.

Feng, Y., Wang, D., and Liu, Q. (2017). Learning to draw samples with amortized Stein variational gradient descent. In *Conference on Uncertainty in Artificial Intelligence*.

Freirich, D., Shimkin, T., Meir, R., and Tamar, A. (2019). Distributional multivariate policy evaluation and exploration with the Bellman GAN. In *International Conference on Machine Learning*.

Gilks, W. R. and Berzuini, C. (2001). Following a moving target—Monte Carlo inference for dynamic Bayesian models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*.

Gretton, A., Borgwardt, K. M., Rasch, M. J., Schölkopf, B., and Smola, A. (2012). A kernel two-sample test. *The Journal of Machine Learning Research*.

Groom, M., Darvariu, V.-A., Kunze, L., Wilson, J., and Hawes, N. (2026). Quantile-coupled flow matching for distributional reinforcement learning. *arXiv preprint arXiv:2605.08515*.

Haarnoja, T., Zhou, A., Abbeel, P., and Levine, S. (2018). Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International Conference on Machine Learning*.

Haussmann, U. G. and Pardoux, É. (1986). Time reversal of diffusions. *The Annals of Probability*.

Hayes, C. F., Rădulescu, R., Bargiacchi, E., Källström, J., Macfarlane, M., Reymond, M., Verstraeten, T., Zintgraf, L. M., Dazeley, R., Heintz, F., Howley, E., Irissappane, A. A., Mannion, P., Nowé, A., Ramos, G., Restelli, M., Vamplew, P., and Roijers, D. M. (2022). A practical guide to multi-objective reinforcement learning and planning. *Autonomous Agents and Multi-Agent Systems*.

Hu, X., Wang, F., Zhang, T., and Cui, Z. (2025). Value diffusion reinforcement learning. In *Advances in Neural Information Processing Systems*.

Karras, T., Aittala, M., Aila, T., and Laine, S. (2022). Elucidating the design space of diffusion-based generative models. In *Advances in Neural Information Processing Systems*.

Kidger, P. (2022). jaxtyping. https://github.com/patrick-kidger/jaxtyping.

Kidger, P. and Garcia, C. (2021). Equinox: Neural networks in JAX via callable PyTrees and filtered transformations. In *Differentiable Programming Workshop at Neural Information Processing Systems*.

Kingma, D. P. and Ba, J. (2015). Adam: A method for stochastic optimization. In *International Conference on Learning Representations*.

Kitagawa, G. (1996). Monte Carlo filter and smoother for non-Gaussian nonlinear state space models. *Journal of Computational and Graphical Statistics*.

Li, Y., Lai, C.-H., Schönlieb, C.-B., Mitsufuji, Y., and Ermon, S. (2024). Bellman diffusion: Generative modeling as learning a linear operator in the distribution space. *arXiv preprint arXiv:2410.01796*.

Liu, Q. (2017). Stein variational gradient descent as gradient flow. In *Advances in Neural Information Processing Systems*.

Liu, Q. and Wang, D. (2016). Stein variational gradient descent: A general purpose Bayesian inference algorithm. In *Advances in Neural Information Processing Systems*.

Liu, Y., Ramachandran, P., Liu, Q., and Peng, J. (2017). Stein variational policy gradient. In *Conference on Uncertainty in Artificial Intelligence*.

Massart, P. (1990). The tight constant in the Dvoretzky-Kiefer-Wolfowitz inequality. *The Annals of Probability*.

McKay, M. D., Beckman, R. J., and Conover, W. J. (1979). A comparison of three methods for selecting values of input variables in the analysis of output from a computer code. *Technometrics*.

Messaoud, S., Mokeddem, B., Xue, Z., Pang, L., An, B., Chen, H., and Chawla, S. (2024). S^2AC: Energy-based reinforcement learning with Stein soft actor critic. In *International Conference on Learning Representations*.

Nguyen-Tang, T., Gupta, S., and Venkatesh, S. (2021). Distributional reinforcement learning via moment matching. *AAAI Conference on Artificial Intelligence*.

Rockafellar, R. T. and Uryasev, S. (2002). Conditional value-at-risk for general loss distributions. *Journal of Banking & Finance*.

Song, Y. and Ermon, S. (2019). Generative modeling by estimating gradients of the data distribution. In *Advances in Neural Information Processing Systems*.

Sun, K., Zhao, Y., Liu, W., Jiang, B., and Kong, L. (2024). Distributional reinforcement learning with regularized Wasserstein loss. In *Advances in Neural Information Processing Systems*.

Towers, M., Kwiatkowski, A., Terry, J., Balis, J. U., De Cola, G., Deleu, T., Goulão, M., Kallinteris, A., Krimmel, M., KG, A., et al. (2024). Gymnasium: A standard interface for reinforcement learning environments. *arXiv preprint arXiv:2407.17032*.

Vauthier, C., Mérigot, Q., and Korba, A. (2026). Highly data parallelizable estimation of the sliced-Wasserstein distance using cumulative distribution functions. *arXiv preprint arXiv:2606.30310*.

Vincent, P. (2011). A connection between score matching and denoising autoencoders. *Neural Computation*.

Wiltzer, H., Farebrother, J., Gretton, A., and Rowland, M. (2024). Foundations of multivariate distributional reinforcement learning. In *Advances in Neural Information Processing Systems*.

Yoo, G., Kim, W. K., and Woo, H. (2026). Knothe-Rosenblatt quantile regression for risk-sensitive multi-

objective reinforcement learning. In *International Conference on Machine Learning*.

Zhang, P., Chen, X., Zhao, L., Xiong, W., Qin, T., and Liu, T.-Y. (2021). Distributional reinforcement learning for multi-dimensional reward functions. In *Advances in Neural Information Processing Systems*.

Supplementary Material

A PROOFS

This section proves the results of Sections 3 and 4 in the order of the main text, together with the corollaries that Section 5 uses. We assume Assumption 1 throughout.

A.1 Notation and Tools

Definition 3 (Wasserstein-p distance)**.** For $p \geq 1$ and laws ν, ν' on $\mathbb{R}^d$ with finite p-th moments, $w_p(\nu, \nu') \triangleq \left(\inf_{\iota \in \Pi(\nu,\nu')} \int \|x - y\|_2^p \, \mathrm{d}\iota\right)^{1/p}$, where $\Pi(\nu, \nu')$ is the set of couplings of ν and ν', and $\bar{w}_p(\eta, \eta') \triangleq \sup_{(x,a)} w_p(\eta(x, a), \eta'(x, a))$ for return distribution functions η, η'.

Posterior Form and Hessian. With φ_σ the density of $\mathcal{N}(0, \sigma^2\mathbf{I})$, where $\sigma > 0$, write $p_\sigma(z) = \int \varphi_\sigma(z - y)\mu(\mathrm{d}y)$ for the density of $\mathsf{H}_\sigma\mu$, and $Z = Y + \sigma\varepsilon$ with $Y \sim \mu$, $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$ independent. Differentiating under the integral with $\nabla_z\varphi_\sigma(z - y) = -\sigma^{-2}(z - y)\varphi_\sigma(z - y)$ and recognizing the posterior weights $\varphi_\sigma(z - y)\mu(\mathrm{d}y)/p_\sigma(z)$ gives

$$s_{\mu,\sigma}(z) = \frac{\nabla p_\sigma(z)}{p_\sigma(z)} = -\frac{z - m_\sigma(z)}{\sigma^2}, \tag{8}$$

where $m_\sigma(z) = \mathbb{E}[Y \mid Z = z]$, and $\mu(\mathrm{d}y \mid z) \propto \varphi_\sigma(z - y)\mu(\mathrm{d}y)$. Since that posterior is an exponential family in z with natural parameter z/σ^2, we have $\nabla_z m_\sigma(z) = \sigma^{-2}\operatorname{Cov}[Y \mid Z = z]$, and differentiating (8) once more gives

$$\nabla^2 \log p_\sigma(z) = -\sigma^{-2}\mathbf{I} + \sigma^{-4}\operatorname{Cov}[Y \mid Z = z]. \tag{9}$$

Denoising Score Matching. The DSM loss $L_{\mathrm{DSM}}(\psi; \sigma) = \mathbb{E}_{Y,\varepsilon}\left[\|s_\psi(Y + \sigma\varepsilon, \sigma) + \varepsilon/\sigma\|^2\right]$ satisfies Vincent's identity (Vincent, 2011)

$$L_{\mathrm{DSM}}(\psi; \sigma) = \mathbb{E}_{Z \sim \mathsf{H}_\sigma\mu}\left[\left\|s_\psi(Z, \sigma) - s_{\mu,\sigma}(Z)\right\|^2\right] + C_{\mathrm{DSM}}(\sigma), \tag{10}$$

with $C_{\mathrm{DSM}}(\sigma)$ independent of ψ, because $-\varepsilon/\sigma = -(Z - Y)/\sigma^2$ has the conditional mean $s_{\mu,\sigma}(z)$ given $Z = z$ by (8), and $C_{\mathrm{DSM}}(\sigma)$ is its conditional variance. So the excess risk of DSM is $\mathcal{E}_\sigma$ of Definition 2.

Gaussian Norms and CVaR. For $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$ the map $\varepsilon \mapsto \|\varepsilon\|$ is 1-Lipschitz with $\mathbb{E}[\|\varepsilon\|] \leq \sqrt{d}$, so Gaussian concentration (Boucheron et al., 2013) gives

$$\Pr(\|\varepsilon\| \geq \sqrt{d} + u) \leq e^{-u^2/2}, \qquad u \geq 0. \tag{11}$$

Lemma 3 (CVaR utility facts)**.** *Let $\alpha \in (0, 1)$ and write the lower-tail CVaR of a law P on $\mathbb{R}$ with finite first moment as $\mathrm{CVaR}_\alpha(P) = \sup_c g_P(c)$, $g_P(c) \triangleq c - \alpha^{-1}\mathbb{E}_P[(c - X)_+]$ (Rockafellar and Uryasev, 2002).* (i) $|\mathrm{CVaR}_\alpha(P_1) - \mathrm{CVaR}_\alpha(P_2)| \leq \alpha^{-1}w_1(P_1, P_2)$. (ii) *If P_1, P_2 are supported in a closed interval of length at most D, then $|\mathrm{CVaR}_\alpha(P_1) - \mathrm{CVaR}_\alpha(P_2)| \leq (D/\alpha)\|F_{P_1} - F_{P_2}\|_\infty$.* (iii) *For laws ν_1, ν_2 on $\mathbb{R}^d$ and $\|\lambda\|_2 \leq 1$, $w_1(\lambda_\#\nu_1, \lambda_\#\nu_2) \leq w_1(\nu_1, \nu_2)$.* (iv) *If P_1, P_2 are supported in a set of diameter D, then $w_1(P_1, P_2) \leq D\,\mathrm{TV}(P_1, P_2)$ and $w_2(P_1, P_2) \leq D\,\mathrm{TV}(P_1, P_2)^{1/2}$.*

Proof. (i) As $x \mapsto (c - x)_+$ is 1-Lipschitz, $|g_{P_1}(c) - g_{P_2}(c)| \leq \alpha^{-1}\mathbb{E}[|X_1 - X_2|]$ for every coupling and every c; take the infimum over couplings. *(ii)* For P supported in $I = [a, b]$ the supremum over c is attained in I, and $\mathbb{E}_P[(c - X)_+] = \int_a^c F_P(x)\,\mathrm{d}x$ for $c \in I$ by Fubini's theorem, which differs between P_1 and P_2 by at most $D\|F_{P_1} - F_{P_2}\|_\infty$. *(iii)* A coupling (Z_1, Z_2) gives the coupling $(\lambda^\top Z_1, \lambda^\top Z_2)$ with $\mathbb{E}\left[|\lambda^\top(Z_1 - Z_2)|\right] \leq \mathbb{E}[\|Z_1 - Z_2\|]$. *(iv)* With the common part $\nu \triangleq P_1 \wedge P_2$, of mass $1 - \mathrm{TV}(P_1, P_2)$, the coupling $(\mathrm{id}, \mathrm{id})_\#\nu + (P_1 - \nu) \otimes (P_2 - \nu)/\mathrm{TV}(P_1, P_2)$ moves only the mass $\mathrm{TV}(P_1, P_2)$, by at most D. □

A.2 The Score Bellman Equation

Proof of Lemma 1. (i) For independent $Y \sim \mu$ and $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$ we have $Y + \frac{\sigma}{\gamma}\varepsilon \sim \mathsf{H}_{\sigma/\gamma}\mu$ and $b_{r,\gamma}(Y + \frac{\sigma}{\gamma}\varepsilon) = (r + \gamma Y) + \sigma\varepsilon \sim \mathsf{H}_\sigma((b_{r,\gamma})_\#\mu)$. *(ii)* For every law ν and $\tau > 0$, $\mathsf{H}_\tau\nu$ has the smooth, strictly positive density $\int \varphi_\tau(\cdot - w)\,\nu(\mathrm{d}w)$, so its score exists everywhere. By (i) and the change of variables $y = r + \gamma u$, the density of $\mathsf{H}_\sigma((b_{r,\gamma})_\#\mu)$ is $y \mapsto \gamma^{-d} p_{\mathsf{H}_{\sigma/\gamma}\mu}((y-r)/\gamma)$; under $\nabla_y \log$ the constant γ^{-d} drops out, and the chain rule gives the factor $1/\gamma$. □

Proof of Corollary 1. By (1), $\mathcal{T}^\pi\eta(x,a)$ is the $q(\cdot \mid x, a)$-mixture of the branch laws $(b_{r_j,\gamma})_\#\eta(x'_j, a'_j)$, and H_σ is linear, so by Lemma 1(i) the density of $\mathsf{H}_\sigma\mathcal{T}^\pi\eta(x,a)$ is $p_\sigma = \int p_{j,\sigma}\, q(\mathrm{d}j \mid x, a)$. Differentiating under the integral,

$$\nabla \log p_\sigma(y) = \int \frac{p_{j,\sigma}(y)\, q(\mathrm{d}j \mid x, a)}{p_\sigma(y)} \nabla \log p_{j,\sigma}(y) = \mathbb{E}_{j \sim q_\sigma(\cdot \mid y)} \left[\nabla \log p_{j,\sigma}(y)\right],$$

and Lemma 1(ii) gives $\nabla \log p_{j,\sigma}(y) = \gamma^{-1} s_{\eta(x'_j,a'_j),\sigma/\gamma}((y - r_j)/\gamma)$. For discrete branches the exchange is a finite sum; in general it follows by dominated convergence for every η, since by Lemma 1(i) $p_{j,\sigma}$ is the convolution of φ_σ with the probability measure $(b_{r_j,\gamma})_\#\eta(x'_j, a'_j)$, so that $p_{j,\sigma} \le \sup \varphi_\sigma$ and $\|\nabla p_{j,\sigma}\| \le \sup \|\nabla \varphi_\sigma\| < \infty$, uniformly in j and y, and $p_\sigma > 0$. □

The score Bellman equation defines the operator that the consistency loss applies to the target network.

Definition 4 (Score Bellman image)**.** Let $\nu \in \mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$ be a bootstrap source, $\mu \triangleq \mathcal{T}^\pi\nu$, and $q_\sigma(\mathrm{d}j \mid y)$ the posterior branch weights of Corollary 1 for the source ν at (x, a). The *score Bellman image* of a measurable field $f(\cdot; x', a', \sigma')$ is

$$(\mathcal{S}_\nu f)(y; x, a, \sigma) \triangleq \int q_\sigma(\mathrm{d}j \mid y)\, \gamma^{-1} f\Big(\frac{y - r_j}{\gamma}; x'_j, a'_j, \frac{\sigma}{\gamma}\Big),$$

so that $\mathcal{S}_\nu s_{\nu,\cdot} = s_{\mu,\cdot}$ by Corollary 1.

Corollary 2 (Commutation under invertible matrix discounting)**.** *Let $\Gamma \in \mathbb{R}^{d\times d}$ be invertible and $b_{r,\Gamma}(z) \triangleq r + \Gamma z$. Then for every $\mu \in \mathcal{P}(\mathbb{R}^d)$ and $\sigma > 0$,*

$$\begin{aligned}
&\textit{(i)}\ (b_{r,\Gamma})_\#\Big(\mu * \mathcal{N}\big(0, \sigma^2(\Gamma^\top\Gamma)^{-1}\big)\Big) = \mathsf{H}_\sigma\big((b_{r,\Gamma})_\#\mu\big),\ \textit{and}\\
&\textit{(ii)}\ s_{(b_{r,\Gamma})_\#\mu,\sigma}(y) = \Gamma^{-\mathsf{T}}\nabla \log\big(\mu * \mathcal{N}(0, \sigma^2(\Gamma^\top\Gamma)^{-1})\big)\big(\Gamma^{-1}(y - r)\big).
\end{aligned}$$

Proof. For independent $Z \sim \mu$ and $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$, $r + \Gamma(Z + \sigma\Gamma^{-1}\varepsilon) = r + \Gamma Z + \sigma\varepsilon$ and $\sigma\Gamma^{-1}\varepsilon \sim \mathcal{N}(0, \sigma^2(\Gamma^\top\Gamma)^{-1})$, which gives (i); (ii) follows from (i) as in the proof of Lemma 1. □

Under matrix discounting the predecessor is smoothed anisotropically, so the multi-level analysis of Section A.4, which pairs the levels σ and σ/γ, is for a scalar γ, whereas the policy-evaluation bound needs no such pairing (Corollary 9).

A.3 One-Sample Bootstrap Targets

Lemma 4 (One-sample DSM is (U))**.** *Fix (x, a) and $\sigma > 0$, let $J = (R, X', A') \sim q(\cdot \mid x, a)$ index the Bellman branches as in Corollary 1, with branch targets $\mu_j \triangleq (b_{r_j,\gamma})_\#\eta(x'_j, a'_j)$ and mixture target $\mu \triangleq \mathcal{T}^\pi\eta(x, a)$, and let $L^{(j)}_{\mathrm{DSM}}(\psi; \sigma)$ be the DSM loss (4) with $Y \sim \mu_j$. Then*

$$\mathbb{E}_{j \sim q(\cdot \mid x,a)}\left[L^{(j)}_{\mathrm{DSM}}(\psi; \sigma)\right] = \mathbb{E}_{Z \sim \mathsf{H}_\sigma\mu}\left[\big\|s_\psi(Z, \sigma) - s_{\mu,\sigma}(Z)\big\|^2\right] + C_{\mathrm{DSM}}(\sigma), \tag{12}$$

with $C_{\mathrm{DSM}}(\sigma)$ independent of ψ. So (i) *the population minimizer of the one-sample bootstrap DSM objective over measurable fields is $s_{\mathcal{T}^\pi\eta(x,a),\sigma}$, the posterior-weighted branch mixture (3), and* (ii) *whenever ∇_ψ commutes with the expectations in (12), the sampled-branch gradient $\nabla_\psi L^{(J)}_{\mathrm{DSM}}(\psi; \sigma)$ is an unbiased estimate of $\nabla_\psi \mathbb{E}_{\mathsf{H}_\sigma\mu}\left[\|s_\psi - s_{\mu,\sigma}\|^2\right]$.*

Proof of Lemma 4. Drawing $J \sim q(\cdot \mid x,a)$ and then $Y \sim \mu_J$ gives $Y \sim \mu$, so by Tonelli's theorem the left side of (12) is the DSM loss for μ, and Vincent's identity (10) gives (12). Its right side is minimized exactly by $s_{\mu,\sigma}$, which Corollary 1 identifies, and (ii) is the exchange of ∇_ψ with the expectation over J; the exchange holds if $\nabla_\psi s_\psi(z,\sigma)$ grows at most polynomially in $\|z\|$, locally uniformly in ψ, since $\mathsf{H}_\sigma\mu$ has Gaussian tails. □

Lemma 5 (One-sample consistency is unbiased)**.** *Fix (x,a), $\sigma>0$, a bootstrap source $\nu \in \mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$, $\mu \triangleq \mathcal{T}^\pi\nu$, and a measurable field f. Draw $j=(r,x',a') \sim q(\cdot \mid x,a)$, then, given j, draw $z \sim \nu(x',a')$ and $\xi \sim \mathcal{N}(0,\mathbf{I})$ independently, and set*

$$u \triangleq z + \tfrac{\sigma}{\gamma}\,\xi, \qquad y \triangleq r+\gamma u, \qquad T \triangleq \gamma^{-1} f\big(u;x',a',\tfrac{\sigma}{\gamma}\big).$$

Then $y \sim \mathsf{H}_\sigma\mu(x,a)$ and $\mathbb{E}[T \mid y] = (\mathcal{S}_\nu f)(y;x,a,\sigma)$ almost surely. If moreover $\mathbb{E}\left[\|T\|^2\right] < \infty$, then for every measurable s with $\mathbb{E}_{\mathsf{H}_\sigma\mu(x,a)}\left[\|s\|^2\right] < \infty$,

$$\mathbb{E}\left[\|s(y)-T\|^2\right] = \mathbb{E}_{\mathsf{H}_\sigma\mu(x,a)}\left[\big\|s-(\mathcal{S}_\nu f)(\cdot;x,a,\sigma)\big\|^2\right] + \mathbb{E}\left[\|T-\mathbb{E}[T\mid y]\,\|^2\right], \tag{13}$$

where the last term does not depend on s and vanishes when $q(\cdot \mid x,a)$ is a point mass.

Proof of Lemma 5. Write $\nu_j \triangleq \nu(x'_j,a'_j)$. Given j, $u \sim \mathsf{H}_{\sigma/\gamma}\nu_j$, so $y = r_j+\gamma u$ has the density p_j of $(b_{r_j,\gamma})_\#\mathsf{H}_{\sigma/\gamma}\nu_j$, and (j,y) has the joint law $p_j(y)\,q(\mathrm{d}j \mid x,a)\,\mathrm{d}y$. By Lemma 1 the y-marginal is $\mathsf{H}_\sigma\mu(x,a)$, and the law of j given y is the posterior branch law $q_\sigma(\mathrm{d}j \mid y)$ of Corollary 1. Since $u=(y-r_j)/\gamma$, T is a function of (j,y), so $\mathbb{E}[T\mid y] = \int q_\sigma(\mathrm{d}j \mid y)\,\gamma^{-1}f((y-r_j)/\gamma;x'_j,a'_j,\sigma/\gamma) = (\mathcal{S}_\nu f)(y;x,a,\sigma)$, and (13) is the orthogonality of $T-\mathbb{E}[T\mid y]$ to the square-integrable functions of y. If $q(\cdot \mid x,a)$ is a point mass, T is a function of y alone. □

Both lemmas obtain the posterior branch weights without evaluating a density, because y is drawn through the sampled branch. A target drawn independently of the branch would estimate the prior-weighted average of the branch scores, which is not the score of $\mathsf{H}_\sigma\mu(x,a)$ unless the posterior weights are constant.

Remark 1 (Target variance of the two losses). The one-sample DSM target $-\xi/\sigma$ at $y=y_0+\sigma\xi$, $y_0\sim\mu(x,a)$, has by Tweedie's formula $\mathbb{E}[\xi\mid y] = -\sigma s_{\mu,\sigma}(y)$ the normalized conditional variance $\sigma^2\,\mathbb{E}\left[\|\xi/\sigma-\mathbb{E}[\xi/\sigma\mid y]\,\|^2\right] = d-\sigma^2\mathcal{I}(\mathsf{H}_\sigma\mu)$, with the Fisher information $\mathcal{I}$ of Definition 2, which tends to d as $\sigma\to0$ whenever $\mathcal{I}(\mu)<\infty$. The consistency target of Lemma 5 with the exact field $f=s_{\nu,\cdot}$ has

$$\sigma^2\,\mathbb{E}\left[\|T-\mathbb{E}[T\mid y]\,\|^2\right] \le \sigma^2\,\mathbb{E}\left[\|T\|^2\right] = \frac{\sigma^2}{\gamma^2}\int \mathcal{I}\big(\mathsf{H}_{\sigma/\gamma}\nu(x'_j,a'_j)\big)\,q(\mathrm{d}j \mid x,a) \le \frac{\sigma^2}{\gamma^2}\sup_j \mathcal{I}(\nu(x'_j,a'_j)),$$

which is $\mathcal{O}(\sigma^2)$ when the successor laws have finite Fisher information; both bounds use $\mathcal{I}(\mathsf{H}_\tau\rho)\le\mathcal{I}(\rho)$, i.e., Jensen's inequality for the smoothed score as the posterior mean of the unsmoothed one. So the two targets stand to each other as a temporal-difference and a Monte Carlo target: at the lowest noise levels the consistency target has a far smaller variance, and its price is the accumulation of residuals in Corollary 3.

A.4 Score-Space Temporal Differences

This subsection states Theorem 2 in full and proves it. For a measurable field f and a law ρ we write $\mathcal{E}_\sigma(f;\rho) \triangleq \mathbb{E}_{\mathsf{H}_\sigma\rho}\left[\|f(\cdot,\sigma)-s_{\rho,\sigma}\|^2\right]$, as in Definition 2, and $\bar{\mathcal{E}}_\sigma(f;\nu) \triangleq \sup_{x',a'}\mathcal{E}_\sigma(f(\cdot;x',a',\sigma);\nu(x',a'))$.

Theorem 5 (Score-space TD is non-expansive in the multi-level risk)**.** *Let $\nu\in\mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$ be a bootstrap source, $\mu\triangleq\mathcal{T}^\pi\nu$, and $\mathcal{S}_\nu$ the score Bellman image of Definition 4.* (i) One level. *For every (x,a) and $\sigma>0$,*

$$\mathcal{E}_\sigma\big(\mathcal{S}_\nu f;\mu(x,a)\big) \le \gamma^{-2}\int \mathcal{E}_{\sigma/\gamma}\big(f;\nu(x'_j,a'_j)\big)\,q(\mathrm{d}j \mid x,a) \le \gamma^{-2}\,\bar{\mathcal{E}}_{\sigma/\gamma}(f;\nu), \tag{14}$$

with equality in the first inequality whenever $q(\cdot \mid x,a)$ is a point mass. (ii) All levels. *On the geometric grid $\sigma_{n+1}=\kappa_{\mathrm{g}}\sigma_n$ with $\kappa_{\mathrm{g}}=\gamma^{1/m}$, extended upward to the levels $\sigma_{-m}>\cdots>\sigma_{-1}>\sigma_0$, one has $h_{n-m}=h_n/\gamma^2$ and*

$$\sum_{n=0}^{N_{\mathrm{g}}-1} h_n\,\mathcal{E}_{\sigma_n}\big(\mathcal{S}_\nu f;\mu(x,a)\big) \le \sum_{n=-m}^{N_{\mathrm{g}}-1-m} h_n\,\bar{\mathcal{E}}_{\sigma_n}(f;\nu). \tag{15}$$

(iii) Boundary. *If $f(z;x',a',\sigma) = -(z-c)/\sigma^2$ with c the center of B_R, then $\sigma^2\mathcal{E}_\sigma(f;\rho)\le D^2/(4\sigma^2)$ for every ρ supported in B_R and every $\sigma>0$.*

Proof of Theorem 5. Part (i). Fix (x,a) and σ, and write $\nu_j \triangleq \nu(x'_j,a'_j)$ and p_j for the density of $(b_{r_j,\gamma})_{\#}\mathsf{H}_{\sigma/\gamma}\nu_j$, so that $\mathsf{H}_\sigma\mu(x,a)$ has the density $p=\int p_j\,q(\mathrm{d}j\mid x,a)$ (Lemma 1), $q_\sigma(\mathrm{d}j\mid y)=p_j(y)\,q(\mathrm{d}j\mid x,a)/p(y)$, and $s_{\mu,\sigma}=\int q_\sigma(\mathrm{d}j\mid\cdot)\nabla\log p_j$ (Corollary 1). With $e_j(y)\triangleq\gamma^{-1}\big(f-s_{\nu_j,\sigma/\gamma}\big)\big((y-r_j)/\gamma;x'_j,a'_j,\sigma/\gamma\big)$ we have $(\mathcal{S}_\nu f)(y)-s_{\mu,\sigma}(y)=\int q_\sigma(\mathrm{d}j\mid y)\,e_j(y)$, and Jensen's inequality for the kernel $q_\sigma(\cdot\mid y)$ gives

$$\mathcal{E}_\sigma(\mathcal{S}_\nu f;\mu(x,a))\le\int\Big(\int\|e_j(y)\|^2p_j(y)\,\mathrm{d}y\Big)q(\mathrm{d}j\mid x,a)=\gamma^{-2}\int\mathcal{E}_{\sigma/\gamma}(f;\nu_j)\,q(\mathrm{d}j\mid x,a),$$

since $p_j(y)\,\mathrm{d}y$ is the law of $r_j+\gamma u$ for $u\sim\mathsf{H}_{\sigma/\gamma}\nu_j$. The second inequality of (14) is the supremum over (x',a'), and for a point mass $q(\cdot\mid x,a)$ Jensen's inequality holds with equality.

Part (ii). With $\kappa_{\mathrm{g}}=\gamma^{1/m}$, $\sigma_{n-m}=\sigma_n/\gamma$ and $h_{n-m}=(1-\kappa_{\mathrm{g}}^2)\sigma_{n-m}^2=h_n/\gamma^2$, so (14) at $\sigma=\sigma_n$ times h_n gives $h_n\mathcal{E}_{\sigma_n}(\mathcal{S}_\nu f;\mu(x,a))\le h_{n-m}\bar{\mathcal{E}}_{\sigma_{n-m}}(f;\nu)$; sum over n.

Part (iii). By (8), $f(z)-s_{\rho,\sigma}(z)=(m_\sigma(z)-c)/\sigma^2$ with the posterior mean $m_\sigma(z)$ in the convex hull of $\operatorname{supp}\rho\subseteq B_R$, a box of circumradius $D/2$ centered at c, so $\|f(z)-s_{\rho,\sigma}(z)\|\le D/(2\sigma^2)$ for every z. □

In the normalized error $\sigma\mathcal{E}_\sigma^{1/2}$, the unit in which scores scale, the operator is an isometry with one branch and a contraction with several, and the level weights $h_n\propto\sigma_n^2$ are the squares of this normalization. This is the sense of Theorem 2(i): the multi-level risk of the image is at most that of the input on the levels shifted up by m.

Corollary 3 (The consistency residual bounds the score error)**.** *In the setting of Theorem 5, let s_ψ be the online and $s_{\bar\psi}$ the target score model, and let*

$$\mathcal{R}_n(x,a)\triangleq\mathbb{E}_{\mathsf{H}_{\sigma_n}\mu(x,a)}\Big[\big\|s_\psi(\cdot;x,a,\sigma_n)-(\mathcal{S}_\nu s_{\bar\psi})(\cdot;x,a,\sigma_n)\big\|^2\Big]$$

be the consistency residual at level n. With the normalized quantities $\tilde e_n(\psi;\rho)\triangleq\sigma_n\bar{\mathcal{E}}_{\sigma_n}(s_\psi;\rho)^{1/2}$ and $\tilde r_n\triangleq\sigma_n\sup_{x,a}\mathcal{R}_n(x,a)^{1/2}$,

$$\tilde e_n(\psi;\mu)\le\tilde r_n+\tilde e_{n-m}(\bar\psi;\nu)\qquad\text{for all } n\ge0.\tag{16}$$

At a fixed point of training, $\bar\psi=\psi$ and $\nu=\mu=\eta^\pi$, and with the prior score of Theorem 5(iii) above σ_0, iterating (16) gives, with $L_n\triangleq\lfloor n/m\rfloor$,

$$\tilde e_n(\psi;\eta^\pi)\le\sum_{i=0}^{L_n}\tilde r_{n-im}+\frac{D}{2\sigma_0}.$$

So exact consistency at every level and every state–action pair determines the score field of η^π up to the normalized error $D/(2\sigma_0)$ at every noise level.

Proof of Corollary 3. The triangle inequality in $L^2(\mathsf{H}_{\sigma_n}\mu(x,a))$ and Theorem 5(i) with $f=s_{\bar\psi}$ give $\mathcal{E}_{\sigma_n}(s_\psi;\mu(x,a))^{1/2}\le\mathcal{R}_n(x,a)^{1/2}+\gamma^{-1}\bar{\mathcal{E}}_{\sigma_{n-m}}(s_{\bar\psi};\nu)^{1/2}$, as $\sigma_n/\gamma=\sigma_{n-m}$; the supremum over (x,a) times σ_n is (16). At the fixed point, iterating (16) L_n+1 times reaches the index $n-(L_n+1)m<0$, a level $\sigma\in(\sigma_0,\sigma_0/\gamma]$ where Theorem 5(iii) bounds the normalized error of the prior score by $D/(2\sigma)\le D/(2\sigma_0)$. □

By Lemma 5 with $f=s_{\bar\psi}$, $\mathcal{R}_n(x,a)$ is the population consistency loss (5) at level n minus a conditional variance that does not depend on ψ, and equals it for point-mass branch laws. So the gradient of the one-sample consistency loss is an unbiased estimate of the gradient of $\mathcal{R}_n$, the residuals need no ground truth and are bounded by the training objective, and Theorem 2(ii) is the iterated bound of Corollary 3.

Corollary 4 (Score-space TD converges in finitely many sweeps)**.** *Let the branch laws be point masses, so that $\mathcal{S}\triangleq\mathcal{S}_\nu$ does not depend on ν. Let f_0 be any field on the grid that equals the prior score of Theorem 5(iii) above σ_0, and let $f_{k+1}\triangleq\mathcal{S}f_k$ on the grid, the field above σ_0 being held fixed, with the normalized per-sweep residuals $\tilde r_n^{(k)}\triangleq\sigma_n\sup_{x,a}\|f_k(\cdot;x,a,\sigma_n)-(\mathcal{S}f_{k-1})(\cdot;x,a,\sigma_n)\|_{L^2(\mathsf{H}_{\sigma_n}\eta^\pi(x,a))}$ of an inexact sweep ($\tilde r\equiv0$ for the exact iteration). Then for every level n and every $k\ge L_n+1$, with $L_n=\lfloor n/m\rfloor$,*

$$\tilde e_n(f_k;\eta^\pi)\le\sum_{i=0}^{L_n}\tilde r_{n-im}^{(k-i)}+\frac{D}{2\sigma_0},$$

whatever f_0 is on the grid. So the exact iteration reaches all levels within $\lceil N_{\mathrm{g}}/m\rceil=\lceil\log(\sigma_0/\sigma_{N_{\mathrm{g}}})/\log(1/\gamma)\rceil$ sweeps, with a normalized error of at most $D/(2\sigma_0)$, and the residual of an inexact sweep is charged at most L_n+1 times, never accumulated over sweeps.

Proof of Corollary 4. With point-mass branch laws the proof of (16) uses only Theorem 5(i), so it applies to the pair (f_k, f_{k-1}) with $\nu = \mu = \eta^\pi$ and gives $\tilde{e}_n(f_k) \le \tilde{r}_n^{(k)} + \tilde{e}_{n-m}(f_{k-1})$. Unrolling it $L_n + 1$ times reaches a level above σ_0, where every f_j is the prior score, whose normalized error is at most $D/(2\sigma_0)$ by Theorem 5(iii). □

For stochastic branch laws $\mathcal{S}_\nu$ moves with the source ν, and Corollary 3 at a fixed point replaces the finite-sweep statement. The statement needs a lagged target: with the online network as its own target, every update moves the target through the parameters shared across noise levels, which breaks the triangular order of the sweeps, and consistency-only training then diverges on a self-looping state–action pair (Section D.6).

A.5 The Sampler

Remark 2 (Why one noise level does not suffice). No isoperimetric constant enters because the sampler is compared with the reference process whose marginals $\mathsf{H}_{\sigma_n}\mu$ are exactly the measures under which DSM certifies s_ψ. At a single level the DSM risk is blind to mode masses: for two modes at separation $\Delta \ge 2\sigma$ whose weights are perturbed by ϵ, the risk at level σ is $\mathcal{O}(\epsilon^2 e^{-\Delta^2/(8\sigma^2)})$ (Proposition 3(ii)). By (6) and Pinsker's inequality, a weight error in the sampler law must instead show up at the levels $\sigma \gtrsim \Delta$ (Section D.6 checks this on trained networks). The grid must start above the spread of the target, $\sigma_0^2 \gg m_2(\mu; z_c)$, and $\varepsilon_{\mathrm{disc}}$ vanishes as $\kappa_{\mathrm{g}} \to 1$.

In this subsection μ is supported in $B_R \subseteq B(0, R_Z)$, $X_0 \sim \mu$, and $(W_u)_{u\ge 0}$ is a standard Brownian motion on $\mathbb{R}^d$ independent of X_0. For $\sigma > 0$ we write $X^{(\sigma)} \triangleq X_0 + W_{\sigma^2}$, so that $X^{(\sigma)} \sim \mathsf{H}_\sigma\mu$.

Lemma 6 (The score along the heat flow is a reverse martingale). *For $\sigma > 0$ let $\mathcal{G}_\sigma$ be the σ-algebra generated by $\{X^{(\sigma')} : \sigma' \ge \sigma\}$. Then, for all $0 < \sigma \le \sigma'$,*

$$\mathbb{E}\left[s_{\mu,\sigma}\big(X^{(\sigma)}\big) \,\middle|\, \mathcal{G}_{\sigma'}\right] = s_{\mu,\sigma'}\big(X^{(\sigma')}\big) \qquad \text{and} \qquad \mathbb{E}\left[\left\|s_{\mu,\sigma}\big(X^{(\sigma)}\big) - s_{\mu,\sigma'}\big(X^{(\sigma')}\big)\right\|^2\right] = \mathcal{I}(\mathsf{H}_\sigma\mu) - \mathcal{I}(\mathsf{H}_{\sigma'}\mu).$$

The map $\sigma \mapsto \mathcal{I}(\mathsf{H}_\sigma\mu)$ is nonincreasing, and $\mathcal{I}(\mathsf{H}_\sigma\mu) \le d/\sigma^2$.

Proof of Lemma 6. By the posterior form (8),

$$s_{\mu,\sigma}\big(X^{(\sigma)}\big) = -\frac{X^{(\sigma)} - \mathbb{E}\left[X_0 \mid X^{(\sigma)}\right]}{\sigma^2} = -\frac{\mathbb{E}\left[W_{\sigma^2} \mid X^{(\sigma)}\right]}{\sigma^2}.$$

Jensen's inequality gives $\mathbb{E}\left[\|s_{\mu,\sigma}(X^{(\sigma)})\|^2\right] \le \mathbb{E}\left[\|W_{\sigma^2}\|^2\right]/\sigma^4 = d/\sigma^2$, so $\mathcal{I}(\mathsf{H}_\sigma\mu) \le d/\sigma^2$ and all second moments below are finite. The σ-algebra $\mathcal{G}_\sigma$ is generated by $X^{(\sigma)}$ and the increments $(W_u - W_{\sigma^2})_{u\ge\sigma^2}$, which are independent of (X_0, W_{σ^2}), so $\mathbb{E}\left[W_{\sigma^2} \mid X^{(\sigma)}\right] = \mathbb{E}\left[W_{\sigma^2} \mid \mathcal{G}_\sigma\right]$.

Now fix $\sigma < \sigma'$ and let $\mathcal{H}$ be the σ-algebra generated by X_0, $W_{\sigma'^2}$ and $(W_u - W_{\sigma'^2})_{u\ge\sigma'^2}$, which contains $\mathcal{G}_{\sigma'}$. The pair $(W_{\sigma^2}, W_{\sigma'^2})$ is independent of X_0 and of the later increments, so the Brownian bridge gives $\mathbb{E}\left[W_{\sigma^2} \mid \mathcal{H}\right] = \mathbb{E}\left[W_{\sigma^2} \mid W_{\sigma'^2}\right] = (\sigma^2/\sigma'^2)\, W_{\sigma'^2}$. We have

$$\begin{aligned}
\mathbb{E}\left[s_{\mu,\sigma}\big(X^{(\sigma)}\big) \,\middle|\, \mathcal{G}_{\sigma'}\right] &= -\frac{\mathbb{E}\left[W_{\sigma^2} \mid \mathcal{G}_{\sigma'}\right]}{\sigma^2} && \text{(tower property, } \mathcal{G}_{\sigma'} \subseteq \mathcal{G}_\sigma\text{)} \\
&= -\frac{\mathbb{E}\left[\mathbb{E}\left[W_{\sigma^2} \mid \mathcal{H}\right] \mid \mathcal{G}_{\sigma'}\right]}{\sigma^2} && \text{(tower property, } \mathcal{G}_{\sigma'} \subseteq \mathcal{H}\text{)} \\
&= -\frac{\mathbb{E}\left[W_{\sigma'^2} \mid \mathcal{G}_{\sigma'}\right]}{\sigma'^2} = s_{\mu,\sigma'}\big(X^{(\sigma')}\big),
\end{aligned}$$

proving the first claim. So the increment $s_{\mu,\sigma}(X^{(\sigma)}) - s_{\mu,\sigma'}(X^{(\sigma')})$ is orthogonal in $L^2(\mathbb{P})$ to the $\mathcal{G}_{\sigma'}$-measurable $s_{\mu,\sigma'}(X^{(\sigma')})$, and Pythagoras' identity with $\mathbb{E}\left[\|s_{\mu,\sigma}(X^{(\sigma)})\|^2\right] = \mathcal{I}(\mathsf{H}_\sigma\mu)$ gives the second claim, which implies the monotonicity. □

The lemma is the heat-flow form of the stochastic-localization argument that Benton et al. (2024, Lemmas 1 to 4) use for the Ornstein–Uhlenbeck process, where the same control of the discretization error is obtained by Itô calculus.[1]

[1] They remark that their analysis extends to the variance-exploding process.

Proof of Theorem 3. Reverse process. Let $T \triangleq \sigma_0^2 - \sigma_{N_g}^2$, $\sigma(t) \triangleq (\sigma_0^2 - t)^{1/2}$ for $t \in [0, T]$, and $t_n \triangleq \sigma_0^2 - \sigma_n^2$, so that $\sigma(t_n) = \sigma_n$ and $t_{n+1} - t_n = h_n$. The process $Y_t \triangleq X^{(\sigma(t))}$ is the heat flow on the variance interval $[\sigma_{N_g}^2, \sigma_0^2]$, a Brownian motion started from $\mathsf{H}_{\sigma_{N_g}}\mu$, read backwards in time. By the time-reversal theorem (Anderson, 1982; Haussmann and Pardoux, 1986) it solves $\mathrm{d}Y_t = s_{\mu,\sigma(t)}(Y_t)\,\mathrm{d}t + \mathrm{d}B_t$ with a Brownian motion B, $Y_0 \sim \mathsf{H}_{\sigma_0}\mu$, and $Y_T \sim \mathsf{H}_{\sigma_{N_g}}\mu$. The sampler of Definition 2 is the Euler–Maruyama scheme of this equation with the learned score frozen on the grid: the process $\mathrm{d}\hat{X}_t = s_\psi(\hat{X}_{t_n}, \sigma_n)\,\mathrm{d}t + \mathrm{d}B_t$ on $[t_n, t_{n+1})$, with $\hat{X}_0 \sim \mathcal{N}(z_c, \sigma_0^2\mathbf{I})$, satisfies $\hat{X}_{t_{n+1}} = \hat{X}_{t_n} + h_n s_\psi(\hat{X}_{t_n}, \sigma_n) + (B_{t_{n+1}} - B_{t_n})$, which is the recursion of Definition 2, so $\hat{X}_T \sim \hat{Q}_\psi$.

Girsanov and data processing. Let P^Y and $P^{\hat{X}}$ be the path laws of Y and $\hat{X}$ on $[0, T]$, and let $P^{\hat{X}}_{\mathrm{ref}}$ be the path law of the scheme started from $\mathsf{H}_{\sigma_0}\mu$ instead of $\mathcal{N}(z_c, \sigma_0^2\mathbf{I})$. The two scheme laws share their transition kernels, so the chain rule for the relative entropy gives

$$\mathrm{KL}\left(P^Y \,\|\, P^{\hat{X}}\right) = \mathrm{KL}\left(\mathsf{H}_{\sigma_0}\mu \,\|\, \mathcal{N}(z_c, \sigma_0^2\mathbf{I})\right) + \mathrm{KL}\left(P^Y \,\|\, P^{\hat{X}}_{\mathrm{ref}}\right).$$

For the second term, Girsanov's theorem for two diffusions with a common initial law, in the form obtained by the approximation argument of Chen et al. (2023, inequality (5.5) and Section 5.2),[2] gives

$$\mathrm{KL}\left(P^Y \,\|\, P^{\hat{X}}_{\mathrm{ref}}\right) \le \frac{1}{2}\sum_{n=0}^{N_g-1} \int_{t_n}^{t_{n+1}} \mathbb{E}\left[\left\|s_{\mu,\sigma(t)}(Y_t) - s_\psi(Y_{t_n}, \sigma_n)\right\|^2\right] \mathrm{d}t,$$

whenever the right side is finite, which the bounds below show. By data processing for the time-T marginals, $\mathrm{KL}\left(\mathsf{H}_{\sigma_{N_g}}\mu \,\|\, \hat{Q}_\psi\right) \le \mathrm{KL}\left(P^Y \,\|\, P^{\hat{X}}\right)$.

Initialization. By joint convexity of the Kullback–Leibler divergence, $\mathrm{KL}\left(\mathsf{H}_{\sigma_0}\mu \,\|\, \mathcal{N}(z_c, \sigma_0^2\mathbf{I})\right) \le \int \mathrm{KL}\left(\mathcal{N}(y, \sigma_0^2\mathbf{I}) \,\|\, \mathcal{N}(z_c, \sigma_0^2\mathbf{I})\right)\, \mu(\mathrm{d}y) = m_2(\mu; z_c)/(2\sigma_0^2)$.

Score and discretization errors. For $t \in [t_n, t_{n+1})$ we have

$$\begin{aligned}
\frac{1}{2}\mathbb{E}\left[\left\|s_{\mu,\sigma(t)}(Y_t) - s_\psi(Y_{t_n}, \sigma_n)\right\|^2\right] &\le \mathbb{E}\left[\left\|s_{\mu,\sigma_n}(Y_{t_n}) - s_\psi(Y_{t_n}, \sigma_n)\right\|^2\right] \\
&\quad + \mathbb{E}\left[\left\|s_{\mu,\sigma(t)}(Y_t) - s_{\mu,\sigma_n}(Y_{t_n})\right\|^2\right] \qquad (\|a+b\|^2 \le 2\|a\|^2 + 2\|b\|^2) \\
&= \mathcal{E}_{\sigma_n}(\psi; \mu) + \mathcal{I}(\mathsf{H}_{\sigma(t)}\mu) - \mathcal{I}(\mathsf{H}_{\sigma_n}\mu) \qquad (\text{Lemma 6}) \\
&\le \mathcal{E}_{\sigma_n}(\psi; \mu) + \mathcal{I}(\mathsf{H}_{\sigma_{n+1}}\mu) - \mathcal{I}(\mathsf{H}_{\sigma_n}\mu).
\end{aligned}$$

The equality uses $Y_{t_n} \sim \mathsf{H}_{\sigma_n}\mu$ for the first term and $(Y_{t_n}, Y_t) = (X^{(\sigma_n)}, X^{(\sigma(t))})$ with $\sigma(t) \le \sigma_n$ for the second, and the last step is the monotonicity in Lemma 6, as $\sigma(t) \ge \sigma_{n+1}$. Integrating over $[t_n, t_{n+1})$, an interval of length h_n, and summing over n proves (6).

Geometric grid. With $\sigma_{n+1} = \kappa_g \sigma_n$ we have $h_n = (1 - \kappa_g^2)\sigma_n^2$, and summation by parts gives

$$\sum_{n=0}^{N_g-1} \sigma_n^2\big(\mathcal{I}(\mathsf{H}_{\sigma_{n+1}}\mu) - \mathcal{I}(\mathsf{H}_{\sigma_n}\mu)\big) \le \sigma_{N_g-1}^2\,\mathcal{I}(\mathsf{H}_{\sigma_{N_g}}\mu) + \sum_{n=1}^{N_g-1} (\sigma_{n-1}^2 - \sigma_n^2)\,\mathcal{I}(\mathsf{H}_{\sigma_n}\mu) \le \frac{d}{\kappa_g^2}\big(1 + (N_g - 1)(1 - \kappa_g^2)\big),$$

by $\mathcal{I}(\mathsf{H}_\sigma\mu) \le d/\sigma^2$ and $\sigma_{n-1}^2 = \sigma_n^2/\kappa_g^2$. As $1 - e^{-2x} \le 2x$, $(N_g - 1)(1 - \kappa_g^2) \le 2N_g \log(1/\kappa_g) = 2\log(\sigma_0/\sigma_{N_g})$, and multiplying by $1 - \kappa_g^2$ gives the bound on $\varepsilon_{\mathrm{disc}}$.

Total variation and w_1. Pinsker's inequality, $\mathrm{TV} \le (\mathrm{KL}/2)^{1/2}$, gives the first part of (7). The projection proj_{B_R} is 1-Lipschitz, as B_R is closed and convex, and it fixes μ, so $w_1(Q_\psi, \mu)$ is at most $w_1\big((\mathrm{proj}_{B_R})_\#\hat{Q}_\psi, (\mathrm{proj}_{B_R})_\#\mathsf{H}_{\sigma_{N_g}}\mu\big) + w_1\big(\mathsf{H}_{\sigma_{N_g}}\mu, \mu\big)$. Both laws in the first term are supported in B_R, of diameter D, so Lemma 3(iv) and data processing bound it by $D\,\mathrm{TV}(\hat{Q}_\psi, \mathsf{H}_{\sigma_{N_g}}\mu) \le D\varepsilon_{\mathrm{TV}}$, and the coupling $(Y + \sigma_{N_g}\varepsilon, Y)$ bounds the second by $\sigma_{N_g}\mathbb{E}\,[\|\varepsilon\|] \le \sigma_{N_g}\sqrt{d}$. □

[2] The argument avoids Novikov's condition. For a diffusion coefficient g instead of their $\sqrt{2}$ the bound is $\frac{1}{2}\mathbb{E}\left[\int \|(b_t^Y - b_t^X)/g\|^2\,\mathrm{d}t\right]$ for the drifts b^Y, b^X; here $g = 1$.

Remark 3 (Score models of linear growth). The score network of Section B is not bounded on $\mathbb{R}^d$: the prior term $-\|z\|^2/(2(\sigma^2+\sigma_{\mathrm{d}}^2))$ of its energy makes its score grow linearly in z. Theorem 3 covers it through a clipped model. For a radius $R_{\mathrm{c}} > 0$ let $\tilde{s}_\psi(z,\sigma) \triangleq s_\psi(\mathrm{proj}_{\bar{B}(z_c,R_{\mathrm{c}})}(z),\sigma)$, which is bounded on $\mathbb{R}^d$ at every grid level, as s_ψ is continuous, and equals s_ψ on the closed ball $\bar{B}(z_c,R_{\mathrm{c}})$. Started from the same draw and driven by the same noise, the two samplers coincide on the event that $\hat{X}_0,\dots,\hat{X}_{N_{\mathrm{g}}-1}$ all lie in this ball, so the coupling inequality bounds the total variation between their output laws by the probability $\varepsilon_{\mathrm{clip}}$ that one of these iterates leaves it, which the sampler's own paths estimate. So (7) holds for s_ψ with $\varepsilon_{\mathrm{TV}}$ computed for $\tilde{s}_\psi$ and $\varepsilon_{\mathrm{clip}}$ added to it, and the multi-level risk of $\tilde{s}_\psi$ differs from that of s_ψ only through the mass of $\mathsf{H}_{\sigma_n}\mu$ outside the ball.

Remark 4 (Estimating the multi-level risk). If the level index $n < N_{\mathrm{g}}$ is drawn with probability $p_n > 0$ and $\ell_n(\psi)$ is the one-sample bootstrap DSM loss (4) at level σ_n, then $\mathbb{E}\left[(h_n/p_n)\,\ell_n(\psi)\right] = \mathcal{E}_{\mathrm{ml}}(\psi;\mathcal{T}^\pi\nu(x,a)) + C$ with C independent of ψ, by Lemma 4 at every level. On a geometric grid $h_n \propto \sigma_n^2$, so uniform draws call for the weight σ_n^2, the likelihood weighting of score-based diffusion models. Our implementation does not use this estimator: it draws one noise level per sample and weights every sample equally (Section B), so its training loss is not an unbiased estimate of $\mathcal{E}_{\mathrm{ml}}$, and Section D.6 measures the multi-level risk and the bound of Theorem 3 on the trained networks instead.

The consistency residuals of Corollary 3 turn Theorem 3 into a certificate that needs no ground truth at a fixed point of training.

Corollary 5 (Score Bellman residual certificate). *In the setting of Corollary 3, at a fixed point of training with the prior score above σ_0, let $\tilde{e}_n^* \triangleq \sum_{i=0}^{L_n}\tilde{r}_{n-im} + D/(2\sigma_0)$. Then for every (x,a),*

$$\mathcal{E}_{\mathrm{ml}}\big(\psi;\eta^\pi(x,a)\big) \le (1-\kappa_{\mathrm{g}}^2)\sum_{n=0}^{N_{\mathrm{g}}-1}(\tilde{e}_n^*)^2,$$

and Theorem 3 with $\mu = \eta^\pi(x,a)$ and this bound on the multi-level risk in $\varepsilon_{\mathrm{TV}}$ gives $w_1(Q_\psi(x,a),\eta^\pi(x,a)) \le D\varepsilon_{\mathrm{TV}} + \sigma_{N_{\mathrm{g}}}\sqrt{d}$. A field with vanishing consistency residuals satisfies $\mathcal{E}_{\mathrm{ml}}(\psi;\eta^\pi(x,a)) \le \log(\sigma_0/\sigma_{N_{\mathrm{g}}})\,D^2/(2\sigma_0^2)$.

Proof of Corollary 5. For each n, $h_n\mathcal{E}_{\sigma_n}(\psi;\eta^\pi(x,a)) \le (1-\kappa_{\mathrm{g}}^2)\sigma_n^2\bar{\mathcal{E}}_{\sigma_n}(s_\psi;\eta^\pi) = (1-\kappa_{\mathrm{g}}^2)\tilde{e}_n(\psi;\eta^\pi)^2 \le (1-\kappa_{\mathrm{g}}^2)(\tilde{e}_n^*)^2$ by Corollary 3, and summing over n gives the first claim. The sampler bound is Theorem 3 for $\mu = \eta^\pi(x,a)$, which is supported in B_R by Assumption 1. With vanishing residuals every $\tilde{e}_n^*$ equals $D/(2\sigma_0)$, and $(1-\kappa_{\mathrm{g}}^2)N_{\mathrm{g}} \le 2N_{\mathrm{g}}\log(1/\kappa_{\mathrm{g}}) = 2\log(\sigma_0/\sigma_{N_{\mathrm{g}}})$ gives the last claim. □

The residuals $\tilde{r}_n$ are bounded by the population consistency loss (Lemma 5), so every term of the certificate is computed from the training objective and the bootstrap targets.

The next proposition is the witness behind Remark 2. It uses the following functional inequalities.

Definition 5 (Log-Sobolev and Poincaré inequalities). Let $\rho \in \mathcal{P}(\mathbb{R}^d)$ and $C > 0$. The law ρ satisfies the *logarithmic Sobolev inequality* $\mathrm{LSI}(C)$ if $\mathrm{KL}(\nu\,\|\,\rho) \le \frac{C}{2}\mathbb{E}_\nu\left[\|\nabla\log(\mathrm{d}\nu/\mathrm{d}\rho)\|^2\right]$ for every $\nu \ll \rho$ with a strictly positive, locally Lipschitz density $\mathrm{d}\nu/\mathrm{d}\rho$, and the *Poincaré inequality* $\mathrm{P}(C)$ if $\mathrm{var}_\rho[f] \le C\mathbb{E}_\rho\left[\|\nabla f\|^2\right]$ for every locally Lipschitz $f \in L^2(\rho)$. We write C_{LSI} and C_P for the smallest admissible constants; $\mathcal{N}(0,\mathbf{I})$ satisfies both with $C = 1$.

Proposition 3 (Single-level DSM is blind to mode masses). *Let $\sigma > 0$ and $\Delta \ge 2\sigma$, let e_1 be the first unit vector of $\mathbb{R}^d$, and let $\mu \triangleq \frac{1}{2}(\delta_{-\Delta e_1/2} + \delta_{\Delta e_1/2})$.* (i) *The log-Sobolev and Poincaré constants of $\mathsf{H}_\sigma\mu$ satisfy*

$$C_{\mathrm{LSI}}(\sigma) \ge C_P(\sigma) \ge \frac{\sqrt{2\pi}}{4}\,\sigma^2\,e^{(\Delta-2\sigma)^2/(8\sigma^2)}. \tag{17}$$

(ii) *For $0 < \epsilon < \frac{1}{2}$ let $\mu_\epsilon \triangleq (\frac{1}{2}+\epsilon)\delta_{-\Delta e_1/2} + (\frac{1}{2}-\epsilon)\delta_{\Delta e_1/2}$ be the same law with perturbed weights, and let the score model be its smoothed score, $s_\psi(\cdot,\sigma) \triangleq s_{\mu_\epsilon,\sigma}$. Then $w_1(\mu_\epsilon,\mu) = w_1(\mathsf{H}_\sigma\mu_\epsilon,\mathsf{H}_\sigma\mu) = \epsilon\Delta$, whereas*

$$\mathcal{E}_\sigma(\psi;\mu) \le \frac{32\,\Delta\,\epsilon^2}{\sqrt{2\pi}\,\sigma^3(1-2\epsilon)^2}\,e^{-\Delta^2/(8\sigma^2)}.$$

Proof of Proposition 3. Under $\mathsf{H}_\sigma\mu$ the first coordinate has the law $\frac{1}{2}\mathcal{N}(-\Delta/2,\sigma^2) + \frac{1}{2}\mathcal{N}(\Delta/2,\sigma^2)$, with density $p_\sigma^{(1)}$, and the other coordinates are independent $\mathcal{N}(0,\sigma^2)$.

Part (i). Let $f(z) \triangleq \min\{1, \max\{-1, z_1/\sigma\}\}$, which is bounded and Lipschitz. We have $\mathbb{E}_{\mathsf{H}_\sigma\mu}[f] = 0$ by symmetry, $f^2 = 1$ on $\{|z_1| \geq \sigma\}$, and $\|\nabla f\|^2 = \sigma^{-2}\mathbb{1}\{|z_1| < \sigma\}$ almost everywhere. On either component $z_1 = \pm\Delta/2 + \sigma G$ with $G \sim \mathcal{N}(0,1)$, so the event $\{|z_1| \geq \sigma\}$ has probability at least $\frac{1}{2}$, as $\Delta \geq 2\sigma$. For $|u| \leq \sigma \leq \Delta/2$ both Gaussian terms of $p_\sigma^{(1)}(u)$ are at most $\frac{1}{2\sqrt{2\pi}\sigma} e^{-(\Delta/2-\sigma)^2/(2\sigma^2)}$. So

$$\operatorname{var}_{\mathsf{H}_\sigma\mu}[f] \geq (\mathsf{H}_\sigma\mu)(\{|z_1| \geq \sigma\}) \geq \frac{1}{2},$$
$$\mathbb{E}_{\mathsf{H}_\sigma\mu}\left[\|\nabla f\|^2\right] = \sigma^{-2}\int_{-\sigma}^{\sigma} p_\sigma^{(1)}(u)\,\mathrm{d}u \leq \frac{2}{\sqrt{2\pi}\,\sigma^2}\, e^{-(\Delta-2\sigma)^2/(8\sigma^2)},$$

and $\mathrm{P}(C_P(\sigma))$ applied to f gives the second inequality of (17). For the first, $\mathrm{LSI}(C)$ implies $\mathrm{P}(C)$: for bounded locally Lipschitz f with $\mathbb{E}_\rho[f] = 0$, take $\mathrm{d}\nu/\,\mathrm{d}\rho \propto (1+\varepsilon f)^2$; then $\mathrm{KL}(\nu \,\|\, \rho) = 2\varepsilon^2 \operatorname{var}_\rho[f] + o(\varepsilon^2)$ and $\mathbb{E}_\nu\left[\|\nabla \log(\mathrm{d}\nu/\,\mathrm{d}\rho)\|^2\right] = 4\varepsilon^2\mathbb{E}_\rho\left[\|\nabla f\|^2\right] + o(\varepsilon^2)$, and letting $\varepsilon \to 0$ in $\mathrm{LSI}(C)$ and truncating gives $\mathrm{P}(C)$.

Part (ii). Write π_1, π_2 for the posterior responsibilities of the two components under $\mathsf{H}_\sigma\mu$, i.e., $\pi_1(z) = (1 + e^{\Delta z_1/\sigma^2})^{-1} = 1 - \pi_2(z)$. The density of $\mathsf{H}_\sigma\mu_\epsilon$ is that of $\mathsf{H}_\sigma\mu$ times $1 + 2\epsilon(\pi_1 - \pi_2)$, and $\nabla(\pi_1 - \pi_2) = -(2\Delta/\sigma^2)\pi_1\pi_2\, e_1$, so

$$s_{\mu_\epsilon,\sigma} - s_{\mu,\sigma} = \nabla \log\big(1 + 2\epsilon(\pi_1 - \pi_2)\big) = -\frac{4\epsilon\Delta}{\sigma^2}\,\frac{\pi_1\pi_2}{1 + 2\epsilon(\pi_1 - \pi_2)}\, e_1.$$

With $1 + 2\epsilon(\pi_1 - \pi_2) \geq 1 - 2\epsilon$ and $\pi_1\pi_2 \leq e^{-\Delta|z_1|/\sigma^2}$ we get

$$\mathcal{E}_\sigma(\psi;\mu) \leq \frac{16\epsilon^2\Delta^2}{\sigma^4(1-2\epsilon)^2}\int_{\mathbb{R}} e^{-2\Delta|u|/\sigma^2}\, p_\sigma^{(1)}(u)\,\mathrm{d}u.$$

Completing the square in each of the four half-line Gaussian integrals and using $\int_x^\infty e^{-t^2/2}\,\mathrm{d}t \leq e^{-x^2/2}/x$, the four terms carry the factors $\frac{1}{3}$, $\frac{1}{5}$, $\frac{1}{3}$ and $\frac{1}{5}$, and the integral is at most $\frac{16}{15}\frac{\sigma}{\sqrt{2\pi}\Delta}e^{-\Delta^2/(8\sigma^2)} \leq \frac{2\sigma}{\sqrt{2\pi}\Delta}e^{-\Delta^2/(8\sigma^2)}$, which gives the bound on $\mathcal{E}_\sigma(\psi;\mu)$. For the distances, the coupling that moves the mass ϵ from $\Delta e_1/2$ to $-\Delta e_1/2$ gives $w_1(\mu_\epsilon, \mu) \leq \epsilon\Delta$, and adding the same Gaussian noise to both of its coordinates gives $w_1(\mathsf{H}_\sigma\mu_\epsilon, \mathsf{H}_\sigma\mu) \leq \epsilon\Delta$. Conversely, the 1-Lipschitz function $z \mapsto z_1$ has the mean $-\epsilon\Delta$ under μ_ϵ and under $\mathsf{H}_\sigma\mu_\epsilon$, and the mean 0 under μ and under $\mathsf{H}_\sigma\mu$, so both distances are at least $\epsilon\Delta$ by the Kantorovich–Rubinstein duality. □

So on this family $w_1(\mathsf{H}_\sigma\mu_\epsilon, \mathsf{H}_\sigma\mu)/\mathcal{E}_\sigma(\psi;\mu)^{1/2}$ grows at least like $e^{\Delta^2/(16\sigma^2)}$, and a bound of w_1 by a function of the risk at the single level σ that holds for all such models carries a factor of this order. By (i), the Poincaré constant, which is the relaxation time of the Langevin diffusion toward $\mathsf{H}_\sigma\mu$, grows exponentially in $(\Delta/\sigma)^2$ as well.

Corollary 6 (Mode masses of the sampler)**.** *In the setting of Theorem 3, let $\mu = \sum_j w_j\mu_j$ and let $\{A_j\}$ be a Borel partition of $\mathbb{R}^d$ with $\operatorname{dist}(\operatorname{supp}\mu_j, A_j^c) \geq m_{\mathrm{cell}} > 0$ for every j. If $\sigma_{N_g} \leq m_{\mathrm{cell}}/\sqrt{d}$, then*

$$\sum_j \left|\hat{Q}_\psi(A_j) - w_j\right| \leq 2\,\varepsilon_{\mathrm{TV}} + 2\,e^{-(m_{\mathrm{cell}}/\sigma_{N_g} - \sqrt{d})^2/2}.$$

Proof of Corollary 6. For every Borel partition, $\sum_j |P_1(A_j) - P_2(A_j)| \leq 2\,\mathrm{TV}(P_1, P_2)$, which with Theorem 3 gives $\sum_j |\hat{Q}_\psi(A_j) - (\mathsf{H}_{\sigma_{N_g}}\mu)(A_j)| \leq 2\varepsilon_{\mathrm{TV}}$. The leakage satisfies $(\mathsf{H}_{\sigma_{N_g}}\mu)(A_j) - w_j = \sum_{i\neq j} w_i(\mathsf{H}_{\sigma_{N_g}}\mu_i)(A_j) - w_j(\mathsf{H}_{\sigma_{N_g}}\mu_j)(A_j^c)$, so $\sum_j |(\mathsf{H}_{\sigma_{N_g}}\mu)(A_j) - w_j| \leq 2\sum_i w_i(\mathsf{H}_{\sigma_{N_g}}\mu_i)(A_i^c)$. For $Y \sim \mu_i$ and $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$, the event $Y + \sigma_{N_g}\varepsilon \notin A_i$ forces $\sigma_{N_g}\|\varepsilon\| \geq m_{\mathrm{cell}}$, and (11) gives $\Pr(\|\varepsilon\| \geq m_{\mathrm{cell}}/\sigma_{N_g}) \leq e^{-(m_{\mathrm{cell}}/\sigma_{N_g} - \sqrt{d})^2/2}$ for $m_{\mathrm{cell}}/\sigma_{N_g} \geq \sqrt{d}$. The triangle inequality completes the proof. □

A.6 Policy Evaluation

Proof of Theorem 4. Fix k and (x, a). By Assumption 1, $\mathcal{T}^\pi$ maps $\mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$ into itself, so $\mu_k(x,a)$ is supported in B_R, and Theorem 3 gives $w_1(Q_{k+1}(x,a), \mu_k(x,a)) \leq \varepsilon_{\mathrm{gen}}$. By Theorem 1 with $p = 1$ and $\eta^\pi = \mathcal{T}^\pi\eta^\pi$, $w_1(\mu_k(x,a), \eta^\pi(x,a)) \leq \gamma\,\bar{w}_1(\nu_k, \eta^\pi)$.

Part (i). For $\nu_k = Q_k$, which is supported in B_R, the triangle inequality gives $w_1(Q_{k+1}(x,a), \eta^\pi(x,a)) \le \varepsilon_{\text{gen}} + \gamma\, \bar{w}_1(Q_k, \eta^\pi)$. Taking the supremum over (x,a) and unrolling the recursion gives $\bar{w}_1(Q_K, \eta^\pi) \le \gamma^K \bar{w}_1(Q_0, \eta^\pi) + \varepsilon_{\text{gen}} \sum_{i<K} \gamma^i$, which implies (i).

Part (ii). For $\nu_k = \hat{\eta}_k$, one more triangle inequality through $Q_{k+1}(x,a)$ gives $w_1(\hat{\eta}_{k+1}(x,a), \eta^\pi(x,a)) \le \varepsilon_{\text{St}} + \varepsilon_{\text{gen}} + \gamma\, \bar{w}_1(\hat{\eta}_k, \eta^\pi)$, which unrolls in the same way. The bound on Q_K is the first paragraph of this proof at $k = K-1$. □

Remark 5 (The smoothing term of Theorem 4). The term $\sigma_{N_g}\sqrt{d}/(1-\gamma)$ inside $\varepsilon_{\text{gen}}/(1-\gamma)$ is not an artifact of the proof. Sampler draws carry the smoothing variance $\sigma^2_{N_g}$ per coordinate, and on a self-loop with a deterministic transition and an otherwise exact sampler the bootstrap compounds it to an excess of $\gamma^2\sigma^2_{N_g}/(1-\gamma^2)$ per coordinate in the variance of the Bellman target at the fixed point, since the target variance v then solves $v = \gamma^2(v + \sigma^2_{N_g})$. Section D.6 measures this accumulation, and the affine moment matching of Section B, followed by the projection onto B_R, removes it. The moment-matched draws are a particle set supported in B_R, so Theorem 4(ii) covers sampler bootstrapping with moment matching, with these draws in the role of the Stein particles and their measured distance to the sampler law as ε_{St}.

Corollary 7 (Supremal Wasserstein-2). *In the setting of Theorem 4, with* $\varepsilon_{\text{gen},2} \triangleq D \sup_{k,(x,a)} \varepsilon_{\text{TV},k}(x,a)^{1/2} + \sigma_{N_g}\sqrt{d}$, *sampler bootstrapping satisfies*

$$\bar{w}_2(Q_K, \eta^\pi) \le \gamma^K\, \bar{w}_2(Q_0, \eta^\pi) + \frac{\varepsilon_{\text{gen},2}}{1-\gamma},$$

and bootstrapping from Stein particles the same bound with $\sup_{k,(x,a)} w_2(\hat{\eta}_{k+1}(x,a), Q_{k+1}(x,a))$ *added to the numerator.*

Proof of Corollary 7. Lemma 3(iv) for $(\text{proj}_{B_R})_\# \hat{Q}_\psi$ and $(\text{proj}_{B_R})_\# \mathsf{H}_{\sigma_{N_g}}\mu$, data processing for the projection, Proposition 2(i) and Theorem 3 give $w_2(Q_\psi, \mu) \le D\varepsilon_{\text{TV}}^{1/2} + \sigma_{N_g}\sqrt{d}$. The rest is the proof of Theorem 4 with w_1 replaced by w_2, using Theorem 1 with $p = 2$. □

The bound needs no dissipativity or mode-separation assumption, and the model error enters through its square root, which is the price of an argument without isoperimetry.

Corollary 8 (Mode masses of the evaluated policy). *In the setting of Theorem 4, fix* (x,a), *write* $\eta^\pi(x,a) = \sum_j w_j \mu_j$, *and let* $\{A_j\}$ *be a Borel partition of* $\mathbb{R}^d$ *with* $\text{dist}(\text{supp}\,\mu_j, A_j^c) \ge m_{\text{cell}} > 0$ *for every* j. *Let* Q *be any law on* B_R *and* $\sigma_c \le m_{\text{cell}}/\sqrt{d}$. *Then the* σ_c*-smoothed cell masses of* Q *satisfy*

$$\sum_j \left|(\mathsf{H}_{\sigma_c} Q)(A_j) - w_j\right| \le \frac{2}{\sigma_c\sqrt{2\pi}}\, w_1\big(Q, \eta^\pi(x,a)\big) + 2\, e^{-(m_{\text{cell}}/\sigma_c - \sqrt{d})^2/2},$$

so with $Q = Q_K(x,a)$ *the* $\bar{w}_1$ *bound of Theorem 4 certifies the mode masses of the evaluated policy at every* (x,a).

Proof of Corollary 8. For $x, y \in \mathbb{R}^d$, $\text{TV}\big(\mathcal{N}(x, \sigma_c^2\mathbf{I}), \mathcal{N}(y, \sigma_c^2\mathbf{I})\big) = 2\Phi\big(\|x-y\|/(2\sigma_c)\big) - 1 \le \|x-y\|/(\sigma_c\sqrt{2\pi})$, since the standard normal density is at most $(2\pi)^{-1/2}$. Joint convexity of TV over the w_1-optimal coupling of Q and $\eta^\pi(x,a)$ gives $\text{TV}(\mathsf{H}_{\sigma_c} Q, \mathsf{H}_{\sigma_c}\eta^\pi(x,a)) \le w_1(Q, \eta^\pi(x,a))/(\sigma_c\sqrt{2\pi})$. Then $\sum_j |(\mathsf{H}_{\sigma_c} Q)(A_j) - (\mathsf{H}_{\sigma_c}\eta^\pi(x,a))(A_j)| \le 2\,\text{TV}(\mathsf{H}_{\sigma_c} Q, \mathsf{H}_{\sigma_c}\eta^\pi(x,a))$, and the leakage bound of the proof of Corollary 6 at the scale σ_c completes the proof. □

Corollary 9 (Matrix discounting). *Let the Bellman target be* $b_{r,\Gamma}(z) = r + \Gamma z$ *with a matrix* $\Gamma \in \mathbb{R}^{d\times d}$ *of operator norm* $\|\Gamma\|_{\text{op}} < 1$, *and let Assumption 1 hold with* B_R *replaced by a box that every* $b_{r,\Gamma}$ *maps into itself. Then Theorems 4 and 6 hold verbatim with* γ *replaced by* $\|\Gamma\|_{\text{op}}$.

Proof of Corollary 9. The affine map $z \mapsto r + \Gamma z$ contracts Euclidean w_1 by the factor $\|\Gamma\|_{\text{op}}$, and mixtures over transitions preserve the contraction, so Theorem 1 holds with this factor. The proofs of Theorems 4 and 6 use γ only through Theorem 1, and Theorem 3 does not involve the discount. □

Under matrix discounting, sliced divergences contract through the max-slicing construction of Debes and Tuytelaars (2026b), whereas this bound needs no such construction, and Corollary 2 gives the matching score-space identity.

A.7 No Global Dissipativity

The Stein particles enter Theorem 4 only through $\varepsilon_{\mathrm{St}}$, which we measure, because the contraction that an unconditional bound would need fails for SVGD. Fix a level $\sigma > 0$ and μ supported in B_R, and let $k(z, z') = \kappa(z - z')$ be a translation-invariant kernel with an even $\kappa \in C^1(\mathbb{R}^d)$ that is bounded with a bounded gradient. We write

$$\Phi_\nu(z) \triangleq \mathbb{E}_{z'\sim\nu}\left[k(z', z)\, s_{\mu,\sigma}(z') + \nabla_{z'} k(z', z)\right], \qquad \mathbb{F}(z_{1:M})_i \triangleq \frac{1}{M}\sum_{j=1}^{M}\left[k(z_j, z_i)\, s_{\mu,\sigma}(z_j) + \nabla_{z_j} k(z_j, z_i)\right]$$

for the mean-field and the joint M-particle exact-score SVGD drifts. The drift Φ is *globally dissipative* with constant $\lambda > 0$ if $\mathbb{E}\left[\langle X - Y, \Phi_\nu(X) - \Phi_{\nu'}(Y)\rangle\right] \le -\lambda\, \mathbb{E}\left[\|X - Y\|^2\right]$ for all $\nu, \nu' \in \mathcal{P}_2(\mathbb{R}^d)$ and all couplings (X, Y) of them, and $\mathbb{F}$ is globally dissipative if $\langle z - y, \mathbb{F}(z) - \mathbb{F}(y)\rangle \le -\lambda |z - y|^2$ for all $z, y \in (\mathbb{R}^d)^M$, with $|\cdot|$ the Euclidean norm of $\mathbb{R}^{dM}$. Either property would make the exact-score SVGD flow a contraction toward its target.

Proposition 4 (No global dissipativity for the exact-score SVGD drift)**.** *Let μ be supported in B_R and $\sigma > 0$.* (i) *For every even kernel $\kappa \in C^1(\mathbb{R}^d)$ that is bounded with a bounded gradient, the mean-field drift Φ is not globally dissipative with any constant $\lambda > 0$.* (ii) *For the RBF kernel $\kappa(u) = e^{-\|u\|^2/(2h^2)}$ and every $M \ge 2$ with*

$$M \ge 1 + \frac{h^2}{\sigma^2},$$

the joint drift $\mathbb{F}$ is not globally dissipative with any constant $\lambda > 0$. With a bandwidth $h \le C_h \sigma$ it therefore fails for every $M \ge 1 + C_h^2$. (iii) *For every $M \ge 2$ and every kernel as in (i), global dissipativity of $\mathbb{F}$ implies that every zero of $\mathbb{F}$ is a collapsed configuration $(c, \dots, c)$ with $s_{\mu,\sigma}(c) = 0$, and that the M-particle exact-score SVGD flow converges exponentially to it from every initial cloud.*

Proof of Proposition 4. Part (i). The density p_σ of $\mathsf{H}_\sigma\mu$ is smooth, positive and has Gaussian tails, so it has a maximizer $\bar{z}$, and $s_{\mu,\sigma}(\bar{z}) = 0$. For every z, the function $g(z') \triangleq \kappa(z' - z)\, p_\sigma(z')$ is integrable and continuously differentiable with the integrable gradient $\nabla g = \kappa(\cdot - z)\, p_\sigma\, s_{\mu,\sigma} + \nabla\kappa(\cdot - z)\, p_\sigma$, as κ and $\nabla\kappa$ are bounded and $\mathbb{E}_{\mathsf{H}_\sigma\mu}\left[\|s_{\mu,\sigma}\|\right] \le \mathcal{I}(\mathsf{H}_\sigma\mu)^{1/2} < \infty$ (Lemma 6). So $\Phi_{\mathsf{H}_\sigma\mu}(z) = \int \nabla g(z')\, \mathrm{d}z' = 0$, which is Stein's identity; the integral of the gradient of an integrable C^1 function with an integrable gradient vanishes, as one sees by integrating against $\chi(\cdot/R)$ for a smooth cutoff χ and letting $R \to \infty$. At the point mass, $\Phi_{\delta_{\bar{z}}}(\bar{z}) = \kappa(0)\, s_{\mu,\sigma}(\bar{z}) + \nabla\kappa(0) = 0$, since $\nabla\kappa(0) = 0$ for an even C^1 function. For $\nu = \mathsf{H}_\sigma\mu$, $\nu' = \delta_{\bar{z}}$ and their only coupling $(X, \bar{z})$, the left side of the dissipativity inequality is $\mathbb{E}\left[\langle X - \bar{z}, 0\rangle\right] = 0$, whereas the right side is $-\lambda\, \mathbb{E}\left[\|X - \bar{z}\|^2\right] < 0$, as $\mathsf{H}_\sigma\mu$ is not a point mass.

Part (ii). With $\bar{z}$ as in part (i), the collapsed configuration $z^* \triangleq (\bar{z}, \dots, \bar{z})$ is a zero of $\mathbb{F}$, since $\mathbb{F}(z^*)_i = \kappa(0)\, s_{\mu,\sigma}(\bar{z}) + \nabla\kappa(0) = 0$. As $\mathbb{F} \in C^1$, dissipativity at the pairs $(z^* + tv, z^*)$ and $t \downarrow 0$ give $\langle v, D\mathbb{F}(z^*)v\rangle \le -\lambda|v|^2$ for every $v \in (\mathbb{R}^d)^M$. Take $v = (\delta, 0, \dots, 0)$ with $\delta \in \mathbb{R}^d$. In the derivative of $\mathbb{F}(z)_1 = \frac{1}{M}\sum_j \left[\kappa(z_j - z_1)\, s_{\mu,\sigma}(z_j) + \nabla\kappa(z_j - z_1)\right]$ with respect to z_1 at z^*, the term $j = 1$ contributes $\frac{1}{M}\kappa(0)\nabla s_{\mu,\sigma}(\bar{z})$, and each term $j \ne 1$ contributes $\frac{1}{M}\left[-s_{\mu,\sigma}(\bar{z})\nabla\kappa(0)^\top - \nabla^2\kappa(0)\right] = -\frac{1}{M}\nabla^2\kappa(0)$. For the RBF kernel $\kappa(0) = 1$ and $\nabla^2\kappa(0) = -h^{-2}\mathbf{I}$, so

$$\begin{aligned}\delta^\top\Big(\frac{\partial \mathbb{F}_1}{\partial z_1}(z^*)\Big)\delta &= \frac{1}{M}\Big[\delta^\top \nabla^2 \log p_\sigma(\bar{z})\, \delta + \frac{M-1}{h^2}\|\delta\|^2\Big] \\ &\ge \frac{1}{M}\Big[\frac{M-1}{h^2} - \frac{1}{\sigma^2}\Big]\|\delta\|^2. \qquad \text{((9) gives } -\nabla^2 \log p_\sigma \preceq \sigma^{-2}\mathbf{I})\end{aligned}$$

The left side is at most $-\lambda\|\delta\|^2 < 0$, which forces $(M-1)/h^2 < 1/\sigma^2$, i.e., $M < 1 + h^2/\sigma^2$, proving (ii).

Part (iii). If $z^* \ne y^*$ were two zeros of $\mathbb{F}$, dissipativity at the pair would give $0 \le -\lambda|z^* - y^*|^2 < 0$, so $\mathbb{F}$ has at most one zero. The zero exists, as $(\bar{z}, \dots, \bar{z})$ is one. The drift is permutation-equivariant, $\mathbb{F}(Pz) = P\mathbb{F}(z)$ for every permutation P of the particles, so the unique zero satisfies $Pz^* = z^*$ for all P, i.e., $z^* = (c, \dots, c)$, and $0 = \mathbb{F}(z^*)_i = \kappa(0)\, s_{\mu,\sigma}(c)$. If $\kappa(0) = 0$, every collapsed configuration would be a zero, so $\kappa(0) \ne 0$ and $s_{\mu,\sigma}(c) = 0$. Along the flow $\dot{z} = \mathbb{F}(z)$ we have $\frac{\mathrm{d}}{\mathrm{d}t}|z(t) - z^*|^2 = 2\langle z(t) - z^*, \mathbb{F}(z(t)) - \mathbb{F}(z^*)\rangle \le -2\lambda|z(t) - z^*|^2$, so $|z(t) - z^*| \le e^{-\lambda t}|z(0) - z^*|$. □

Near a collapsed cloud the $M - 1$ repulsive terms $-\nabla^2\kappa(0) = h^{-2}\mathbf{I}$ outweigh the attraction $\nabla^2 \log p_\sigma \succeq -\sigma^{-2}\mathbf{I}$, so the collapsed zero is unstable, whereas a dissipative drift has a single zero that attracts every cloud. For instance,

for $\mu = \delta_0$, $\sigma = h = 1$ and $M = 2$, both $(0,0)$ and $(a,-a)$ with $e^{-2a^2} = 1/3$ are zeros of $\mathbb{F}$, and SVGD converges to the second. Near equilibrium, exponential convergence of SVGD fails for every kernel in $C^{1,1}$, the RBF kernel included, under a mild integrability condition on the target (Duncan et al., 2023, Lemma 39 and Corollary 42).

A.8 Risk Queries and the Resolution Floor

Theorem 6 (CVaR from the sampler)**.** *Fix (x,a), a bootstrap source $\nu \in \mathcal{P}(B_R)^{\mathcal{X}\times\mathcal{A}}$, the Bellman target $\mu \triangleq \mathcal{T}^\pi\nu(x,a)$, a score model ψ with sampler law Q_ψ and bound $\varepsilon_{\mathrm{TV}}$ from Theorem 3, and $\lambda \in \mathbb{R}^d$ with $\|\lambda\|_2 \le 1$. Let $\hat{F}_N$ be the empirical distribution function of N i.i.d. draws from $\lambda_\# Q_\psi$. Then, for $\alpha, \beta \in (0,1)$, with probability at least $1-\beta$,*

$$\left|\mathrm{CVaR}_\alpha(\hat{F}_N) - \mathrm{CVaR}_\alpha\big(\lambda_\#\eta^\pi(x,a)\big)\right| \le \frac{D}{\alpha}\Big(\varepsilon_{\mathrm{TV}} + \sqrt{\frac{\log(2/\beta)}{2N}}\Big) + \frac{\sigma_{N_g}\sqrt{d} + \gamma\,\bar{w}_1(\nu,\eta^\pi)}{\alpha},$$

where CVaR_α is the lower-tail CVaR, $\mathrm{CVaR}_\alpha(P) = \frac{1}{\alpha}\int_0^\alpha F_P^{-1}(u)\,\mathrm{d}u$, and $\mathrm{CVaR}_\alpha(\hat{F}_N)$ is the plug-in estimate.

Proof of Theorem 6. Let I be the smallest closed interval containing $\lambda^\top B_R$; its length is at most D, since $|\lambda^\top z - \lambda^\top z'| \le \|\lambda\|_2\|z - z'\| \le D$ for $z, z' \in B_R$. The triangle inequality splits the error into four terms,

$$\begin{aligned}\left|\mathrm{CVaR}_\alpha(\hat{F}_N) - \mathrm{CVaR}_\alpha(\lambda_\#\eta^\pi(x,a))\right| \le{}& \underbrace{\left|\mathrm{CVaR}_\alpha(\hat{F}_N) - \mathrm{CVaR}_\alpha(\lambda_\# Q_\psi)\right|}_{T_1}\\ &+ \underbrace{\left|\mathrm{CVaR}_\alpha(\lambda_\# Q_\psi) - \mathrm{CVaR}_\alpha\big(\lambda_\#(\mathrm{proj}_{B_R})_\#\mathsf{H}_{\sigma_{N_g}}\mu\big)\right|}_{T_2}\\ &+ \underbrace{\left|\mathrm{CVaR}_\alpha\big(\lambda_\#(\mathrm{proj}_{B_R})_\#\mathsf{H}_{\sigma_{N_g}}\mu\big) - \mathrm{CVaR}_\alpha(\lambda_\#\mu)\right|}_{T_3}\\ &+ \underbrace{\left|\mathrm{CVaR}_\alpha(\lambda_\#\mu) - \mathrm{CVaR}_\alpha(\lambda_\#\eta^\pi(x,a))\right|}_{T_4}.\end{aligned}$$

Both $\lambda_\# Q_\psi$ and the empirical law of the draws are supported in I, so Lemma 3(ii) and the Dvoretzky–Kiefer–Wolfowitz inequality with Massart's constant (Massart, 1990) give $T_1 \le \frac{D}{\alpha}\sqrt{\log(2/\beta)/(2N)}$ with probability at least $1-\beta$. Both laws in T_2 are supported in I, so Lemma 3(i), (iv) and data processing for $\lambda \circ \mathrm{proj}_{B_R}$ give $T_2 \le \frac{D}{\alpha}\mathrm{TV}(\hat{Q}_\psi, \mathsf{H}_{\sigma_{N_g}}\mu) \le \frac{D}{\alpha}\varepsilon_{\mathrm{TV}}$. For T_3, Lemma 3(i), (iii) and the last step of the proof of Theorem 3 give $T_3 \le \frac{1}{\alpha}w_1\big((\mathrm{proj}_{B_R})_\#\mathsf{H}_{\sigma_{N_g}}\mu, \mu\big) \le \frac{1}{\alpha}\sigma_{N_g}\sqrt{d}$. For T_4, Lemma 3(i), (iii) and Theorem 1 with $p=1$ give $T_4 \le \frac{1}{\alpha}w_1\big(\mathcal{T}^\pi\nu(x,a), \mathcal{T}^\pi\eta^\pi(x,a)\big) \le \frac{\gamma}{\alpha}\bar{w}_1(\nu,\eta^\pi)$. □

With $\nu = Q_K$ or $\nu = \hat{\eta}_K$, Theorem 4 bounds the last term. No term depends on M, and the estimation term decays with the number of draws N alone; the factor $1/\alpha$ is the exact Lipschitz constant of CVaR_α in w_1.

Corollary 10 (One set of sampler draws bounds the error for every scalarization and every risk level)**.** *In the setting of Theorem 6, let $\hat{F}_N^\lambda$ be the empirical distribution function of $\lambda^\top\hat{X}_1, \dots, \lambda^\top\hat{X}_N$ for the same N sampler draws $\hat{X}_i \sim Q_\psi$, and let $\alpha_0 \in (0,1)$. Then, with probability at least $1-\beta$, simultaneously for all λ in the Euclidean unit ball and all $\alpha \in [\alpha_0, 1)$,*

$$\left|\mathrm{CVaR}_\alpha(\hat{F}_N^\lambda) - \mathrm{CVaR}_\alpha\big(\lambda_\#\eta^\pi(x,a)\big)\right| \le \frac{D}{\alpha_0}\Big(\varepsilon_{\mathrm{TV}} + \sqrt{\frac{d\log(3\sqrt{N}) + \log(2/\beta)}{2N}} + \frac{1}{\sqrt{N}}\Big) + \frac{\sigma_{N_g}\sqrt{d} + \gamma\,\bar{w}_1(\nu,\eta^\pi)}{\alpha_0}.$$

Proof of Corollary 10. For P supported in $B(0,R_Z)$ and λ, λ' in the unit ball, the identity coupling gives $w_1(\lambda_\# P, \lambda'_\# P) \le \mathbb{E}_P\left[|(\lambda - \lambda')^\top Z|\right] \le \|\lambda - \lambda'\|R_Z$. So by Lemma 3(i) the map $\lambda \mapsto \mathrm{CVaR}_\alpha(\lambda_\# P)$ is (R_Z/α)-Lipschitz, for $P = \eta^\pi(x,a)$ and for the empirical law of the draws alike. Let Λ be an ϵ-net of the unit ball with $\epsilon = N^{-1/2}$ and $|\Lambda| \le (3/\epsilon)^d$. Theorem 6 with confidence $\beta/|\Lambda|$ at each $\lambda \in \Lambda$ and a union bound give its bound at every net point with $\log(2/\beta)$ replaced by $d\log(3\sqrt{N}) + \log(2/\beta)$, simultaneously for all α, since the Dvoretzky–Kiefer–Wolfowitz inequality is uniform in α and $1/\alpha \le 1/\alpha_0$. For an arbitrary λ and its nearest net point λ', the two Lipschitz bounds add $2\epsilon R_Z/\alpha_0 = D/(\alpha_0\sqrt{N})$. □

Corollary 11 (Conditional queries by rejection). *In the setting of Theorem 3, let $A \subseteq \mathbb{R}^d$ be a Borel set with $(\mathsf{H}_{\sigma_{N_g}}\mu)(A) \geq p_A > \varepsilon_{\mathrm{TV}}$ and let $f\colon \mathbb{R}^d \to [-F, F]$ be measurable. Then*

$$\left|\mathbb{E}_{\hat{Q}_\psi}[f \mid A] - \mathbb{E}_{\mathsf{H}_{\sigma_{N_g}}\mu}[f \mid A]\right| \leq \frac{3F\,\varepsilon_{\mathrm{TV}}}{p_A - \varepsilon_{\mathrm{TV}}}.$$

Let $u_N \triangleq \sqrt{\log(4/\beta)/(2N)}$ and $p_A > \varepsilon_{\mathrm{TV}} + u_N$. The rejection estimator from N sampler draws differs from $\mathbb{E}_{\hat{Q}_\psi}[f \mid A]$ by at most $3F\,u_N/(p_A - \varepsilon_{\mathrm{TV}} - u_N)$ with probability at least $1 - \beta$.

Proof of Corollary 11. For a law ρ and a bounded measurable g we write $\rho(g) \triangleq \int g\,\mathrm{d}\rho$, and we let $\epsilon \triangleq \mathrm{TV}(\hat{Q}_\psi, \mathsf{H}_{\sigma_{N_g}}\mu)$, which is at most $\varepsilon_{\mathrm{TV}}$ by Theorem 3. For every measurable g with $|g| \leq G$ we have $|\hat{Q}_\psi(g) - (\mathsf{H}_{\sigma_{N_g}}\mu)(g)| \leq 2G\epsilon$, and $|\hat{Q}_\psi(A) - (\mathsf{H}_{\sigma_{N_g}}\mu)(A)| \leq \epsilon$. With $|(\mathsf{H}_{\sigma_{N_g}}\mu)(f\mathbb{1}_A)| \leq F\,(\mathsf{H}_{\sigma_{N_g}}\mu)(A)$ and $\hat{Q}_\psi(A) \geq p_A - \epsilon$ we get

$$\begin{aligned}\left|\frac{\hat{Q}_\psi(f\mathbb{1}_A)}{\hat{Q}_\psi(A)} - \frac{(\mathsf{H}_{\sigma_{N_g}}\mu)(f\mathbb{1}_A)}{(\mathsf{H}_{\sigma_{N_g}}\mu)(A)}\right| &\leq \frac{|\hat{Q}_\psi(f\mathbb{1}_A) - (\mathsf{H}_{\sigma_{N_g}}\mu)(f\mathbb{1}_A)|}{\hat{Q}_\psi(A)} + \frac{|(\mathsf{H}_{\sigma_{N_g}}\mu)(f\mathbb{1}_A)|}{(\mathsf{H}_{\sigma_{N_g}}\mu)(A)}\,\frac{|(\mathsf{H}_{\sigma_{N_g}}\mu)(A) - \hat{Q}_\psi(A)|}{\hat{Q}_\psi(A)} \\ &\leq \frac{2F\epsilon}{p_A - \epsilon} + \frac{F\epsilon}{p_A - \epsilon},\end{aligned}$$

and the right side is increasing in ϵ, which proves the first claim. For the estimator, Hoeffding's inequality gives $|\hat{Q}_N(f\mathbb{1}_A) - \hat{Q}_\psi(f\mathbb{1}_A)| \leq 2F\,u_N$ and $|\hat{Q}_N(A) - \hat{Q}_\psi(A)| \leq u_N$ for the empirical law $\hat{Q}_N$ of the draws, each with probability at least $1 - \beta/2$. On both events $\hat{Q}_N(A) \geq \hat{Q}_\psi(A) - u_N \geq p_A - \varepsilon_{\mathrm{TV}} - u_N > 0$, and the same quotient bound with $|\hat{Q}_\psi(f\mathbb{1}_A)| \leq F\,\hat{Q}_\psi(A)$ gives the second claim. □

The lower bounds on particles use the optimal quantizers of a two-point law.

Proposition 5 (Optimal quantizers of a two-point law). *Let $M \geq 1$, $\alpha \in (0, 1/M)$, $z_0, \lambda \in \mathbb{R}^d$ with $\|\lambda\|_2 = 1$, and $T > 0$. For $t \in (0, T]$ let $\eta_t \triangleq (1 - \frac{\alpha}{2})\delta_{z_0} + \frac{\alpha}{2}\delta_{z_0 - t\lambda}$, and let $\mathcal{F}_M$ be the equal-weight M-atom laws of Section 2, atoms allowed to coincide.* (i) *For every $t \in (0, T]$ and every $Q = \frac{1}{M}\sum_i \delta_{y_i} \in \mathcal{F}_M$,*

$$w_1(\eta_t, Q) \geq \frac{\alpha t}{2} + \frac{1 - \alpha M}{M}\sum_{i=1}^{M}\|y_i - z_0\|, \tag{18}$$

so δ_{z_0} is the unique w_1-optimal quantizer of η_t in $\mathcal{F}_M$, for every t. (ii) *The w_2-optimal quantizer of η_t in $\mathcal{F}_M$ places $M - 1$ atoms at z_0 and one atom at $z_0 - \frac{\alpha M}{2}t\lambda$.* (iii) *Under Assumption 1, the choice $z_0 = z_{\max}$, $\lambda = z_{\max}/\|z_{\max}\|_2$ and $T = D$ places every η_t inside B_R.*

Proof of Proposition 5. Part (i). We take $z_0 = 0$, which is no loss of generality as w_1 and $\mathcal{F}_M$ are translation invariant. Fix t and $Q = \frac{1}{M}\sum_i \delta_{y_i}$. Splitting the mass at a point that several atoms share equally among their labels, every coupling of (η_t, Q) moves masses $a_i, b_i \geq 0$ from 0 and from $-t\lambda$ to the atom y_i, with $a_i + b_i = 1/M$, $\sum_i a_i = 1 - \alpha/2$ and $\sum_i b_i = \alpha/2$, at the cost $\sum_i \big(a_i\|y_i\| + b_i\|y_i + t\lambda\|\big)$, and every such (a, b) is a coupling; so $w_1(\eta_t, Q)$ is the infimum of this cost over admissible (a, b). The triangle inequality gives $\|y_i + t\lambda\| \geq t - \|y_i\|$, and $a_i - b_i = \frac{1}{M} - 2b_i \geq \frac{1}{M} - \alpha > 0$, so

$$a_i\|y_i\| + b_i\|y_i + t\lambda\| \geq b_i t + (a_i - b_i)\|y_i\| \geq b_i t + \Big(\frac{1}{M} - \alpha\Big)\|y_i\|,$$

and summing over i gives (18). As $w_1(\eta_t, \delta_0) = \frac{\alpha t}{2}$ and $1 - \alpha M > 0$, every $Q \neq \delta_0$ is strictly worse than δ_0.

Part (ii). Again $z_0 = 0$, and the description of the couplings by (a, b) in (i) gives $w_2^2(\eta_t, Q) = \inf_{(a,b)}\sum_i \big(a_i\|y_i\|^2 + b_i\|y_i + t\lambda\|^2\big)$. For fixed (a, b) and i, the quadratic $y \mapsto a_i\|y\|^2 + b_i\|y + t\lambda\|^2$ has the unique minimizer $y_i = -Mb_i t\lambda$ and the minimum $Ma_i b_i t^2$, as $a_i + b_i = 1/M$. Minimizing over Q and (a, b) jointly gives

$$\min_{Q \in \mathcal{F}_M} w_2^2(\eta_t, Q) = \min_{(a,b)} Mt^2 \sum_i b_i\Big(\frac{1}{M} - b_i\Big) = \min_{(a,b)} t^2\Big(\frac{\alpha}{2} - M\sum_i b_i^2\Big).$$

As $\sum_i b_i^2 \le \max_i b_i \sum_i b_i \le (\alpha/2)^2$, with equality if and only if one b_i equals $\alpha/2$ and the others vanish, which is admissible since $\alpha/2 < 1/M$, the minimum is attained exactly there, with the atoms $-\frac{\alpha M}{2} t\lambda$ for that i and 0 for the other $M-1$.

Part (iii). For $z_0 = z_{\max}$, $\lambda = z_{\max}/\|z_{\max}\|_2$ and $T = D = 2\|z_{\max}\|_2$ we have $z_0 - t\lambda = (1 - t/\|z_{\max}\|_2) z_{\max}$ with $|1 - t/\|z_{\max}\|_2| \le 1$ for $t \in [0, D]$, so both support points of η_t lie in B_R. □

Proposition 6 (What particles cannot answer, full statement of Proposition 1)**.** *For a unit vector $\lambda \in \mathbb{R}^d$ and $\tau \in \mathbb{R}$ let $A_{\lambda,\tau} \triangleq \{z : \lambda^\top z \le \tau\}$, write $\bar{G}(u) \triangleq \Pr(\xi > u)$ for $\xi \sim \mathcal{N}(0,1)$, and read every query from an equal-weight law $Q = \frac{1}{M}\sum_{i=1}^M \delta_{y_i}$ by plug-in, i.e., as the same query on Q.* (i) Every particle set, every event. *For every Borel A and every law η with $p \triangleq \eta(A) \in (0, \frac{1}{2M})$, either $Q(A) = 0$ or $Q(A) \ge \frac{1}{M} > 2p$; in both cases $|Q(A) - p| \ge p$.* (ii) Every particle set, one law. *Let $z_0 \triangleq z_{\max}$, $\lambda \triangleq z_{\max}/\|z_{\max}\|_2$, $t \in (0, D]$, $p \in (0, \frac{1}{2M})$, $\eta_t \triangleq (1-p)\delta_{z_0} + p\,\delta_{z_0 - t\lambda} \in \mathcal{P}(B_R)$, and $A \triangleq A_{\lambda, \lambda^\top z_0 - t/2}$, so $\eta_t(A) = p$. For every Q: (a) the decision "take the risky action if and only if $Q(A) \le \delta$" is wrong for every $\delta \in [0, p)$ if $Q(A) = 0$, and for every $\delta \in [p, \frac{1}{M})$ otherwise, a set of thresholds of length at least p; (b) $\max_{\alpha \in \{p, 1/M\}} \left|\mathrm{CVaR}_\alpha(\lambda_\# Q) - \mathrm{CVaR}_\alpha(\lambda_\# \eta_t)\right| \ge \frac{t}{2}(1 - 2pM)$, and if $Q(A) = 0$ the error exceeds $\frac{t}{2}$ at every $\alpha \in (0, p]$; (c) the w_1-optimal and the w_2-optimal quantizers of η_t in $\mathcal{F}_M$ and the quantile representation of $\lambda_\# \eta_t$ at the levels $(2i-1)/(2M)$ all satisfy $Q(A) = 0$; so for them the decision is wrong for every $\delta \in [0, p)$, the conditional law given A is undefined, and, for the w_1-optimal quantizer, the entropic risk $-\frac{1}{\beta}\log \mathbb{E}\left[e^{-\beta \lambda^\top Z}\right]$, $\beta > 0$, errs by $\frac{1}{\beta}\log(1 - p + pe^{\beta t}) \ge t - \frac{1}{\beta}\log\frac{1}{p}$.* (iii) Sampler, one scalarization. *Let μ be supported in B_R, let ψ be a score model with sampler law $\hat{Q}_\psi$ and bound $\varepsilon_{\mathrm{TV}}$ from Theorem 3, and let $\hat{Q}_N$ be the empirical measure of N i.i.d. draws from $\hat{Q}_\psi$. For a fixed unit λ and $\beta \in (0,1)$, with probability at least $1 - \beta$, simultaneously for all $\tau \in \mathbb{R}$ and $w > 0$,*

$$\left|\hat{Q}_N(A_{\lambda,\tau}) - \mu(A_{\lambda,\tau})\right| \le \varepsilon_{\mathrm{TV}} + \mu\big(\{y : |\lambda^\top y - \tau| < w\}\big) + \bar{G}\Big(\frac{w}{\sigma_{N_g}}\Big) + \sqrt{\frac{\log(2/\beta)}{2N}}. \tag{19}$$

(iv) Sampler, all scalarizations. *With probability at least $1 - \beta$, (19) holds simultaneously for all unit λ, all τ and all $w > 0$ with its last term replaced by $\sqrt{32\big((d+1)\log(N+1) + \log(8/\beta)\big)/N}$.* (v) Propagation. *Suppose that at (x_j, a_j) the reward is 0 and the successor pair (x_{j+1}, a_{j+1}) is reached deterministically under π. Then for every bootstrap source ν and Borel B, $(\mathcal{T}^\pi \nu)(x_j, a_j)(\gamma B) = \nu(x_{j+1}, a_{j+1})(B)$ and $\mathrm{TV}\big((\mathcal{T}^\pi \nu)(x_j, a_j), (\mathcal{T}^\pi \nu')(x_j, a_j)\big) = \mathrm{TV}\big(\nu(x_{j+1}, a_{j+1}), \nu'(x_{j+1}, a_{j+1})\big)$. So along a chain of such pairs, every Bellman target built from M downstream particles gives the pulled-back event a mass in $\{0, \frac{1}{M}, \dots, 1\}$, the particles' own mass of the downstream event, and every target built from a downstream sampler inherits the event error of that sampler in (iii) and (iv) unchanged.*

Proof of Proposition 6. Part (i). $Q(A)$ is a multiple of $\frac{1}{M}$. If $Q(A) = 0$, the error is p; otherwise $Q(A) \ge \frac{1}{M} > 2p$ and $Q(A) - p > p$.

Part (ii)(a). The true decision takes the risky action if and only if $p \le \delta$. If $Q(A) = 0$, the decision computed from Q takes the risky action for every $\delta \ge 0$, which is wrong for $\delta \in [0, p)$. Otherwise $Q(A) \ge \frac{1}{M}$, and the decision computed from Q takes the safe action for every $\delta < \frac{1}{M}$, which is wrong for $\delta \in [p, \frac{1}{M})$, an interval of length $\frac{1}{M} - p > p$.

Part (ii)(b). Let $v_{(1)}$ be the smallest of the scalarized atoms $\lambda^\top y_i$. The lower-tail CVaR of an equal-weight law at a level $\alpha \le \frac{1}{M}$ is its smallest atom, since the lowest α of its mass lies in one atom, and it is at least the smallest atom at every level. The scalarized target $\lambda_\# \eta_t = (1-p)\delta_{\lambda^\top z_0} + p\,\delta_{\lambda^\top z_0 - t}$ has $\mathrm{CVaR}_\alpha(\lambda_\# \eta_t) = \lambda^\top z_0 - t$ for $\alpha \le p$ and $\mathrm{CVaR}_{1/M}(\lambda_\# \eta_t) = \lambda^\top z_0 - pMt$, by the quantile integral of Theorem 6. If $Q(A) = 0$, then $v_{(1)} > \lambda^\top z_0 - \frac{t}{2}$, and the error at every $\alpha \le p$ is at least $v_{(1)} - (\lambda^\top z_0 - t) > \frac{t}{2}$. Otherwise $v_{(1)} \le \lambda^\top z_0 - \frac{t}{2}$, and the error at $\alpha = \frac{1}{M}$ is at least $\lambda^\top z_0 - pMt - v_{(1)} \ge \frac{t}{2} - pMt = \frac{t}{2}(1 - 2pM)$.

Part (ii)(c). With $\alpha \triangleq 2p < \frac{1}{M}$, η_t is the law of Proposition 5, and $\eta_t \in \mathcal{P}(B_R)$ by its part (iii). By its part (i) the w_1-optimal quantizer is δ_{z_0}, and by its part (ii) the w_2-optimal quantizer has $M-1$ atoms at z_0 and one at $z_0 - pMt\lambda$, whose scalarization $\lambda^\top z_0 - pMt$ exceeds $\lambda^\top z_0 - \frac{t}{2}$ since $pM < \frac{1}{2}$. The quantile levels $(2i-1)/(2M) \ge \frac{1}{2M} > p$ all sit at $\lambda^\top z_0$. So all three put no mass on A, (a) applies, and the conditional law given A is undefined. For the entropic risk, $\mathbb{E}_{\eta_t}\left[e^{-\beta \lambda^\top Z}\right] = e^{-\beta \lambda^\top z_0}(1 - p + pe^{\beta t})$, so δ_{z_0} errs by $\frac{1}{\beta}\log(1 - p + pe^{\beta t}) \ge \frac{1}{\beta}\log(pe^{\beta t}) = t - \frac{1}{\beta}\log\frac{1}{p}$.

Part (iii). Write $A = A_{\lambda,\tau}$. The triangle inequality through $\hat{Q}_\psi$ and $\mathsf{H}_{\sigma_{N_g}}\mu$ splits the error into three terms. The events $A_{\lambda,\tau}$ are the sublevel sets of the scalar $\lambda^\top X$, so $|\hat{Q}_N(A) - \hat{Q}_\psi(A)|$ is the deviation of the empirical distribution function of $\lambda^\top \hat{X}_1, \dots, \lambda^\top \hat{X}_N$ from its population version at τ, and the Dvoretzky–Kiefer–Wolfowitz inequality with Massart's constant (Massart, 1990) bounds it by $\sqrt{\log(2/\beta)/(2N)}$ for all τ at once, with probability at least $1-\beta$. The second term is at most $\mathrm{TV}(\hat{Q}_\psi, \mathsf{H}_{\sigma_{N_g}}\mu) \le \varepsilon_{\mathrm{TV}}$ by Theorem 3. For the third, let $X = Y + \sigma_{N_g}\xi$ with $Y \sim \mu$ and $\xi \sim \mathcal{N}(0, \mathbf{I})$ independent, so that $|(\mathsf{H}_{\sigma_{N_g}}\mu)(A) - \mu(A)| \le \Pr(\mathbb{1}_A(X) \neq \mathbb{1}_A(Y))$. The two indicators differ only if $|\lambda^\top Y - \tau| < w$, or if $\lambda^\top Y$ lies at distance at least w from τ and the scalar $\sigma_{N_g}\lambda^\top\xi \sim \mathcal{N}(0, \sigma_{N_g}^2)$ crosses this distance in the one direction that changes the indicator, which has conditional probability at most $\bar{G}(w/\sigma_{N_g})$. Only the first term is random, so the bound holds for all τ and w at once.

Part (iv). The same split applies, with the first term bounded uniformly over all closed half-spaces by the Vapnik–Chervonenkis inequality $\Pr\{\sup_A |\hat{Q}_N(A) - \hat{Q}_\psi(A)| > \epsilon\} \le 8\,\mathcal{S}(N) e^{-N\epsilon^2/32}$, where the shatter coefficient of the closed half-spaces of $\mathbb{R}^d$ is $\mathcal{S}(N) \le (N+1)^{d+1}$, as their Vapnik–Chervonenkis dimension is $d+1$ (Devroye et al., 1996, Theorem 12.5 and Chapter 13). Solving $8(N+1)^{d+1}e^{-N\epsilon^2/32} = \beta$ gives the stated term, and the other two terms hold for every λ, τ and w.

Part (v). The map $b_{0,\gamma}$ is a bijection with $b_{0,\gamma}^{-1}(\gamma B) = B$, which gives the first identity, and a pushforward by a bijection preserves the total variation, which gives the second. Iterating both along the chain gives the last claim. □

Since $\bar{G}(u) \le \frac{1}{2}e^{-u^2/2}$ for $u \ge 0$, (iii) and (iv) give the informal statement of Proposition 1(iii) whenever the slab of half-width w around the event boundary carries no target mass. For $\mu = \eta_t$, $\tau = \lambda^\top z_0 - t/2$ and $w = t/2$ this slab is empty, so the decision of (ii)(a) computed from $\hat{Q}_N$ is correct for every δ with $|\delta - p|$ larger than the right side of (19).

A.9 Smoothing Bias and Score Regularity

Proof of Proposition 2. Let $Y \sim \eta$ and $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$ be independent. *(i)* The pair $(Y + \sigma\varepsilon, Y)$ couples $\mathsf{H}_\sigma\eta$ and η, so $w_2^2(\mathsf{H}_\sigma\eta, \eta) \le \mathbb{E}\left[\|\sigma\varepsilon\|^2\right] = \sigma^2 d$. *(ii)* Let $A \triangleq \{y_k\}$ be the set of atoms and (X, Y') any coupling of $(\mathsf{H}_\sigma\eta, \eta)$. Since $Y' \in A$ almost surely, $\|X - Y'\| \ge \mathrm{dist}(X, A)$, so $w_2^2(\mathsf{H}_\sigma\eta, \eta) \ge \mathbb{E}\left[\mathrm{dist}(X, A)^2\right]$ for $X \sim \mathsf{H}_\sigma\eta$, and we may take $X = Y + \sigma\varepsilon$. If $\sigma\|\varepsilon\| < \Delta/2$, then every atom $y_j \neq Y$ satisfies $\|X - y_j\| \ge \|Y - y_j\| - \sigma\|\varepsilon\| > \Delta/2 > \|X - Y\|$, so $\mathrm{dist}(X, A) = \sigma\|\varepsilon\|$. Therefore

$$\mathbb{E}\left[\mathrm{dist}(X, A)^2\right] \ge \sigma^2\,\mathbb{E}\left[\|\varepsilon\|^2 \mathbb{1}\{\|\varepsilon\|^2 < \Delta^2/(4\sigma^2)\}\right] = \sigma^2 d\left(1 - \Pr(\chi^2_{d+2} \ge \Delta^2/(4\sigma^2))\right),$$

where the equality uses $x f_d(x) = d\, f_{d+2}(x)$ for the density f_d of χ^2_d. For a single atom, $\mathrm{dist}(X, A) = \sigma\|\varepsilon\|$ always, so $w_2^2 \ge \sigma^2 d$, and (i) gives equality. *(iii)* We have $\mathsf{H}_\sigma\mathcal{N}(\mu_k, S_k) = \mathcal{N}(\mu_k, S_k + \sigma^2\mathbf{I})$. With $\xi \sim \mathcal{N}(0, \mathbf{I})$, the pair $\left(\mu_k + (S_k + \sigma^2\mathbf{I})^{1/2}\xi, \mu_k + S_k^{1/2}\xi\right)$ couples the two components, and since S_k and $S_k + \sigma^2\mathbf{I}$ commute,

$$w_2^2\big(\mathcal{N}(\mu_k, S_k + \sigma^2\mathbf{I}), \mathcal{N}(\mu_k, S_k)\big) \le \sum_{i \le d}\left(\sqrt{\lambda_{k,i} + \sigma^2} - \sqrt{\lambda_{k,i}}\right)^2 \le \frac{d\,\sigma^4}{4s^2},$$

with $\lambda_{k,i} \ge s^2$ the eigenvalues of S_k, because $\sqrt{\lambda + \sigma^2} - \sqrt{\lambda} = \sigma^2/(\sqrt{\lambda + \sigma^2} + \sqrt{\lambda}) \le \sigma^2/(2\sqrt{\lambda})$. Drawing the component index $k \sim w$ first and then coupling the two components as above couples the mixtures, so $w_2^2(\mathsf{H}_\sigma\eta, \eta) \le d\sigma^4/(4s^2)$, and taking square roots gives the claim. □

Proof of Lemma 2. Let $Z = Y + \sigma\varepsilon$ with $Y \sim \eta$ and $\varepsilon \sim \mathcal{N}(0, \mathbf{I})$ independent. The conditional law π_z of Y given $Z = z$ has a density proportional to $\exp\left(\langle z, y\rangle/\sigma^2 - \|y\|^2/(2\sigma^2)\right)$ with respect to η, an exponential family in z/σ^2, and (8) and (9) read

$$s_\sigma(z) = \frac{\mathbb{E}_{\pi_z}[Y] - z}{\sigma^2}, \qquad \nabla s_\sigma(z) = \frac{\mathrm{Cov}_{\pi_z}[Y] - \sigma^2\mathbf{I}}{\sigma^4}.$$

(i) The covariance is positive semidefinite, which gives the lower bound. For a unit vector u, $\langle u, Y\rangle$ lies in an interval of length at most D under π_z, since π_z is supported in $\mathrm{supp}\,\eta$, so $\mathrm{Var}_{\pi_z}(\langle u, Y\rangle) \le D^2/4$ by Popoviciu's inequality, which gives the upper bound. *(ii)* Let y_1, y_2 be atoms at the smallest distance δ,

$m \triangleq (y_1+y_2)/2$, $u \triangleq (y_1-y_2)/\delta$ and $z_t \triangleq m+tu$. On this line the log-odds of y_1 against y_2 under π_{z_t} is $\log(w_1/w_2)+t\delta/\sigma^2$, which vanishes at $t_\sigma \triangleq -\sigma^2\log(w_1/w_2)/\delta$. By the parallelogram identity, every other atom satisfies $\|y_k-m\|^2 = \frac{1}{2}\|y_k-y_1\|^2+\frac{1}{2}\|y_k-y_2\|^2-\frac{\delta^2}{4} \geq \frac{3}{4}\delta^2$. So for $|t|\leq\delta/8$ we have $\|z_t-y_k\| \geq (\frac{\sqrt{3}}{2}-\frac{1}{8})\delta$ and $\|z_t-y_i\| \leq \frac{5}{8}\delta$ for $i\in\{1,2\}$, and since $(\frac{\sqrt{3}}{2}-\frac{1}{8})^2-(\frac{5}{8})^2 > \frac{1}{8}$, the total mass ϵ_σ of the other atoms under π_{z_t} is at most $\frac{K}{\min(w_1,w_2)}e^{-\delta^2/(16\sigma^2)}$, uniformly in $|t|\leq\delta/8$. For σ small enough, $|t_\sigma|\leq\delta/8$, and $\pi_{z_{t_\sigma}}$ is the mixture of the equal-weight two-point law on $\{y_1,y_2\}$, with weight $1-\epsilon_\sigma$, and a law on the other atoms, so the law of total variance gives $\mathrm{Var}_{\pi_{z_{t_\sigma}}}(\langle u,Y\rangle) \geq (1-\epsilon_\sigma)\delta^2/4$. So $\sigma^4\mathrm{Lip}(s_\sigma) \geq \sigma^4\langle u,\nabla s_\sigma(z_{t_\sigma})u\rangle \geq (1-\epsilon_\sigma)\delta^2/4-\sigma^2$, which tends to $\delta^2/4$ as $\sigma\to 0$. *(iii)* Here $\mathsf{H}_\sigma\eta = \sum_k w_k\mathcal{N}(\mu_k,\Sigma)$ with $\Sigma \triangleq S+\sigma^2\mathbf{I} \succeq (s^2+\sigma^2)\mathbf{I}$. The posterior $\pi_z(k) \propto w_k\exp\big(\langle\Sigma^{-1}z,\mu_k\rangle-\frac{1}{2}\langle\mu_k,\Sigma^{-1}\mu_k\rangle\big)$ is an exponential family in $\Sigma^{-1}z$, and

$$s_\sigma(z) = \Sigma^{-1}\big(\mathbb{E}_{\pi_z}[\mu_K]-z\big), \qquad \nabla s_\sigma(z) = \Sigma^{-1}\mathrm{Cov}_{\pi_z}[\mu_K]\Sigma^{-1}-\Sigma^{-1}.$$

The first term is positive semidefinite with norm at most $\|\Sigma^{-1}\|^2\Delta^2/4 \leq \Delta^2/(4(s^2+\sigma^2)^2)$, by Popoviciu's inequality for the means, and the second is positive definite with norm at most $1/(s^2+\sigma^2)$. The eigenvalues of a difference $A-B$ of positive semidefinite matrices lie in $[-\|B\|,\|A\|]$, so $\mathrm{Lip}(s_\sigma) \leq \max\{(s^2+\sigma^2)^{-1},\Delta^2/(4(s^2+\sigma^2)^2)\} \leq \max\{s^{-2},\Delta^2/(4s^4)\}$. □

The constants are sharp: on random one-dimensional mixtures of two to four components with a shared width $s\in[0.05,2]$, means in $[-5,5]$ and σ from 0 to 3, the largest ratio of $\mathrm{Lip}(s_\sigma)$ to the bound of (iii) is one up to grid error, and for two atoms the supremum in (ii) matches $\sigma^{-4}(\delta^2/4-\sigma^2)$ numerically.

The regularity decides the stability of the sampler. On a geometric grid $h_n = (1-\kappa_{\mathrm{g}}^2)\sigma_n^2$, so by (ii) $h_n\mathrm{Lip}(s_{\sigma_n}) \geq (1-\kappa_{\mathrm{g}}^2)\big((1-\epsilon_{\sigma_n})\delta^2/(4\sigma_n^2)-1\big)$ grows without bound on atoms, while by (iii) $h_n\mathrm{Lip}(s_{\sigma_n}) \leq (1-\kappa_{\mathrm{g}}^2)\sigma_n^2\max\{s^{-2},\Delta^2/(4s^4)\} \to 0$ on mixtures with densities. An explicit Euler step of the sampler is unstable once $h_n\mathrm{Lip}$ exceeds a constant of order one, and the sampler of the atomic chain shows such escapes, which disappear with reward noise (Section D.3).

B METHOD DETAILS

This section completes Section 3 with the details that reproduce the reported runs. The tabular variant runs on the synthetic MDPs, the rare-event chain, the chain of Debes and Tuytelaars (2026b) and resource gathering, and SAC-SVDRL on MuJoCo; Section B.1 lists every setting.

Stratified Bellman Target Set. The M successors of an update are drawn by systematic resampling (Kitagawa, 1996), so a branch of probability p receives $\lfloor Mp\rfloor$ or $\lceil Mp\rceil$ targets. Within the block of n_k draws of one successor, the standard normal reward noise is a Latin hypercube sample (McKay et al., 1979): each coordinate takes each of the n_k equal-probability strata of $\mathcal{N}(0,1)$ exactly once, at a uniform position within the stratum, with the strata permuted independently per coordinate. Each draw is still marginally a successor and reward of the model, so the targets stay unbiased, while their empirical law matches the branch masses up to $1/M$ and spreads the noise evenly over the quantiles of each coordinate. The successor particle of a draw is one of the current particles at the successor pair.

Resample-Move Update. Each update replaces $\mathrm{round}(\rho M)$ particles by targets and runs the Stein move on the cloud (Algorithm 1). The relaxation $z \leftarrow z+\alpha(\tilde{z}-z)$ then moves the cloud toward the transported cloud $\tilde{z}$. The annealing schedule has ten geometric levels σ_ℓ from 5 to 0.1 in PopArt-normalized units, and the Stein move takes $T_\ell = 10$ SVGD steps at every level $\sigma_\ell \leq \sigma_{\mathrm{start}}$. The SVGD direction uses the RBF kernel with the bandwidth recomputed at every step, $h_\ell = \mathrm{clip}(h_{\mathrm{med}}, 0.1\,\sigma_\ell, 2\,\sigma_\ell)$, where h_{med}^2 is the median of the squared pairwise distances divided by $2\log(M+1)$, so that a balanced bimodal cloud at a low level keeps h_ℓ of the order of σ_ℓ and not of the mode separation. The steps are preconditioned by Adam (Kingma and Ba, 2015), with step size 0.1 and a state per particle and coordinate that is carried across the levels of one update and reset at the next, and every step is followed by the coordinatewise projection onto the return box B_R. The single configuration replaces the whole cloud, $\rho = 1$, with $\alpha = 0.001$ and $\sigma_{\mathrm{start}} = 1$; it is the selected configuration on MDP2, MDP3, MDP4, MDP6 and MDP7. On MDP1 the selected configuration uses $\alpha = 0.1$, and on the unimodal MDP5 and MDP8 it transports the previous cloud through the whole schedule toward the stratified target set, $\rho = 0$ and $\alpha = 1$.

Algorithm 2 Reverse-diffusion sampler, analyzed in Theorem 3

Require: s_ψ, (x, a), DSM grid $\sigma_0 > \cdots > \sigma_{N_g}$ with $\sigma_0^2 \gg m_2(\mu; z_c)$ (Remark 2), sample count N (vectorized)
1: $z \sim \mathcal{N}(z_c, \sigma_0^2 \mathbf{I})$, with $z_c \in B_R$ the standardizing mean at (x, a), fixed before sampling (Definition 2)
2: **for** $n = 0, \ldots, N_g - 1$ **do**
3: $\quad h \leftarrow \sigma_n^2 - \sigma_{n+1}^2; \quad z \leftarrow z + h\, s_\psi(z, x, a, \sigma_n) + \sqrt{h}\,\xi, \quad \xi \sim \mathcal{N}(0, \mathbf{I})$
4: **end for**
Ensure: $\mathrm{proj}_{B_R}(z)$ (coordinatewise clip), denormalized

Moment Matching. After the Stein move the cloud is mapped affinely onto the moments of the target set and projected onto B_R, $z_i \leftarrow \mathrm{proj}_{B_R}\big(m_y + (z_i - \bar{z}) \odot s_y / s_z\big)$, with the per-dimension means $m_y, \bar{z}$ and standard deviations s_y, s_z of the targets and of the cloud; on MDP8 the map matches the full covariance instead. It corrects the spread, since SVGD with finitely many particles underestimates the variance of its target, increasingly so in higher dimensions (Ba et al., 2022), and on a self-looping state–action pair the bootstrap compounds a per-update variance bias by the factor $1/(1-\gamma^2)$, as in Remark 5. Sampler bootstrapping applies the same map to the sampler draws. On the chain the draws are instead shifted, at evaluation only, to a running mean of the bootstrap targets at (x, a), an exponential average with weight $\max(1/128, 1/t)$ at update t; we call this running-mean moment matching, and given to the competitors' particles it is the second of the two components that transfer (Section D.3).

Score Network. The tabular variant keeps one score network per state–action pair, an MLP with three hidden layers of 128 Swish units in the energy parameterization $E_\psi(z, \sigma) = c_{\mathrm{out}} f_\psi(c_{\mathrm{in}} z, c_{\mathrm{noise}}) - \|z\|^2/(2(\sigma^2 + \sigma_d^2))$ with the EDM scalings (Karras et al., 2022), whose score is $s_\psi = \nabla_z E_\psi$ and whose noise level enters every hidden layer by feature-wise affine modulation (FiLM); the prior term makes the score grow linearly in z, and Remark 3 applies Theorem 3 to it. It reads its input in standardized coordinates $u = (z - m)/s$, where m is the mean of the bootstrap targets at (x, a) and s one scale for all dimensions, the largest standard deviation of the targets, both exponential averages with weight 0.02 on the new value; the EDM scalings assume $\sigma_d \approx 1$, which holds only after this standardization. Each Bellman step takes two Adam steps (learning rate 10^{-3}) on a batch of 256 samples of the DSM loss (4) and the consistency loss (5) with $\lambda_{\mathrm{cons}} = 1$ and the weight $\lambda(\sigma) = \sigma^2$, each sample with its own noise level, drawn log-uniformly from $[0.002, 80]$ in standardized units. The target of (5) is the EMA network of the successor pair, and the online network under a stop-gradient where the successor is the pair being updated. The level grid of Theorem 2, with ratio $\gamma^{1/m}$, is a device of the analysis; the network is conditioned on a continuous noise level.

PopArt and the Return Box. Returns are normalized per reward dimension by PopArt, with exponential averages of weight $3 \cdot 10^{-4}$, a scale floor, and an output correction that preserves the predictions. A normalized return $\hat{y} = (y - m)/s$ satisfies $\hat{y} = \hat{r} + \gamma \hat{z}$ with $\hat{r} = (r - (1-\gamma)m)/s$, the same affine bootstrap with the same γ, so Lemma 1 and (5) hold verbatim in normalized space, which a nonlinear transform such as symlog would break. The return box B_R has the half-width $r_{\max}/(1-\gamma)$ per reward dimension; in the tabular variant $r_{\max} = \max_c(|m_c| + 5 s_c)$ over the reward channels c with mean m_c and standard deviation s_c, and in SAC-SVDRL it is the running maximum of the absolute soft reward, which keeps the soft bootstrap inside B_R.

Sampler and Estimators. The sampler is the reverse diffusion of Definition 2 on the EMA score network (Algorithm 2). On the synthetic MDPs and resource gathering it takes Euler steps on 512 geometric levels from $\sigma_u = 80$ to 0.002 in standardized units and returns $N = 4096$ draws. In the rare-event study SVDRL bootstraps from moment-matched draws of its EMA sampler, $4M$ per target set, and answers every query with $N = 10^5$ draws. On the chain the sampler takes 256 Euler steps with a projection onto B_R at every step, its mean is matched to a debiased exponential average (weight 1/128) of the target means, and it returns 4096 or 128 draws. Event probabilities, CVaR and decisions are read from the draws by plug-in, and conditional queries by rejection (Corollary 11).

Hybrid Rule. The rule of Definition 1 runs on the particles at (x, a) without ground truth. First, particles closer than the field's smallest noise level are merged, i.e., connected components of the graph with an edge whenever $\|(z_i - z_j)/\varepsilon\| < 1$, with ε the smallest noise level in return units per dimension, each component replaced by its mean. Single-linkage clusters are then found top-down on the Euclidean minimum spanning tree

of the merged points: a cluster is split at its longest edge whenever that edge exceeds three times the larger standard deviation of its two sides, the standard deviation being the square root of the largest eigenvalue of the count-weighted covariance. The ratio r is the shortest accepted cut over the largest within-cluster standard deviation, and $r = 0$ for a single cluster. The rule reports the particles if $r > 6$ or if a cluster of at least two particles has a per-dimension standard deviation below ε in every dimension, and the sampler draws otherwise. The threshold 6 was fixed in advance, and the thresholds 4 and 10 were evaluated on tuning seeds only. The second clause was added after resource gathering exposed atoms of coincident particles, and it changes no entry of Table 2.

SAC-SVDRL. SAC-SVDRL uses two distributional critics with M particles each and a score network per critic, all four networks with EMA targets. The bootstrap targets use the rank-wise minimum of the two target critics' sorted particles, with the soft reward $r - \gamma\alpha \log \pi(a' \mid x')$ and one-step targets. A moment head regressed onto the mean and the central second moment of the bootstrap targets sets the standardization of the score network, which is conditioned on (σ, x, a) by FiLM, and the moment-matching target; the noise levels and the SVGD step are scaled by its predicted standard deviation, so that the annealing acts on a law of unit spread. The Stein move takes $T_\ell = 2$ steps at each of the ten levels, and each critic is regressed onto its transported cloud by matching ranks,

$$L_{\text{critic}}(\theta) = \textstyle\sum_i \left\| z^{\theta}_{(i)}(x, a) - \text{sg}\left[z^{\text{SVGD}}_{(i)}(x, a) \right] \right\|^2, \tag{20}$$

where $z_{(i)}$ is the i-th smallest particle of the one-dimensional return, which is the w_2-optimal matching in $d = 1$. For a single Euler step of the transport the gradient of (20) is, by the chain rule, $-2\varepsilon \sum_i J_i^\top \hat{\phi}^*(z_i^\theta)$ with the Jacobians J_i of the critic's particles, the direction of amortized SVGD (Feng et al., 2017). The policy minimizes $\mathbb{E}_x \left[\mathbb{E}_{a \sim \pi(\cdot|x)} \left[\alpha \log \pi(a|x) - Q(x, a) \right] \right]$ with Q the smaller of the two critics' particle means, the entropy coefficient α is adapted as in SAC, and every head uses the same running observation normalization. HalfCheetah uses this default configuration; on Walker2d, Hopper and Ant moment matching takes its mean from the bootstrap targets of the sampled transitions instead of the moment head, which gives the critic the temporal-difference signal directly. A variant with a log-variance moment head removes the floor of the linear variance head that made one Ant seed diverge (Section D.5).

B.1 Hyperparameters

Configurations were chosen by their performance on tuning seeds: five for the synthetic MDPs, the chain and resource gathering, and two for the rare-event study (Section C). Every reported tabular, chain and rare-event number comes from seeds that the selection did not use. For MuJoCo, SVDRL uses the default configuration on HalfCheetah and the variant that takes the moment-matching mean from the bootstrap targets on Walker2d, Hopper and Ant. On Walker2d and Ant both variants ran on five seeds, and the one with the higher IQM return of Table 3 was declared and extended to ten seeds; the other is reported in Tables 17 and 18. Hopper ran the declared variant only. Table 4 lists the most important settings of the reported runs per experiment, read from their configurations. All agents are implemented in JAX (Bradbury et al., 2018) with Equinox (Kidger and Garcia, 2021), and jaxtyping (Kidger, 2022) checks the shapes and dtypes of their arrays. Code that reproduces every run will be released upon acceptance.

C PROTOCOLS

This section specifies the environments, the statistics and the metrics of Section 5, and the measured compute. Gymnasium and MO-Gymnasium are released under the MIT license and MuJoCo under the Apache 2.0 license; we reimplement the chain of Debes and Tuytelaars (2026b) and the maze of Zhang et al. (2021) from their published definitions.

Synthetic MDPs. All synthetic MDPs have $\gamma = 0.95$ and Gaussian rewards, deterministic at the root states, so their return laws are Gaussian mixtures. MDP1 has a single absorbing state with reward $\mathcal{N}(1, 0.5^2)$. In MDP2 the root has reward 0 and branches with probabilities 0.3 and 0.7 into absorbing states with rewards $\mathcal{N}(1, 0.3^2)$ and $\mathcal{N}(2, 0.4^2)$. MDP3 is two-dimensional with three modes, reached from the root with probabilities 0.4, 0.35 and 0.25. MDP4 is five-dimensional with three modes reached with the same probabilities, whose mean rewards are 3 times standard normal draws and whose covariances have eigenvalues drawn uniformly from $[0.05, 0.5]$ in

Table 4: *Hyperparameters.* Settings of the reported runs per experiment, read from their configurations. Configurations: the number of distinct ones, tuning and ablations included; a list gives the values they take; –: not a setting of the method. An annealing start at the highest level runs the whole schedule.

		Synthetic MDPs	Rare events	Chain	Resource gathering	MuJoCo	Pixel maze
SVDRL							
Configurations		62	51	14	18	9	6
Particles	M	$64d^a$	32, 128, 512, 2048	128	32, 128, 512	64	200
Sampler draws	N	4096	10^5	128, 4096	4096, 10^5	–	–
Resampled fraction	ρ	0, 1	0, 0.5	0, 1	0, 1	–	–
Relaxation	α	0, 10^{-3}, 0.1, 1	1	0, 0.1, 1	10^{-3}, 1	–	–
Annealing start	σ_{start}	1, 5	0.6, 5	1, 5	1, 5	5	–
SVGD steps per level	T_ℓ	10	10	10	10	2	10
Annealing levels		10	10	10	10	10	10
Level range	σ_ℓ	$[0.1, 5]$	$[0.1, 5]$	$[0.1, 5]$	$[0.1, 5]$	$[0.1, 5]$	$[0.1, 5]$
Consistency weight	λ_{cons}	1	1	1	1	1	0
Score learning rate		10^{-3}	10^{-3}	10^{-3}	10^{-3}	10^{-3}	10^{-3}
Score batch size		32, 256	32, 256	32	256	–	–
Competitors							
Methods		10	7	6	7	4	3
Configurations		191	115	100	51	16	9
Particles	M	$64d^a$	32, 128, 512, 2048	128, 512, 2048, 4096	32, 128	64	200
Quantiles		64	1000	–	1000	32	–
C51 atoms		101	–	–	–	–	–
Learning rate†		$2.5 \cdot 10^{-3}$ to 0.3^b	–	$2.5 \cdot 10^{-3}$	–	–	–
All methods							
Discount	γ	0.95	0.95	0.9	0.95	0.99	0.99
Batch size		–	–	–	–	256	–
Critic learning rate		–	–	–	–	$3 \cdot 10^{-4}$	10^{-5}

[a] $M = 64d$ on an MDP of dimension d; the compute-matched comparison adds $M \in \{512, 4096\}$ and the dimension sweep $M \in \{140, 200, 520\}$. [b] The values $\{2.5 \cdot 10^{-3}, 3 \cdot 10^{-3}, 5 \cdot 10^{-3}, 0.01, 0.03, 0.1, 0.3\}$. † The variants of Debes and Tuytelaars (2026b). Multi-objective MuJoCo runs IQN with 32 quantiles and the other heads with $M = 64$ particles, and the defaults of the code otherwise.

a random orthonormal basis, all drawn with the fixed random key 42, which is part of the specification; the dimension sweep builds the MDP4 family in the same way for $d \in \{2, 3, 5, 8, 10\}$.

MDP5 to MDP8 have one decision state whose actions lead to absorbing reward channels with i.i.d. Gaussian rewards from $t = 0$ on, so each channel's return is Gaussian with $1/(1-\gamma)$ times the mean and $1/(1-\gamma^2)$ times the covariance of its reward. MDP5 ($d = 2$) has four deterministic actions with mean rewards $(3, 0.5)$, $(2, 2)$, $(0.5, 3)$ and the dominated $(1, 1)$, all with covariance $0.04\,\mathbf{I}$. MDP6 ($d = 2$) has a safe action with mean reward $(2, 1)$, a risky action that reaches $(4, 0.8)$ or $(0.2, 1)$ with probability $\frac{1}{2}$ each and variance 0.09 in the first objective, and an action with $(0.5, 3)$; all other variances are 0.04. MDP7 ($d = 3$) has five deterministic actions with mean rewards $(3, 0.5, 0.5)$, $(0.5, 3, 0.5)$, $(0.5, 0.5, 3)$, $(2, 2, 0.3)$ and $(0.3, 2, 2)$ and full covariances, with diagonal entries 0.03 to 0.06 and correlations of both signs. MDP8 ($d = 2$) has two actions with the same mean reward $(2, 2)$ and the covariances $\left(\begin{smallmatrix} 0.5 & \mp 0.45 \\ \mp 0.45 & 0.5 \end{smallmatrix}\right)$, hedged and concentrated, and a low-variance action with mean reward $(0.5, 3)$ and covariance $0.04\,\mathbf{I}$. The multimodal evaluated state–action pairs are those of MDP2, MDP3, MDP4 and MDP6. Every run takes 8000 Bellman steps on MDP1 and MDP2 and 16 000 on the other MDPs, with an evaluation every 800 steps up to 80 % of training and every 100 steps afterwards.

Choice of the Synthetic MDPs. Each MDP probes one property of the return law on which the critic families of Table 1 differ, and every return law is a Gaussian mixture, so every error is measured against the exact law. MDP1 and MDP2 are one-dimensional. MDP1 bootstraps through a self-loop, where any per-update bias in the spread of the cloud is amplified by the factor $1/(1-\gamma^2)$, and MDP2 has two modes with unequal masses and widths. MDP3 and MDP4 have three separated modes in two and five dimensions, with masses 0.4, 0.35 and 0.25, which $M = 64d$ particles hold exactly only at $d = 5$ and $d = 10$, so the MDP4 family separates the resolution floor of Proposition 1 from the growth of the error with the dimension. MDP5 to MDP8 have one decision state each and differ in what separates its actions. In MDP5 three Pareto-optimal actions and a dominated one differ

only in their means, the baseline for decisions. In MDP6 a safe and a risky action have almost the same mean, $(2, 1)$ against $(2.1, 0.9)$, and differ in their tails. MDP7 has correlated objectives in three dimensions. In MDP8 two actions have the same mean and the same marginals and differ only in the sign of the correlation, which no model of the per-dimension marginals can resolve. Together the eight MDPs cover unimodal laws, on which the within-mode shape decides (MDP1, MDP5, MDP7 and MDP8), and separated modes, on which the mode masses decide (MDP2, MDP3, MDP4 and MDP6), in one to five dimensions and up to ten in the dimension sweep.

Gaussian rewards make the returns unbounded, so Assumption 1 holds for the synthetic MDPs only up to a truncation. The return box of Section B keeps every channel's return within $5\big((1+\gamma)/(1-\gamma)\big)^{1/2}$ standard deviations of its mean in each coordinate, about 31 at $\gamma = 0.95$, so the true laws and their truncations, which the theory covers, coincide for every practical purpose.

Rare-Event Chain. The rare-event chain has $d = 2$, $\gamma = 0.95$ and i.i.d. Gaussian per-step rewards with covariance $0.04\,\mathbf{I}$ in every channel. At the root, the risky action reaches a bulk channel with mean reward $(3, 1)$ with probability $1 - p$ and a rare channel with mean reward $(-4, 2)$ with probability p, and the safe action reaches a channel with mean reward $(2.5, 1.5)$. The rare channel is adverse in the first coordinate and shifted in the second, so conditional queries on the catastrophe are not trivial. The branching state is preceded by $D \in \{0, 2\}$ deterministic transient states with reward 0, so for $D = 2$ the rare branch is decided two bootstrapping steps below the root, whose return law is γ^D times that of the branching state. At $D = 0$ the return means are $(60, 20)$ for the bulk, $(-80, 40)$ for the rare and $(50, 30)$ for the safe channel, with standard deviation 0.64 per coordinate. We use $p \in \{0.002, 0.005, 0.01\}$, so $p < 1/M$ already at $M = 128$, and the event $Z_1 \le -10\gamma^D$, between the bulk and the catastrophe. The ground truth is in closed form and agrees with $2 \cdot 10^5$ Monte Carlo returns. Every run takes 16 000 Bellman steps for $D = 0$ and 28 800 for $D = 2$, with 50 evaluations.

Chain. The chain of Debes and Tuytelaars (2026b) has $K = 10$ reward dimensions, forward probability 0.9, $\gamma = 0.9$ and the rewards $e_{S'} + \sigma_r \varepsilon$ with $\sigma_r \in \{0, 0.05, 0.2, 0.5\}$. Every method uses $M = 128$ particles and near-exact TD targets for $5 \cdot 10^4$ updates, and the ground truth uses 10^6 Monte Carlo rollouts per start state. The six variants of Debes and Tuytelaars (2026b) run on their own target sets and on ours, with RAdam at learning rate $2.5 \cdot 10^{-3}$, gradient clipping at 0.1 and 64 projections; with more particles, sliced w_2 and sliced Cramér at $M \in \{512, 2048, 4096\}$ and MMD at $M = 512$, for $\sigma_r \in \{0.05, 0.2\}$. SVDRL replaces the cloud by the stratified target set with $\alpha = 0.1$ and trains the score network of Section B alongside.

Resource Gathering. The resource-gathering environment of MO-Gymnasium (Felten et al., 2023) is a 5×5 grid with $d = 3$, which we reimplement with the attack probability as a parameter, $p \in \{0.1, 0.02, 0.005\}$, where the package fixes 0.1. We use $\gamma = 0.95$ and three fixed policies from value iteration (safe, mixed and risky), and the ground truth comes from exact enumeration. SVDRL runs its single tabular configuration without tuning at $M \in \{32, 128\}$, reported as particles, sampler and hybrid, and the rare-event configuration with $N = 10^5$ sampler draws; the particle critics run at $M \in \{32, 128\}$ (Sinkhorn at 32 only), and a quantile critic on the oracle scalarization with its learning rate tuned on the five tuning seeds. Every run takes 15 680 Bellman steps with 50 evaluations.

MuJoCo. HalfCheetah, Walker2d, Hopper and Ant (v5) (Towers et al., 2024) run for 10^6 environment steps with $\gamma = 0.99$ and $M = 64$ particles, with an evaluation of ten deterministic episodes every 10^4 steps. The heads are SAC and SAC with IQN, MD3QN, energy MMD and SVDRL critics on the same trunk. For the critic shape, the policy at 10^6 steps is frozen, and 64 fixed probes each receive 1000 Monte Carlo rollouts of 1000 steps of the soft return; on Walker2d only the probes that restart exactly are used.

Statistics and Selection. In the tabular experiments and on the chain a seed's error is its mean over the evaluations in the last 20 % of training (on the chain, the last eight of 40), and we report the interquartile mean (IQM) or the median over seeds with 95 % percentile-bootstrap intervals. Configurations are selected on tuning seeds and reported on final seeds that the selection did not use: five tuning and ten final seeds for the synthetic MDPs, the chain and resource gathering, and two tuning and ten final seeds for the rare-event study, which adds ten fresh seeds. On the synthetic MDPs and the chain a comparison is a win or a loss when the unpaired Welch t statistic of SVDRL's errors against the competitor's is at most -2 or at least 2, and a tie otherwise; the best competitor is the row with the lowest seed mean, per MDP, block and framing. In the rare-event study SVDRL is counted as more accurate when its IQM error is smaller than that of the best of the six particle critics at the

same M, answering from its particles or from a Gaussian kernel density estimate on them (Scott's bandwidth), without a significance test. The three framings of the synthetic MDPs compare SVDRL with the competitors on their own target sets, with every row including the competitors on our target set (marked *), and with the best competitor at $M \in \{512, 4096\}$ on the four unimodal MDPs (Table 9). No competitor is limited by the training budget: on the eight MDPs none improves significantly over the last 20 % of its updates, and SVDRL reaches its plateau early in training. On MuJoCo a run's return is its mean evaluation return over 800k to 1M steps (21 evaluations), and we report per environment and head the IQM over five to ten seeds with a 95 % percentile-bootstrap interval; a run stopped by the divergence guard counts with the evaluations it logged.

Metrics. On the synthetic MDPs, mass is the total variation between the mode masses of a cloud and the ground-truth weights, next to its floor $\min_k \frac{1}{2}\sum_i |k_i/M - w_i|$ over integer allocations k, and shape is the within-mode error $\sum_i w_i w_2^2$, in closed form for $d = 1$, by exact assignment against 4096 ground-truth draws for $d = 2$, and by quantile sliced w_2^2 with 128 projections for $d > 2$. The full w_2^2 charges a mass error ϵ across a gap Δ as $\epsilon\Delta^2$, so on separated modes it measures mass rounding rather than shape (Table 10 and Figure 6). Sliced w_2 has two independent estimators, the quantile estimator and a CDF-based estimator that integrates the analytic projected distribution functions of the ground truth (Vauthier et al., 2026), and a disagreement between them fails the asset pipeline. In the rare-event study the queries are CVaR down to $\alpha = 0.001$, $|\log \hat{P}/P|$ of the event (with a floor of 10^{-6} on $\hat{P}$), chance-constrained decisions over thresholds δ, conditional tails, entropic risk and the expected scalarized return, ten rare-event and two bulk query columns, i.e., 60 rare-event and 12 bulk comparisons over the six environments. On the chain we measure w_2 by exact assignment, sliced and per-coordinate w_2^2, the mean error and the dispersion against 10^4 fresh Monte Carlo rollouts, and the squared energy distance against the 10^6-rollout bank. On resource gathering we measure $|\log \hat{P}/P|$ of death, the error rates of the chance-constrained and the CVaR decisions, CVaR errors, the error of $\mathbb{E}\,[\text{treasure} \mid \text{alive}]$, the squared energy distance and sliced w_2^2, and $|\log \hat{P}/P|$ of the time-averaged $\hat{P}$. On MuJoCo the critic shape is the w_2 distance between the critic's return law and the Monte Carlo law, each centered at its own mean, over the Monte Carlo standard deviation, exact in $d = 1$; per run we take the median over the two critics and then over the probes, and per environment and head the IQM over seeds with a 5000-resample bootstrap. Every head's mean is biased by the twin-critic minimum, so the centered distance is the one that compares the heads. Every evaluation draws fresh randomness, and a size mismatch between a cloud and its reference fails loudly instead of subsampling.

C.1 Compute

The Stein move costs at most $(L+1)T_\ell M$ score evaluations per state–action pair and update, plus $\mathcal{O}(M^2)$ kernel evaluations per step, independently of the horizon. A sampler run of N draws costs NN_g score evaluations, vectorized over the draws, and it runs only at evaluation time and inside sampler bootstrapping. Acting takes one forward pass of the critic, which returns the M Stein particles. Every GPU run used an NVIDIA GeForce RTX 4090, and training runs in single precision.

D ADDITIONAL RESULTS

This section gives the full tables and figures behind Section 5, in its order, and the checks of the theory on trained networks.

D.1 Rare Events

Table 5 lists every estimator behind the comparisons of Section 5 at $p = 0.005$ and $D = 2$. SVDRL's sampler is more accurate than every particle critic in every column, from their atoms or from a kernel density estimate. The quantile critic on the oracle scalarization matches it at $\text{CVaR}_{0.01}$ of that scalarization but not at $\text{CVaR}_{0.001}$, and it answers no other scalarization. SVDRL's own Stein particles share the floor of the particle critics, so the gain comes from the sampler.

Particle critics resolve the catastrophe only with about $10/p$ particles (Figure 4b). With $M = 2048$ sliced w_2 reaches the $\text{CVaR}_{0.001}$ error 0.30, below SVDRL's 6.9, and answers the conditional tail more accurately, while SVDRL keeps the more accurate mass, chance-constrained decisions and bulk (Section 5). SVDRL's CVaR error at the lowest levels reflects the width of the far rare mode (Figure 5), whose mass the sampler gets right but which the score network spreads too wide, 13 to 15 times in a static fit to draws of the true law (Section E).

The catastrophe survives the recursion only under sampler bootstrapping (Figure 4a), as Proposition 6(v) predicts: a target built from downstream particles gives the catastrophe a multiple of $1/M$, which is zero for $p < 1/M$ in most updates, whereas a target built from the downstream sampler inherits its event error. At $p = 0.002$ and $M = 32$ the reported configuration loses the catastrophe as well.

D.2 Eight Synthetic MDPs

Table 6 separates the objects behind the hybrid of Table 2. The hybrid reports the sampler draws on the unimodal MDPs and mostly the particles on the separated modes of MDP2, MDP3, MDP4 and MDP6, where the particles have the lower mass error. The particles are also better in shape on MDP4 and MDP6, whereas on MDP2 and MDP3 the sampler would be. The separation test alone ($r > 6$) makes the same choices on every MDP.

Table 7 gives the total variation of the mode masses on the four MDPs with separated modes, next to the within-mode shape of Table 2.

Table 8 names the best competitor behind each ratio of Tables 2 and 7, with its target set, learning rate and error.

On the unimodal MDPs the hybrid's wins come from $N = 4096$ sampler draws, and they vanish at matched compute (Table 9): competitors with $M = 4096$ particles on our target set beat the sampler on all four. SVDRL's own particles without the Stein move beat the sampler at $M = 4096$ as well, so at equal size the particle sets win on unimodal laws.

Table 10 reports the full w_2^2, which charges mass and shape together; on separated modes it is dominated by mass rounding, which every method shares (Figure 6), so Table 2 reports mass and shape apart.

Mass rounding also decides the full error across the reward dimension (Figure 6). The MDP4 family has the mode masses 0.40, 0.35 and 0.25 for every d, which are multiples of $1/M$ for $M = 64d$ only at $d = 5$ and $d = 10$, and exactly there the full error of SVDRL and of the competitors on our target set is small. SVDRL's ratio to the best competitor shows no trend with d (Figure 7).

The relaxation α of the Stein move and the stratified target set are ablated on every MDP in Table 11, and the measured gap $\varepsilon_{\mathrm{St}}$ of Theorem 4 in Table 12. On the unimodal MDPs $\varepsilon_{\mathrm{St}}$ equals the particles' true error within 1 to 12 %; on separated modes it flags a disagreement between the particles and the sampler but can exceed both of their errors.

The Stein move is specific where the selected configuration uses it (Table 13): Langevin dynamics along the same score, a Gaussian jitter of the same size and no move are all worse, with every interval of the ratio above one. Without repulsion the cloud collapses onto few distinct particles on MDP5 and MDP8.

Removing the consistency loss raises the error significantly on MDP1, MDP3, MDP4 and MDP8 and never lowers it significantly (Table 14), in line with the lower variance of the consistency target at small noise levels (Remark 1).

D.3 Chain

On the atomic chain, $\sigma_r = 0$, the persistent particles of Debes and Tuytelaars (2026b) are the most accurate in w_2 (Table 15), as Lemma 2 predicts, and the sampler's Euler steps escape near the atoms, where the score is stiff (Section A.9). With reward noise the moment-matched sampler with 4096 draws has the lowest energy distance and mean error at every $\sigma_r \geq 0.05$. The stratified target set is the first component that transfers: on it, every variant of Debes and Tuytelaars (2026b) has a lower energy distance and mean error than on its own target set at every noise level.

With more particles the competitors win shape at matched compute (Figure 8): with $M = 4096$ their energy distances and w_2^2 errors lie below those of the sampler with 4096 draws on either target set, while the sampler keeps the lower mean error.

Running-mean moment matching is the second component that transfers. Given to the variants of Debes and Tuytelaars (2026b) at evaluation, the shift of Section B to the running target mean lowers their mean error at every particle count and both noise levels, e.g., for sliced w_2 on its own target set with $M = 2048$ from 0.020 to 0.0012 at $\sigma_r = 0.05$ and from 0.025 to 0.0018 at $\sigma_r = 0.2$, below the sampler's 0.0050 and 0.010.

D.4 Resource Gathering

Table 16 gives every method, particle count and estimator behind Section 5 at the three attack probabilities. The returns are atomic, so SVDRL's sampler fails here, as Lemma 2 predicts, and the hybrid rule reports the particles, whose stratified target set holds each atom at its mass up to $1/M$. At $p = 0.005 < 1/M$ no equal-weight set holds the death probability at a single evaluation, but averaged over the evaluations and seeds SVDRL's particle estimate is 0.0053 for $P = 0.005$, since full replacement by the stratified target set rounds the mass without bias.

Table 16: *Resource gathering.* Attack probability $p \in \{0.1, 0.02, 0.005\}$ (top to bottom), every method at each particle count M, median over the ten final seeds. Death probability: $|\log \hat{P}/P|$ of the final and of the time-averaged $\hat{P}$ (avg.); safe: the absolute error under the safe policy. CVaR: errors at the preferences $(1, 1, 1)$ and $(0, 1, 1)$. Decision: error rates of the chance-constrained and the CVaR decision. Treasure: error of $\mathbb{E}[\text{treasure} \mid \text{alive}]$. Law: squared energy distance (ED^2) and sliced w_2^2. The last three columns are in units of 10^{-3}. SVDRL's rows give the hybrid and its two objects, the particles and the sampler ($N = 4096$ draws); $r > 6$ only: the hybrid rule without its sub-resolution clause; superseded†: an earlier rule whose sub-resolution clause also counted single particles, kept for reference; rare event: the rare-event configuration, which bootstraps from its sampler. Particle critics answer from their atoms. QR: the quantile critic on the oracle scalarization, with the settings of the rare-event study or with its learning rate (lr) tuned on the tuning seeds.

Errors (↓)		Death probability			CVaR		Decision		Treasure	Law	
Method	M	$\lvert\log \frac{\hat{P}}{P}\rvert$	avg.	safe	$(1,1,1)$	$(0,1,1)$	chance	CVaR	if alive	ED^2	w_2^2
Attack probability $p = 0.1$											
SVDRL, hybrid	32	0.093	0.028	0	0.019	0.013	0.030	0	1.0	2.8	3.1
SVDRL, hybrid, $r > 6$ only	32	0.52	0.41	0	0.15	0.12	0.48	0	68	73	23
SVDRL, hybrid, superseded†	32	0.093	0.028	0	0.019	0.013	0.030	0	1.0	2.8	3.1
SVDRL, particles	32	0.093	0.028	0	0.011	0.0063	0.030	0	0.012	1.1	3.1
SVDRL, sampler	32	1.0	1.0	0	0.28	0.18	0.80	0	103	128	36
SVDRL, hybrid	128	0.028	0.0083	0	0.014	0.0098	0	0	0.97	2.4	1.1
SVDRL, hybrid, $r > 6$ only	128	0.030	0.0094	0	0.015	0.011	0	0	1.1	2.8	1.1
SVDRL, hybrid, superseded†	128	0.028	0.0083	0	0.013	0.0090	0	0	0.88	2.4	1.1
SVDRL, particles	128	0.028	0.0083	0	0.0035	0.0018	0	0	0.0078	0.82	1.0
SVDRL, sampler	128	1.0	1.0	0	0.28	0.18	0.80	0	102	127	36
SVDRL, rare event, sampler	32	1.0	0.99	0	0.30	0.18	0.80	0	106	120	33
SVDRL, rare event, EMA sampler	32	0.98	0.97	0	0.29	0.18	0.80	0	105	127	36
SVDRL, rare event, sampler	128	0.96	0.96	0	0.30	0.18	0.80	0	104	121	33
SVDRL, rare event, EMA sampler	128	0.95	0.95	0	0.28	0.18	0.80	0	104	127	36
MMD-RBF	32	2.5	0.78	0	0.45	0.24	0.66	0	94	100	24
MMD-RBF	128	0.59	0.44	0	0.28	0.18	0.46	0	73	52	17
MD3QN	32	1.5	0.60	0	0.34	0.20	0.61	0	56	73	17
MD3QN	128	0.43	0.36	0	0.22	0.14	0.35	0	46	46	11
Energy MMD	32	0.43	0.12	0	0.090	0.060	0.30	0	6.7	9.0	11
Energy MMD	128	0.18	0.076	0	0.065	0.055	0.14	0	7.1	4.8	4.8
Sliced Cramér	32	0.43	0.16	0	0.099	0.067	0.34	0	9.6	11	10
Sliced Cramér	128	0.19	0.086	0	0.072	0.064	0.16	0	7.9	4.9	4.7
Sliced w_2	32	0.44	0.11	0	0.053	0.032	0.28	0	0.034	5.7	11
Sliced w_2	128	0.19	0.082	0	0.030	0.019	0.14	0	0.014	1.4	5.3
Sinkhorn	32	0.97	0.14	0	0.12	0.077	0.34	0	9.8	22	9.8
QR, rare event	1000	0.59	0.61	0.27	1.1	1.5	0.58	0.52	–	–	–
QR, rare event, 156 800 steps	1000	0.75	0.59	0.35	1.7	2.0	0.65	0.96	–	–	–
QR, oracle λ, lr 0.03	1000	0.19	0.19	0.029	0.78	1.2	0.20	0.57	–	–	–
Attack probability $p = 0.02$											
SVDRL, hybrid	32	1.8	0.21	0	0.16	0.089	0.48	0.20	1.8	7.5	3.7
SVDRL, hybrid, $r > 6$ only	32	1.8	0.21	0	0.14	0.074	0.48	0	3.3	12	4.1
SVDRL, hybrid, superseded†	32	1.8	0.21	0	0.16	0.089	0.48	0.20	1.8	7.5	3.7

Continued on the next page

Table 16: *Resource gathering*, continued.

Errors (↓)		Death probability			CVaR		Decision		Treasure	Law	
Method	M	$\lvert\log \frac{\hat{P}}{P}\rvert$	avg.	safe	$(1,1,1)$	$(0,1,1)$	chance	CVaR	if alive	ED^2	w_2^2
SVDRL, particles	32	1.8	0.21	0	0.18	0.10	0.48	0.30	0.0057	0.75	3.6
SVDRL, sampler	32	9.1	7.8	0	0.50	0.21	0.80	0	21	80	13
SVDRL, hybrid	128	0.15	0.055	0	0.022	0.015	0.13	0.058	0.97	2.3	0.91
SVDRL, hybrid, $r > 6$ only	128	6.8	2.8	0	0.41	0.16	0.80	0	17	76	13
SVDRL, hybrid, superseded†	128	0.15	0.055	0	0.022	0.015	0.13	0.058	0.93	2.2	0.91
SVDRL, particles	128	0.15	0.055	0	0.015	0.0089	0.13	0.058	0.0062	0.60	0.91
SVDRL, sampler	128	7.9	7.4	0	0.47	0.19	0.80	0	20	87	15
SVDRL, rare event, sampler	32	7.5	7.3	0	0.51	0.22	0.80	0	21	74	12
SVDRL, rare event, EMA sampler	32	7.4	7.3	0	0.50	0.21	0.80	0	21	81	13
SVDRL, rare event, sampler	128	7.1	7.0	0	0.48	0.20	0.80	0	21	78	13
SVDRL, rare event, EMA sampler	128	7.0	6.9	0	0.47	0.19	0.80	0	20	86	15
MMD-RBF	32	8.8	5.5	0	0.71	0.39	0.75	0.81	26	52	11
MMD-RBF	128	7.9	5.8	0	0.54	0.27	0.80	0.22	23	39	8.1
MD3QN	32	9.5	6.0	0	0.68	0.39	0.78	0.74	25	53	11
MD3QN	128	4.5	1.6	0	0.43	0.24	0.79	0.26	20	44	7.6
Energy MMD	32	4.0	0.28	0	0.45	0.27	0.57	0.70	3.6	6.5	5.4
Energy MMD	128	0.74	0.28	0	0.19	0.12	0.50	0.15	4.5	4.5	2.5
Sliced Cramér	32	4.0	0.34	0	0.43	0.25	0.55	0.70	4.7	7.6	5.4
Sliced Cramér	128	0.74	0.29	0	0.17	0.11	0.48	0.14	4.4	4.7	2.3
Sliced w_2	32	4.0	0.31	0	0.38	0.23	0.57	0.70	0.0077	1.6	5.8
Sliced w_2	128	0.45	0.18	0	0.060	0.038	0.41	0.17	0.0098	0.57	2.6
Sinkhorn	32	6.0	0.61	0	0.54	0.32	0.65	0.70	10	25	7.3
QR, rare event	1000	1.9	2.1	0.27	1.3	1.6	0.67	0.60	–	–	–
QR, rare event, 156 800 steps	1000	1.9	2.2	0.35	1.8	2.3	0.70	0.97	–	–	–
QR, oracle λ, lr 0.03	1000	0.54	0.54	0.029	1.0	1.4	0.40	0.67	–	–	–
Attack probability $p = 0.005$											
SVDRL, hybrid	32	6.8	0.51	0	0.51	0.29	0.72	0.67	5.0	18	3.0
SVDRL, hybrid, $r > 6$ only	32	6.8	0.51	0	0.51	0.29	0.72	0.67	5.0	18	3.0
SVDRL, hybrid, superseded†	32	6.8	0.51	0	0.51	0.29	0.72	0.67	5.0	18	3.0
SVDRL, particles	32	6.8	0.51	0	0.57	0.34	0.72	0.73	0.0055	0.40	2.8
SVDRL, sampler	32	8.9	8.9	0	0.63	0.35	0.80	0.68	6.8	24	2.4
SVDRL, hybrid	128	1.6	0.21	0	0.18	0.10	0.48	0.12	1.7	7.8	1.2
SVDRL, hybrid, $r > 6$ only	128	1.6	0.21	0	0.18	0.10	0.48	0.12	1.8	7.8	1.2
SVDRL, hybrid, superseded†	128	1.6	0.21	0	0.18	0.10	0.48	0.12	1.7	7.3	1.1
SVDRL, particles	128	1.6	0.21	0	0.19	0.11	0.48	0.31	0.0057	0.28	0.92
SVDRL, sampler	128	8.9	8.9	0	0.56	0.30	0.80	0.14	6.1	46	4.4
SVDRL, rare event, sampler	32	8.9	8.9	0	0.59	0.33	0.80	0.49	6.1	32	3.0
SVDRL, rare event, EMA sampler	32	8.9	8.9	0	0.59	0.33	0.80	0.49	6.9	36	3.3
SVDRL, rare event, sampler	128	8.9	8.9	0	0.56	0.30	0.80	0.050	5.6	43	4.2
SVDRL, rare event, EMA sampler	128	8.9	8.9	0	0.55	0.30	0.80	0.042	6.0	47	4.7
MMD-RBF	32	8.9	8.9	0	0.68	0.40	0.80	0.78	6.8	16	3.6
MMD-RBF	128	8.9	8.9	0	0.61	0.35	0.80	0.68	5.5	18	2.8
MD3QN	32	8.9	8.9	0	0.69	0.41	0.80	0.78	6.8	21	3.8
MD3QN	128	8.9	8.9	0	0.62	0.36	0.80	0.70	5.5	24	3.1
Energy MMD	32	7.1	0.51	0	0.60	0.36	0.73	0.78	1.5	3.5	3.1
Energy MMD	128	5.3	0.75	0	0.47	0.29	0.67	0.63	3.4	7.8	1.9
Sliced Cramér	32	7.0	0.72	0	0.60	0.36	0.73	0.78	1.8	4.6	3.1
Sliced Cramér	128	5.7	0.61	0	0.47	0.28	0.68	0.63	3.4	7.7	1.9
Sliced w_2	32	6.9	0.55	0	0.57	0.34	0.72	0.78	0.0064	0.59	3.0
Sliced w_2	128	4.2	0.29	0	0.35	0.21	0.60	0.59	0.0059	0.31	1.5

Continued on the next page

Table 16: *Resource gathering*, continued.

Errors (↓)		Death probability			CVaR		Decision		Treasure	Law	
Method	M	$\lvert\log \frac{\hat{P}}{P}\rvert$	avg.	safe	$(1,1,1)$	$(0,1,1)$	chance	CVaR	if alive	ED^2	w_2^2
Sinkhorn	32	8.9	8.9	0	0.68	0.41	0.80	0.78	6.2	16	3.7
QR, rare event	1000	2.9	3.1	0.27	1.4	1.7	0.72	0.65	–	–	–
QR, rare event, 156 800 steps	1000	3.2	3.5	0.35	2.1	2.4	0.68	0.82	–	–	–
QR, oracle λ, lr 0.03	1000	1.6	1.6	0.029	1.2	1.5	0.60	0.79	–	–	–

D.5 MuJoCo

Table 17 gives the returns with their intervals, as the mean over 800k to 1M steps of Table 3 and as the mean of the last five evaluations, and Figure 9 the learning curves. Each curve point is a run's mean over its last five evaluations, because single evaluations on Hopper are bimodal for every head. The divergence guard stopped one Ant seed of SVDRL when the linear variance of its moment head reached its floor; a log-variance head cannot reach the floor, and on HalfCheetah it trains without late collapses (the rows marked *).

Table 18 compares every head's critic with Monte Carlo returns of the frozen policy beyond the centered shape of Table 3: the ratio of standard deviations, the mean error, the full w_2 and the coverage of the central 80 % interval. Every head's mean is biased by the twin-critic minimum, so only the centered metrics separate the heads, and the sampler row is a diagnostic, not a claim.

D.6 Checks of the Theory

We evaluate each bound of Section 4 on the trained networks of MDP1–3, with the ground truth only to measure the left sides (Tables 19 and 20). The checks of the consistency loss add the 3-cycle, a deterministic cycle of three states with rewards 1, -1 and 0.5 and $\gamma = 0.9$, whose return laws are point masses.

The bound of Theorem 3 holds in 80 of 80 (seed, state–action pair) combinations, but loosely: it exceeds the measured total variation 4.4 to 34 times and the measured w_1 by a median factor of 177, the gap between the two being the factor D of $w_1 \le D\,\mathrm{TV}$ on the return box. The discretization term is at least 67 % of the bound in every combination, as Remark 2 predicts for the point-mass targets that the particles produce.

Several noise levels are necessary (Remark 2). A network trained at the bottom level alone misallocates the mode masses, with errors of 0.32 on MDP2 and 0.46 on MDP3, against 0.056 and 0.099 for the multi-level network trained by DSM alone, and a single level of the order of the mode separation recovers the masses on MDP3 but not on MDP2.

Sampler bootstrapping without moment matching accumulates the smoothing of the sampler over the horizon (Remark 5), and moment matching removes the accumulation: with it, sampler bootstrapping and sampler-seeded particles match the Stein particles on MDP1–3 in $\bar{w}_1$ and in the mode masses, all differences within their intervals (Table 19).

The level-by-level bound of Corollary 3 holds at every level in 87 of 87 seeds (Table 20). On the stochastic MDP1, consistency-only training with a lagged target reaches a lower multi-level risk than DSM alone, as Lemma 5 and Remark 1 predict, and with the online network as its own target it diverges on the self-looping state–action pair, which is the lag condition after Corollary 4.

E LIMITATIONS IN DETAIL

SVDRL is not the most accurate critic on every return law, and this section collects the evidence for the limits of Section 7. The sampler fails on atomic and widely separated laws, on the atomic chain (Section D.3) and on resource gathering (Section D.4), as Lemma 2 predicts, and the hybrid rule reports the particles there. At equal size or matched compute, fitted particle clouds beat SVDRL's sampler on the unimodal MDPs and on the chain with reward noise (Table 9 and Figure 8). The measured gap $\varepsilon_{\mathrm{St}}$ is undefined for the rare-event configuration,

which keeps no particles, and on separated modes it flags a disagreement but can exceed both errors (Table 12). In SAC-SVDRL the linear variance of the moment head reached its floor and made one Ant seed diverge, a floor that a log-variance head cannot reach (Section D.5), and we make no claim about a sampler on MuJoCo.

A light rare mode far from the bulk comes out too wide, an extrapolation limit of the score network where the data are sparse. A static fit of the selected score network to fixed draws of the true root law of the rare-event chain, without bootstrapping, consistency loss or EMA, reproduces the width, with a standard deviation of the rare component 13 to 15 times the true one (Table 21).

E.1 Pixel Maze

On the maze of Zhang et al. (2021) under the pixel protocol of Debes and Tuytelaars (2026b), the competitors' benchmark from pixels, SVDRL is third of four heads in every variant (Table 22). The benchmark evaluates the uniform random policy from pixels with a PQN backbone in three variants, and every head is compared with Monte Carlo returns at the start state; SVDRL runs there without the consistency loss (Table 4). In SVDRL's first configuration nothing that trained the trunk depended on the transition, so the trunk was blind to the state and the action. Taking the moment-matching mean from each transition's bootstrap targets, as on Walker2d, Hopper and Ant, together with the log-variance moment head, removes this blindness and improves SVDRL in two of three variants, but it stays behind the sliced Cramér and multiquadric MMD critics.

E.2 Multi-Objective MuJoCo

The standardization by the moment head limits SAC-SVDRL in distributional policy evaluation of fixed policies on mo-halfcheetah and mo-hopper (Table 23). With its selected configuration, SVDRL's particles lie between MD3QN and energy MMD in energy distance after 10^6 steps on both environments, and behind both after $5 \cdot 10^5$ steps. The remaining gap is under-dispersion: the EMA target cloud is narrower than the online one, which the bootstrap compounds by up to the factor $1/(1-\gamma^2)$, about 50 at $\gamma = 0.99$. The sampler of the selected configuration fails its ground-truth-free gate on three of four policies.

A float32 cancellation in PopArt's variance accumulator affected the low-coefficient-of-variation targets of the first multi-objective runs. We reran the HalfCheetah baselines with a centered accumulator, and these rows, together with SVDRL's selected configuration, which uses it, are the comparison there; the uncentered HalfCheetah rows are kept for reference. A check of the logged PopArt statistics of every other reported run finds no comparable effect.

Table 5: *Rare-event queries.* Rare-event chain, $p = 0.005$, $D = 2$, $M = 128$: IQM over the ten final seeds with its 95 % interval below; bold: inside the interval of the best IQM. Queries: CVaR of the first coordinate and of the equal-weight scalarization, the catastrophe's probability, the conditional means of Z_2 given $Z_1 \leq \tau$ and at the rare mode, the entropic risk, and the accuracies of the chance-constrained decision and of the decision by expected scalarized return (ESR). SVDRL's sampler answers with $N = 10^5$ draws, and Langevin dynamics on the partial score answers the conditional mean at the rare mode (quadrature of the energy: 1.7 for particle bootstrapping); *: the reported configuration. Particle critics answer from their atoms or from a kernel density estimate (KDE) on them.

	CVaR error (↓)			Event (↓)	Conditional mean error (↓)		Entropic (↓)	Decision accuracy (↑)	
Estimator	$\alpha = 0.001$	$\alpha = 0.01$	$\alpha = 0.001$, $(\frac{1}{2}, \frac{1}{2})$	$\lvert \log \hat{P}/P \rvert$	$Z_1 \leq \tau$	$Z_1 = z_{\text{rare}}$	$\beta = 0.1$	chance	ESR
SVDRL sampler, bootstrapping from									
its draws, EDM*	**6.9** [5.8, 8]	9.7 [6.3, 15]	**2.9** [2.4, 3.4]	**0.16** [0.11, 0.28]	2.6 [2.3, 2.8]	2.9 [2.5, 3.6]	13 [11, 15]	**1** [0.99, 1]	0.57 [0.57, 0.57]
its draws	9.2 [6.4, 12]	8 [5.7, 13]	4.8 [3.2, 5.9]	0.32 [0.19, 0.49]	**1.6** [1.3, 1.8]	–	**7.7** [5.6, 9.5]	0.95 [0.8, 1]	0.28 [0.28, 0.28]
its draws, moment-matched	11 [8, 16]	18 [11, 27]	4.8 [3.6, 7]	0.48 [0.28, 0.87]	2.6 [2, 3.4]	–	17 [10, 25]	0.89 [0.62, 1]	0.27 [0.27, 0.27]
its draws, running moments	36 [34, 38]	**3.7** [1.9, 5.5]	7.1 [6.1, 8.2]	0.29 [0.17, 0.53]	22 [21, 22]	–	35 [34, 36]	0.94 [0.75, 1]	**0.88** [0.88, 0.88]
particles	15 [11, 17]	16 [10, 22]	6.6 [4.8, 7.4]	0.45 [0.26, 0.78]	**1.8** [1.4, 2.5]	**1.6** [1.2, 2]	14 [10, 17]	0.91 [0.72, 1]	0.35 [0.35, 0.35]
particles, resample-move	31 [16, 56]	32 [27, 38]	13 [6.6, 23]	1.6 [1.1, 3.1]	2.7 [2.2, 6.1]	–	27 [20, 41]	0.64 [0.04, 1]	0 [0, 0]
Particles and kernel density estimates									
SVDRL Stein particles	83 [71, 94]	52 [49, 54]	35 [30, 39]	6.5 [5.6, 7.3]	13 [11, 15]	–	57 [50, 63]	0.56 [0.28, 0.92]	0.28 [0.28, 0.28]
MMD-RBF, atoms	120 [120, 120]	56 [54, 57]	51 [50, 51]	8.5 [8.5, 8.5]	18 [18, 18]	–	69 [69, 70]	0.5 [0, 1]	0 [0, 0]
MMD-RBF, KDE	110 [110, 110]	49 [47, 51]	46 [45, 47]	8.5 [8.5, 8.5]	33 [29, 38]	60 [52, 67]	69 [68, 69]	0.5 [0, 1]	0 [0, 0]
MD3QN, atoms	100 [97, 110]	41 [37, 44]	42 [39, 44]	7.9 [7.2, 8.4]	18 [18, 18]	–	66 [64, 67]	0.54 [0.09, 1]	0.09 [0.09, 0.09]
MD3QN, KDE	98 [92, 100]	40 [35, 44]	40 [36, 42]	6.7 [6.2, 7.5]	18 [15, 19]	20 [18, 22]	66 [63, 67]	0.56 [0.1, 1]	0.1 [0.1, 0.1]
Energy MMD, atoms	71 [58, 80]	53 [50, 56]	30 [25, 34]	4.8 [4, 5.4]	11 [9.3, 12]	–	44 [36, 50]	0.71 [0.48, 1]	0.48 [0.48, 0.48]
Energy MMD, KDE	72 [60, 82]	54 [51, 57]	31 [26, 36]	4.9 [4.1, 5.4]	11 [8.6, 14]	11 [8.5, 15]	44 [37, 50]	0.7 [0.48, 1]	0.48 [0.48, 0.48]
Sliced Cramér, atoms	70 [58, 80]	53 [50, 56]	30 [25, 34]	4.8 [4, 5.4]	11 [9.3, 12]	–	43 [36, 49]	0.71 [0.48, 0.99]	0.48 [0.48, 0.48]
Sliced Cramér, KDE	72 [60, 82]	54 [51, 56]	31 [26, 35]	4.8 [4, 5.4]	11 [9.2, 13]	11 [8.8, 14]	44 [37, 50]	0.71 [0.48, 0.99]	0.48 [0.48, 0.48]
Sliced w_2, atoms	67 [55, 76]	53 [50, 56]	29 [24, 32]	4.8 [4, 5.4]	9.8 [8.1, 11]	–	42 [35, 47]	0.7 [0.48, 1]	0.48 [0.48, 0.48]
Sliced w_2, KDE	71 [59, 79]	55 [52, 58]	30 [25, 34]	4.8 [4, 5.4]	9.9 [8.4, 11]	9.9 [8.1, 12]	43 [36, 48]	0.7 [0.48, 1]	0.48 [0.48, 0.48]
Sinkhorn, atoms	88 [80, 98]	47 [39, 53]	38 [35, 43]	5.6 [4.9, 6.5]	12 [12, 14]	–	53 [50, 59]	0.59 [0.38, 0.89]	0.38 [0.38, 0.38]
Sinkhorn, KDE	84 [77, 93]	47 [40, 52]	36 [33, 40]	5.5 [4.8, 6.5]	17 [14, 22]	30 [24, 38]	53 [50, 58]	0.6 [0.38, 0.95]	0.38 [0.38, 0.38]
Quantile critic									
oracle λ, 1000 quantiles	62 [59, 64]	10 [6.6, 13]	–	1.8 [0.49, 4.6]	–	–	59 [56, 61]	0.71 [0.26, 1]	0.02 [0.02, 0.02]

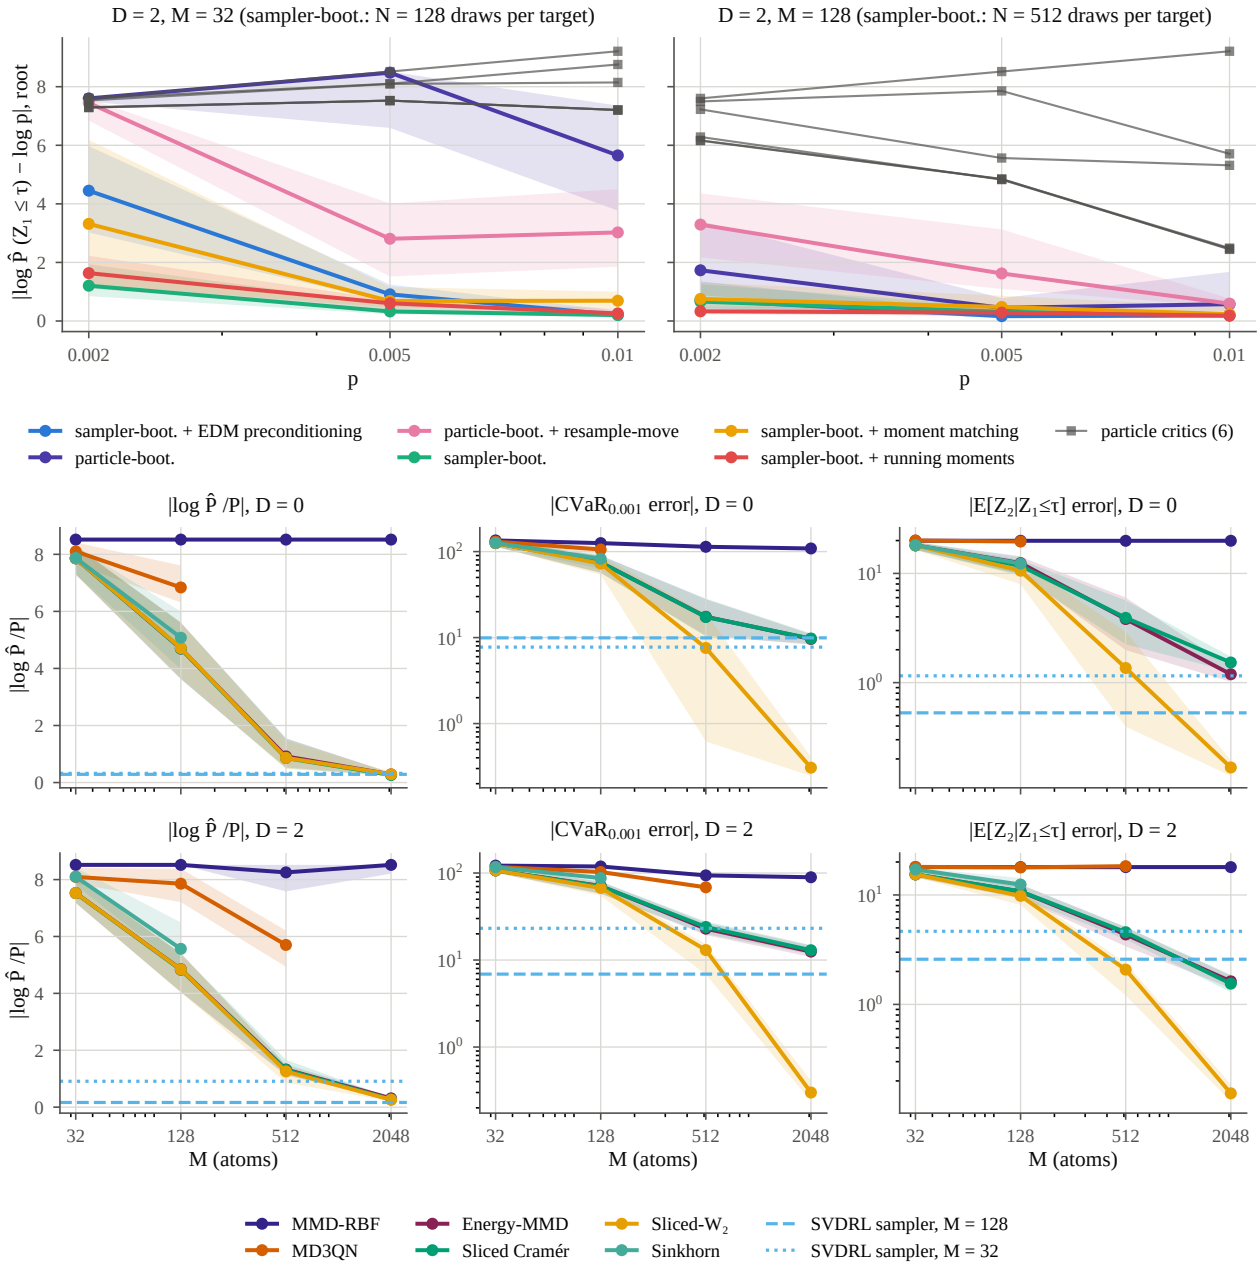

Figure 4: *Propagation and particle count.* Rare-event chain, final seeds. (a) Error of the catastrophe's probability at the root against p for $D = 2$, with $M = 32$ and $M = 128$, for SVDRL's variants and the six particle critics: with sampler bootstrapping in the reported configuration (moment-matched, EDM preconditioning) the catastrophe survives the recursion for $p \geq 0.005$, while particle bootstrapping and the particle critics lose it. (b) Errors of the particle critics against their particle count M at $D = 0$ (top) and $D = 2$ (bottom), with SVDRL's sampler at $M = 128$ and $M = 32$ as horizontal lines: sliced w_2 needs about $10/p = 2048$ particles to resolve the location and the conditional tail of the catastrophe.

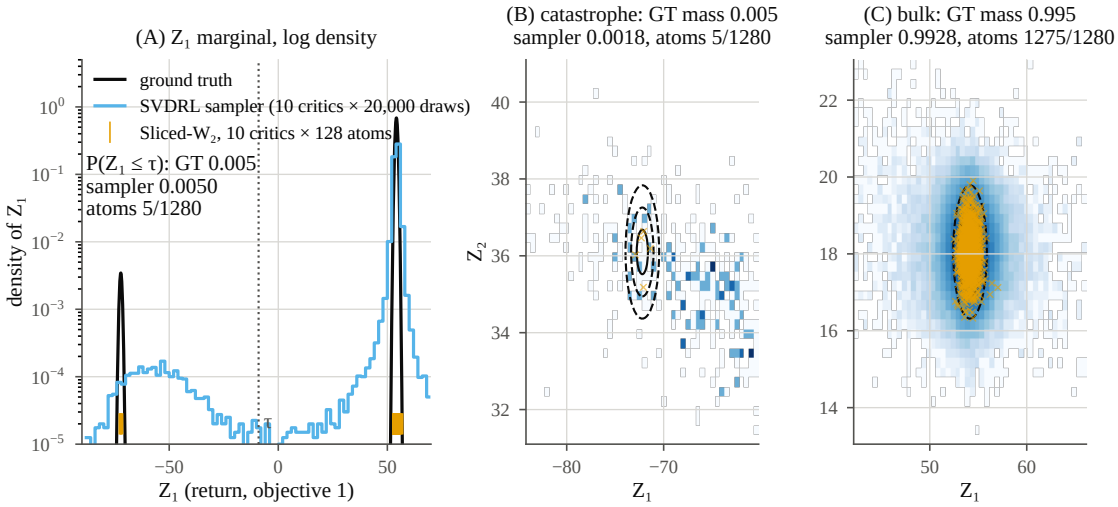

Figure 5: *Return law of the risky root action.* In the rare-event chain, a catastrophe of mass $p = 0.005 < 1/M$ lies two steps below the root. (A) Log density of the first coordinate: the true law, the sampler law of ten SVDRL critics and the atoms of ten sliced w_2 critics ($M = 128$). (B) The catastrophe and (C) the bulk. The sampler gets the probability of the catastrophe right but spreads the rare mode too wide; particles can only hold whole atoms of mass $1/M$ there.

Table 6: *The objects behind the hybrid.* Synthetic MDPs, final seeds: every SVDRL object behind Table 2, per MDP and error (mass: total variation of the mode masses; shape: within-mode w_2^2), in units of 10^{-3}. The hybrid, the sampler ($N = 4096$), the particles ($M = 64d$), the object chosen on the tuning seeds, and the hybrid with the separation test $r > 6$ alone, each with its ratio to the best competitor of all rows, including those on our target set; bold: a win, underlined: a tie by Welch's test; –: the competitor's error is zero. The hybrid reports the particles where their separation ratio exceeds 6 after merging particles closer than the smallest noise level, or where a cluster of at least two particles is narrower than it, and the sampler draws otherwise; the last column is the object it reports. Wins/ties/losses over the MDPs against the competitors as specified (own) and including those on our target set (all).

			Hybrid		Sampler		Particles		Tuning choice		Test $r > 6$ only		
MDP	d	Error (↓)	10^{-3}	ratio	10^{-3}	ratio	10^{-3}	ratio	10^{-3}	ratio	10^{-3}	ratio	Reports
MDP1	1	shape	5.7	**0.27**	5.7	**0.27**	64	3.1	5.7	**0.27**	5.7	**0.27**	sampler
MDP2	1	mass	4.7	1.5	7.7	2.5	4.8	1.5	7.7	2.5	4.7	1.5	sampler
		shape	81	1.6	13	**0.25**	83	1.7	13	**0.25**	81	1.6	sampler
MDP3	2	mass	2.3	1.5	12	8.0	2.3	1.5	12	8.0	2.3	1.5	sampler
		shape	389	**0.80**	125	**0.26**	389	**0.80**	125	**0.26**	389	**0.80**	sampler
MDP4	5	mass	0.065	–	13	–	0	1.0	0	1.0	0.065	–	particles
		shape	62	**0.097**	808	1.3	55	**0.086**	55	**0.086**	62	**0.097**	particles
MDP5	2	shape	6.1	**0.20**	6.1	**0.20**	37	1.2	6.1	**0.20**	6.1	**0.20**	sampler
MDP6	2	mass	0.060	–	11	–	0	1.0	0	1.0	0.060	–	particles
		shape	50	**0.12**	154	**0.37**	73	**0.17**	73	**0.17**	50	**0.12**	particles
MDP7	3	shape	1.5	**0.27**	1.5	**0.27**	5.5	**0.97**	1.5	**0.27**	1.5	**0.27**	sampler
MDP8	2	shape	38	**0.23**	38	**0.23**	160	0.98	38	**0.23**	38	**0.23**	sampler
Wins/ties/losses, selected configuration													
mass, own			0/2/2		0/0/4		0/2/2		0/2/2		0/2/2		
mass, all			0/2/2		0/0/4		0/2/2		0/2/2		0/2/2		
shape, own			8/0/0		8/0/0		6/1/1		8/0/0		8/0/0		
shape, all			7/0/1		7/1/0		4/1/3		8/0/0		7/0/1		
Wins/ties/losses, single configuration													
mass, own			0/2/2		0/0/4		0/2/2		0/2/2		0/2/2		
mass, all			0/2/2		0/0/4		0/2/2		0/2/2		0/2/2		
shape, own			8/0/0		8/0/0		5/1/2		8/0/0		8/0/0		
shape, all			7/0/1		7/1/0		4/0/4		8/0/0		7/0/1		

Table 7: *Mode masses on the synthetic MDPs.* Total variation of the mode masses on the four MDPs with separated modes, in units of 10^{-3}, averaged over the last 20 % of training; floor: the $1/M$ limit of the mass error. Ratios to the best competitor as in Table 2. Underlined: a tie by Welch's test; –: the competitor's error is zero.

		Mode masses, total variation (↓)			
				SVDRL / best	
MDP	d	SVDRL	floor	own	all
MDP2	1	4.7	3.1	1.3	1.5
MDP3	2	2.3	1.6	1.4	1.5
MDP4	5	0.065	0	0.11	–
MDP6	2	0.060	0	–	–
Wins/ties/losses				0/2/2	0/2/2
its particles alone				0/2/2	0/2/2

Table 8: *Best competitors.* The competitor behind each ratio of Tables 2 and 7, per MDP, error and comparison, with its error in units of 10^{-3}. Own: the competitors as specified; all: also the competitors given our target set; $M = 4096$: the compute-matched comparison on the unimodal MDPs. †: an objective of Debes and Tuytelaars (2026b), implemented in our agent; base: the learning rate of the competitor's base configuration (Table 4).

MDP	Error	Comparison	Best competitor	Target set	Learning rate	Its error (10^{-3}) (↓)
MDP1	shape	own	multiquadric MMD†	own	base	24.3
		all	multiquadric MMD†	ours	base	20.8
		M=4096	multiquadric MMD†	ours	base	0.563
MDP2	shape	own	sliced w_2	own	base	98.4
		all	sliced Cramér†	ours	$3 \cdot 10^{-2}$	50.1
	mass	own	multiquadric MMD†	own	$3 \cdot 10^{-2}$	3.73
		all	sliced Cramér†	ours	$3 \cdot 10^{-2}$	3.13
MDP3	shape	own	sliced Cramér†	own	$3 \cdot 10^{-2}$	1950
		all	sliced Cramér†	ours	$3 \cdot 10^{-2}$	488
	mass	own	sliced Cramér†	own	$3 \cdot 10^{-2}$	1.68
		all	sliced Cramér†	ours	$3 \cdot 10^{-2}$	1.56
MDP4	shape	own	multiquadric MMD†	own	$3 \cdot 10^{-2}$	3080
		all	multiquadric MMD†	ours	$3 \cdot 10^{-2}$	641
	mass	own	multiquadric MMD†	own	$3 \cdot 10^{-2}$	0.597
		all	multiquadric MMD†	ours	$3 \cdot 10^{-2}$	0
MDP5	shape	own	sliced w_2†	own	10^{-1}	36.7
		all	sliced w_2†	ours	10^{-1}	30.0
		M=4096	sliced w_2†	ours	base	5.17
MDP6	shape	own	multiquadric MMD†	own	$3 \cdot 10^{-2}$	925
		all	multiquadric MMD†	ours	$3 \cdot 10^{-2}$	419
	mass	own	multiquadric MMD†	own	$3 \cdot 10^{-2}$	0
		all	multiquadric MMD†	own	$3 \cdot 10^{-2}$	0
MDP7	shape	own	sliced w_2	own	base	8.17
		all	sliced w_2	ours	base	5.60
		M=4096	sliced w_2†	ours	base	0.375
MDP8	shape	own	sliced w_2†	own	10^{-1}	205
		all	sliced w_2†	ours	10^{-1}	163
		M=4096	sliced w_2†	ours	base	16.9

Table 9: *Compute-matched comparison.* Within-mode shape on the unimodal MDPs: the hybrid's sampler (4096 draws) against the best competitor on our target set and against SVDRL's transport-free particles ($\alpha = 0$), at $M \in \{64d, 512, 4096\}$; †: an objective of Debes and Tuytelaars (2026b), implemented in our agent. Error mean ± standard error over the final seeds, and the hybrid's ratio to each row with its Welch t (negative: the hybrid is better); bold: the lowest mean error of each MDP.

MDP	row	M	error (mean ± SE) (↓)	hybrid ÷ row (Welch t)
MDP1	SVDRL hybrid (sampler)	4096	0.00559±0.00052	
MDP1	multiquadric MMD†	64	0.0208±0.00056	0.269 (-19.9)
MDP1	multiquadric MMD†	512	0.00286±0.000070	1.95 (+5.2)
MDP1	multiquadric MMD†	4096	**0.000563**±0.000061	9.94 (+9.7)
MDP1	SVDRL particles, $\alpha = 0$	512	0.0118±0.00028	0.474 (-10.5)
MDP1	SVDRL particles, $\alpha = 0$	4096	0.00157±0.000069	3.56 (+7.7)
MDP5	SVDRL hybrid (sampler)	4096	0.00609±0.000068	
MDP5	sliced w_2†	128	0.0300±0.00022	0.203 (-103.5)
MDP5	sliced w_2†	512	0.0112±0.00027	0.544 (-18.2)
MDP5	sliced w_2†	4096	0.00517±0.00024	1.18 (+3.6)
MDP5	SVDRL particles, $\alpha = 0$	512	0.0166±0.00011	0.366 (-82.3)
MDP5	SVDRL particles, $\alpha = 0$	4096	**0.00489**±0.000022	1.24 (+16.6)
MDP7	SVDRL hybrid (sampler)	4096	0.00155±0.000030	
MDP7	sliced w_2†	192	0.00730±0.00010	0.213 (-53.6)
MDP7	sliced w_2†	512	0.00284±0.000041	0.548 (-25.2)
MDP7	sliced w_2†	4096	0.000375±0.0000059	4.14 (+38.1)
MDP7	SVDRL particles, $\alpha = 0$	512	0.00219±0.000012	0.709 (-19.6)
MDP7	SVDRL particles, $\alpha = 0$	4096	**0.000305**±0.0000055	5.10 (+40.5)
MDP8	SVDRL hybrid (sampler)	4096	0.0378±0.0011	
MDP8	sliced w_2†	128	0.163±0.0031	0.232 (-37.8)
MDP8	sliced w_2†	512	0.0469±0.00061	0.806 (-7.1)
MDP8	sliced w_2†	4096	**0.0169**±0.00028	2.24 (+18.2)
MDP8	SVDRL particles, $\alpha = 0$	512	0.0787±0.0015	0.481 (-21.4)
MDP8	SVDRL particles, $\alpha = 0$	4096	0.0228±0.00031	1.66 (+12.9)

Table 10: *Mass and shape together.* Full w_2^2 per MDP: SVDRL's selected and single configuration against the best competitor on its own target set and among every row, with the selected configuration's ratio and the Welch decision; †: an objective of Debes and Tuytelaars (2026b), implemented in our agent, with its learning rate where it is not the base one; *: on our target set.

	SVDRL		best competitor, own target set		best competitor, every row	
MDP	selected (↓)	single (↓)	error (row) (↓)	selected ÷	error (row) (↓)	selected ÷
MDP1 (d=1)	0.0635	0.0768	0.0243 (multiquadric MMD†)	×2.6 L	0.0208 (multiquadric MMD†*)	×3.1 L
MDP2 (d=1)	1.05	1.05	0.832 (multiquadric MMD†, $3 \cdot 10^{-2}$)	×1.3 L	0.504 (sliced w_2†*, $3 \cdot 10^{-2}$)	×2.1 L
MDP3 (d=2)	12.9	12.9	11.8 (sliced Cramér†, $3 \cdot 10^{-2}$)	×1.1 T	11.8 (sliced Cramér†, $3 \cdot 10^{-2}$)	×1.1 T
MDP4 (d=5)	0.0526	0.0526	3.72 (multiquadric MMD†, $3 \cdot 10^{-2}$)	×0.014 W	0.613 (multiquadric MMD†*, $3 \cdot 10^{-2}$)	×0.086 W
MDP5 (d=2)	0.0368	0.0459	0.0367 (sliced w_2†, 0.1)	×1.0 T	0.0300 (sliced w_2†*, 0.1)	×1.2 L
MDP6 (d=2)	9.68	9.68	9.77 (multiquadric MMD†, $3 \cdot 10^{-2}$)	×0.99 T	9.73 (multiquadric MMD†*, $3 \cdot 10^{-2}$)	×0.99 T
MDP7 (d=3)	0.00545	0.00545	0.00817 (sliced w_2)	×0.67 W	0.00560 (sliced w_2*)	×0.97 W
MDP8 (d=2)	0.171	0.219	0.205 (sliced w_2†, 0.1)	×0.83 W	0.163 (sliced w_2†*, 0.1)	×1.1 T

W/T/L, selected configuration per MDP: 3/3/2 (own target set) / 2/3/3 (every row)
W/T/L, single configuration: 2/3/3 (own target set) / 2/2/4 (every row)

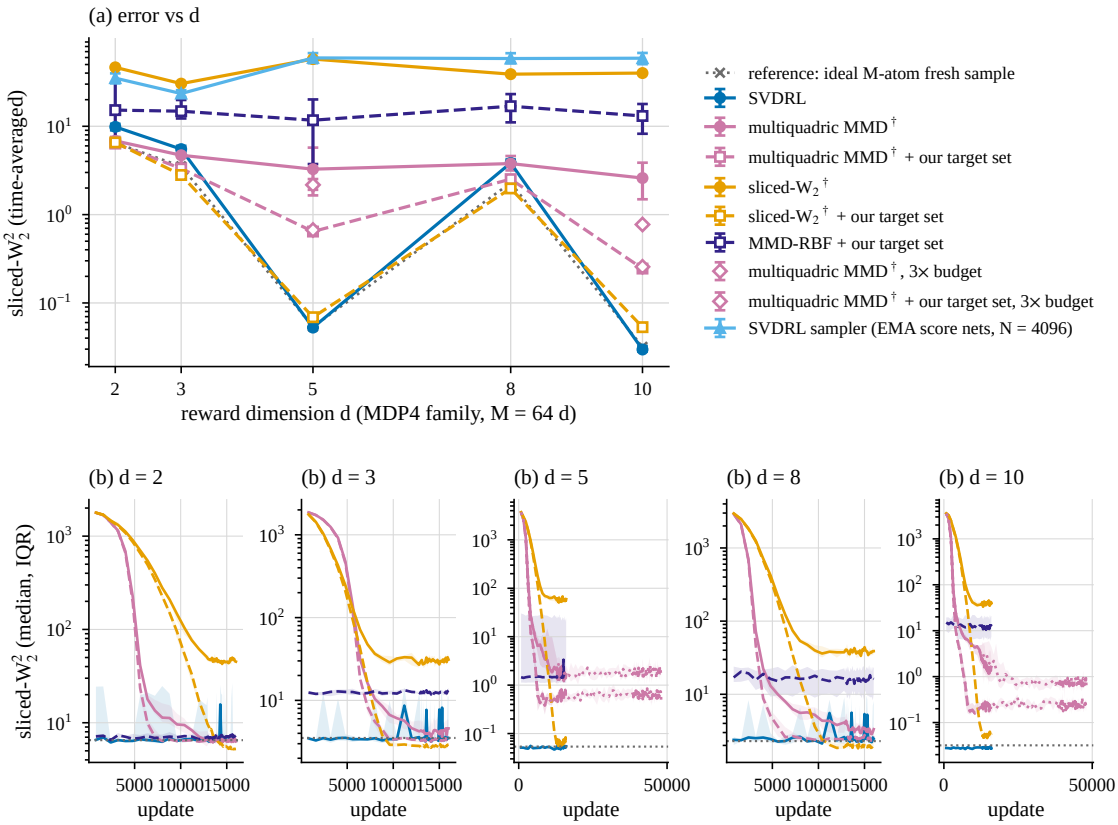

Figure 6: *Scaling with the reward dimension.* MDP4 family across the reward dimension d, $M = 64d$, ten final seeds. (a) Time-averaged sliced w_2^2 against d for SVDRL and the competitors on their own and on our target set. (b) Learning curves per d, median and interquartile range. The full error is small where the mode masses are multiples of $1/M$, at $d = 5$ and $d = 10$, and dominated by mass rounding elsewhere. †: an objective of Debes and Tuytelaars (2026b), implemented in our agent; 3× budget: three times the training steps.

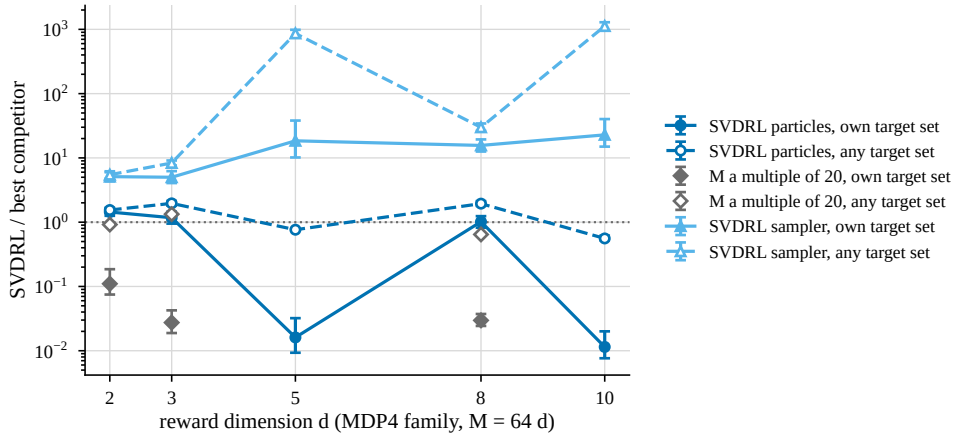

Figure 7: *SVDRL against the best competitor across the reward dimension.* MDP4 family, $M = 64d$, ten final seeds: SVDRL's error over the best competitor's (ratio of means with its 95 % interval; below one, SVDRL is better), for its particles and its sampler, against the competitors on their own target set and on any target set. M a multiple of 20: the allocation control, at which every mode mass is a multiple of $1/M$.

Table 11: *Relaxation and target set.* Ablation of the relaxation α of the Stein move and the stratified target set on the synthetic MDPs: within-mode shape error (every MDP; the full error where every evaluated state–action pair is unimodal) and total variation of the mode masses (mass TV; MDPs with separated modes). Fully annealed: SVDRL with fully annealed transport; + strat.: the stratified target set (systematic successors + Latin-hypercube rewards); + full-cov. matching: full-covariance moment matching; $\rho = 1$: the whole cloud replaced by the stratified target set, with relaxation α (no running-mean moment matching). IQM over the ten final seeds; †: run on the five tuning seeds only, IQM over those; –: not run. Time-averaged over the last 20 % of the steps.

					$\rho = 1$		
MDP	block	fully annealed	fully annealed + strat.	+ full-cov. matching	$\alpha = 10^{-3}$	$\alpha = 0.01$	$\alpha = 0.1$
MDP1	shape (↓)	0.127	0.109†	–	0.0770	–	0.0635
MDP2	shape (↓)	2.52	–	–	0.0815	–	0.327†
	mass TV (↓)	0.0374	–	–	0.00487	–	0.00496†
MDP3	shape (↓)	20.0	–	–	0.387	–	2.94†
	mass TV (↓)	0.0550	–	–	0.00244	–	0.00251†
MDP4 (d=5)	shape (↓)	59.6	–	–	0.0549	0.0575†	1.20†
	mass TV (↓)	0.0420	–	–	0	0†	0†
MDP5	shape (↓)	0.0437	0.0369	–	0.0460	–	0.115†
MDP6	shape (↓)	16.4	–	–	0.0727	–	0.461
	mass TV (↓)	0.0440	–	–	0	–	0
MDP7	shape (↓)	0.0241	0.0213†	0.00631†	0.00545	–	0.0380†
MDP8	shape (↓)	1.23	1.16†	0.171	0.217	–	0.263†

Table 12: *Ground-truth-free gap.* Measured Stein-particle gap $\varepsilon_{\mathrm{St}}$: the distance between SVDRL's particle cloud (M=64d) and its own reverse-diffusion sampler (N=4096 draws, Euler on 512 geometric levels of the state–action pair's EMA score net), computed without ground truth; next to it, the ground-truth error of each object in the same distance. d=1: exact W_1 and W_2; d=2: exact W_2 (GT side: 4096 draws); d>2: sliced W_2 (quantile estimator). Per seed: W_2 per evaluated state–action pair, mean over pairs, then over the tail evaluations (last 20 % of the steps; every 4th at d=2); IQM over the ten final seeds [bootstrap 95 % CI]. Configurations: the selected configuration of each MDP and the single configuration ($\rho = 1$, $\alpha = 10^{-3}$).

MDP	d	config	M	$\varepsilon_{\mathrm{St}}$ (W_2)	$\varepsilon_{\mathrm{St}}$ (W_1)	particles–GT (↓)	sampler–GT (↓)
MDP1	1	ρ=1, α=0.1 (selected)	64	0.240$_{[0.23,\,0.25]}$	0.168$_{[0.16,\,0.18]}$	0.248$_{[0.24,\,0.26]}$	0.0692$_{[0.065,\,0.075]}$
		ρ=1, α=10^{-3} (single)	64	0.264$_{[0.25,\,0.27]}$	0.189$_{[0.18,\,0.20]}$	0.271$_{[0.26,\,0.28]}$	0.0695$_{[0.065,\,0.075]}$
MDP2	1	ρ=1, α=10^{-3} (selected = single)	64	1.26$_{[1.2,\,1.4]}$	0.342$_{[0.32,\,0.37]}$	0.983$_{[0.91,\,1.0]}$	1.08$_{[1.0,\,1.2]}$
MDP3	2	ρ=1, α=10^{-3} (selected = single)	128	4.12$_{[3.8,\,4.9]}$	–	3.32$_{[3.0,\,3.7]}$	4.44$_{[3.8,\,5.1]}$
MDP4 (d=5)	5	ρ=1, α=10^{-3} (selected = single)	320	7.95$_{[7.4,\,9.0]}$	–	0.229$_{[0.23,\,0.23]}$	7.95$_{[7.4,\,8.9]}$
MDP5	2	fully annealed + strat. (selected)	128	0.196$_{[0.19,\,0.20]}$	–	0.191$_{[0.19,\,0.19]}$	0.0774$_{[0.077,\,0.079]}$
		ρ=1, α=10^{-3} (single)	128	0.217$_{[0.22,\,0.22]}$	–	0.214$_{[0.21,\,0.22]}$	0.0777$_{[0.077,\,0.079]}$
MDP6	2	ρ=1, α=10^{-3} (selected = single)	128	2.41$_{[2.1,\,2.7]}$	–	1.77$_{[1.6,\,2.0]}$	2.44$_{[2.2,\,2.7]}$
MDP7	3	ρ=1, α=10^{-3} (selected = single)	192	0.0817$_{[0.081,\,0.083]}$	–	0.0733$_{[0.073,\,0.074]}$	0.0377$_{[0.037,\,0.039]}$
MDP8	2	+ full-cov. matching (selected)	128	0.374$_{[0.37,\,0.38]}$	–	0.377$_{[0.37,\,0.38]}$	0.179$_{[0.18,\,0.19]}$
		ρ=1, α=10^{-3} (single)	128	0.443$_{[0.43,\,0.45]}$	–	0.436$_{[0.43,\,0.45]}$	0.177$_{[0.17,\,0.18]}$

Table 13: *Stein move ablation.* The Stein move against other moves on MDP1, MDP5 and MDP8, whose selected configurations use it: Langevin dynamics along the same score, a Gaussian jitter of matched size, the Stein move without repulsion and no move. Within-mode shape error (mean ± standard error over the final seeds), its ratio to the Stein move with the Welch t, the fraction of distinct particles, and the nearest-neighbor distance over the cloud's standard deviation.

MDP	move	shape error (↓)	÷ Stein	Welch t	distinct	NN / std
MDP1	Stein (selected)	$0.0638_{\pm 0.0026}$	–	–	1.0	0.022
	Langevin	$0.0762_{\pm 0.0034}$	×1.19	+2.9	1.0	0.021
	jitter	$0.0768_{\pm 0.0031}$	×1.20	+3.2	1.0	0.021
	no repulsion	$0.0771_{\pm 0.0031}$	×1.21	+3.2	1.0	0.021
	no move	$0.0771_{\pm 0.0032}$	×1.21	+3.2	1.0	0.021
MDP5	Stein (selected)	$0.0374_{\pm 0.0004}$	–	–	1.0	0.17
	Langevin	$0.0458_{\pm 0.0005}$	×1.23	+13.2	1.0	0.16
	jitter	$0.0453_{\pm 0.0004}$	×1.21	+13.9	1.0	0.16
	no repulsion	$1.07_{\pm 0.01}$	×28.6	+85.7	0.030	0.0
	no move	$0.0452_{\pm 0.0008}$	×1.21	+8.8	1.0	0.16
MDP8	Stein (selected)	$0.164_{\pm 0.002}$	–	–	1.0	0.13
	Langevin	$0.215_{\pm 0.004}$	×1.31	+11.6	1.0	0.12
	jitter	$0.228_{\pm 0.004}$	×1.39	+13.2	1.0	0.12
	no repulsion	$33.5_{\pm 1.6}$	×204	+20.6	0.035	0.0
	no move	$0.224_{\pm 0.008}$	×1.36	+6.8	1.0	0.12

Table 14: *Consistency ablation.* Full w_2^2 of SVDRL with the fully annealed transport, with ($\lambda_{\text{cons}} = 1$) and without ($\lambda_{\text{cons}} = 0$) the consistency loss, mean over ten seeds, and the paired difference with its 95 % Student-t interval; exact w_2^2 for $d \leq 2$ and CDF-based sliced w_2^2 for MDP4 and MDP7. A negative difference favors the consistency loss.

MDP	$\lambda_{\text{cons}} = 1$ (↓)	$\lambda_{\text{cons}} = 0$ (↓)	difference [95 % interval]
MDP1	0.136	0.162	−0.027 [−0.053, −0.0009]
MDP2	10.1	9.4	0.67 [−1.65, 2.99]
MDP3	67.4	101.2	−33.8 [−57.9, −9.6]
MDP4	230.8	270.5	−39.7 [−70.4, −9.0]
MDP5	0.0421	0.0428	−0.0006 [−0.0036, 0.0023]
MDP6	71.8	103.5	−31.7 [−102.8, 39.3]
MDP7	0.0241	0.0242	−0.0001 [−0.0003, 0.0001]
MDP8	1.21	2.12	−0.91 [−1.05, −0.78]

Table 15: *Chain with reward noise.* The chain of Debes and Tuytelaars (2026b) with $K = 10$ reward dimensions and reward noise σ_r, near-exact TD, $M = 128$ particles, final seeds: w_2 at the start state (exact assignment; not defined for 4096 draws against the 128-draw reference), squared energy distance (ED^2) in units of 10^{-3} and mean error in units of 10^{-2}; bold: the lowest error of each column among the methods. SVDRL's particles, its moment-matched sampler with 4096 and 128 draws and its fully annealed transport; the six variants of Debes and Tuytelaars (2026b) on their own target sets and on ours, and i.i.d. samples of the true law as references. All methods have a higher w_2 here than Debes and Tuytelaars (2026b) publish at $\sigma_r = 0$, because they use $M = 500$: with 128 particles against 128 reference draws, even an i.i.d. sample of the true law lies above the published values.

Errors (↓)	$\sigma_r = 0$			$\sigma_r = 0.05$			$\sigma_r = 0.2$			$\sigma_r = 0.5$		
Method	w_2	ED^2	mean	w_2	ED^2	mean	w_2	ED^2	mean	w_2	ED^2	mean
SVDRL particles	0.482	19.0	4.31	0.572	6.79	5.44	1.35	12.7	9.91	2.96	25.9	19.2
SVDRL sampler, $N = 4096$	–	12.5	**0.413**	–	**1.05**	**0.497**	–	**0.778**	**1.03**	–	**1.27**	**2.10**
SVDRL sampler, $N = 128$	0.520	22.1	8.39	0.600	11.5	8.78	1.37	18.7	14.6	3.00	38.8	30.5
SVDRL, fully annealed	0.661	77.7	15.3	0.775	70.1	18.9	1.45	48.2	25.5	2.95	75.0	53.7
Variants of Debes and Tuytelaars (2026b), as specified												
sliced w_2	**0.438**	14.7	2.20	**0.521**	5.27	2.49	1.26	6.09	3.34	2.80	11.9	4.73
sliced Cramér	0.452	5.18	2.25	0.530	2.58	2.68	1.27	5.64	3.51	2.81	11.8	4.87
MMD, multiquadric	0.452	5.13	2.02	0.530	2.60	2.80	1.27	5.65	3.68	2.81	11.8	4.97
MMD, inv. multiquadric	0.452	5.26	2.20	0.530	2.62	2.78	1.27	5.66	3.67	2.81	11.8	4.95
max-sliced w_2	0.456	21.2	3.68	0.539	7.56	3.96	**1.26**	7.66	4.64	**2.78**	14.0	5.75
max-sliced Cramér	0.468	10.6	3.02	0.557	3.72	3.62	1.28	6.78	4.78	2.80	13.8	6.69
The same variants, on our target set												
sliced w_2	0.442	4.73	1.17	0.525	2.40	1.60	1.27	5.13	1.82	2.83	10.6	2.49
sliced Cramér	0.454	4.66	1.72	0.530	2.33	2.16	1.26	5.25	2.14	2.81	11.4	2.87
MMD, multiquadric	0.452	2.77	1.62	0.530	2.40	2.43	1.26	5.29	2.64	2.82	11.1	3.19
MMD, inv. multiquadric	0.452	**2.75**	1.60	0.529	2.40	2.44	1.26	5.33	2.67	2.82	11.2	3.27
max-sliced w_2	0.446	6.96	1.59	0.534	2.62	1.70	1.27	5.81	2.19	2.80	12.4	2.48
max-sliced Cramér	0.460	9.59	2.25	0.551	2.59	1.94	1.28	5.87	2.38	2.80	12.9	3.53
i.i.d. sample, $M = 128$	0.536	12.9	9.71	0.599	9.64	8.78	1.36	19.2	15.4	2.97	35.9	28.0
i.i.d. sample, 4096	–	0.298	1.63	–	0.289	1.53	–	0.584	2.69	–	1.36	5.62

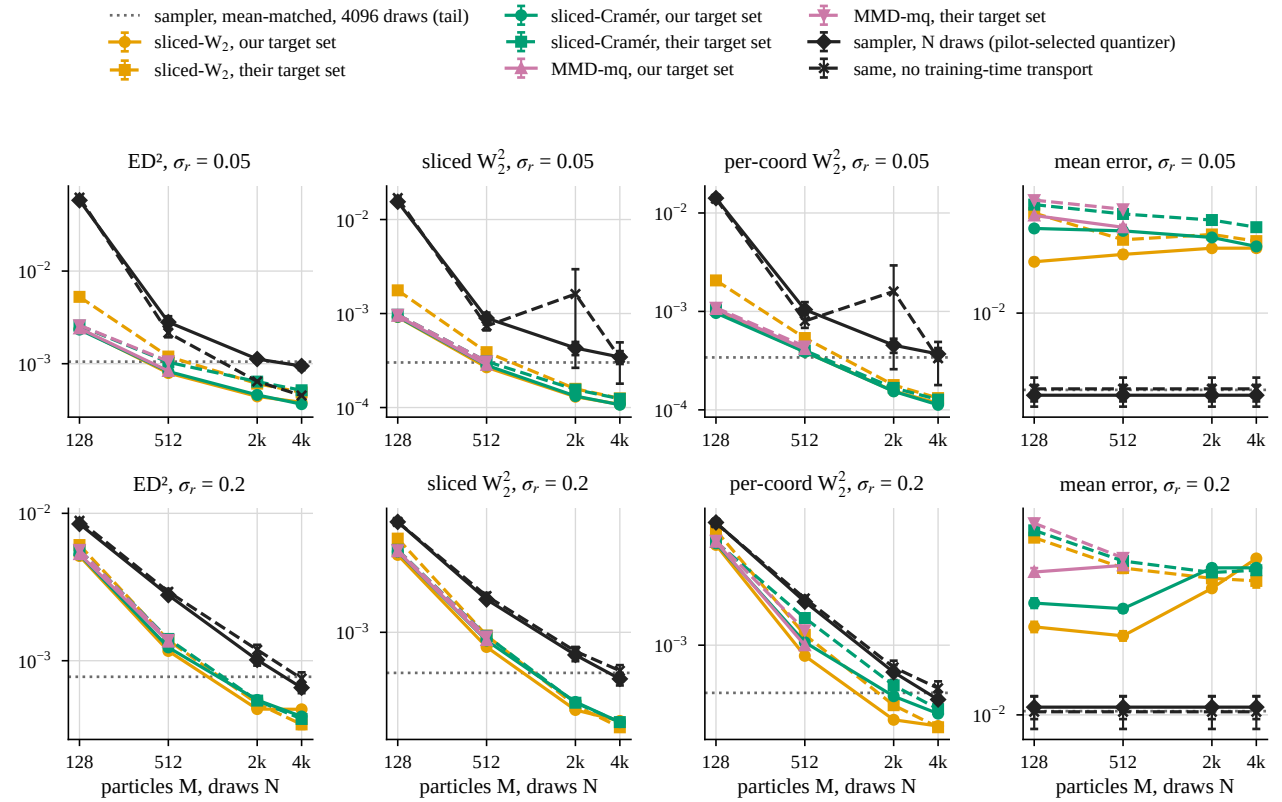

Figure 8: *More particles for the competitors.* The chain with reward noise $\sigma_r = 0.05$ (top) and 0.2 (bottom): the two strongest variants of Debes and Tuytelaars (2026b) and multiquadric MMD (up to $M = 512$) with more particles, on their own and on our target set, at their specified learning rate; squared energy distance, sliced and per-coordinate w_2^2, and mean error against the particle count M of the competitors and the number of draws N of SVDRL's sampler. Dotted: SVDRL's moment-matched sampler with 4096 draws; black: SVDRL's sampler with N draws. Mean ± standard error over the final seeds.

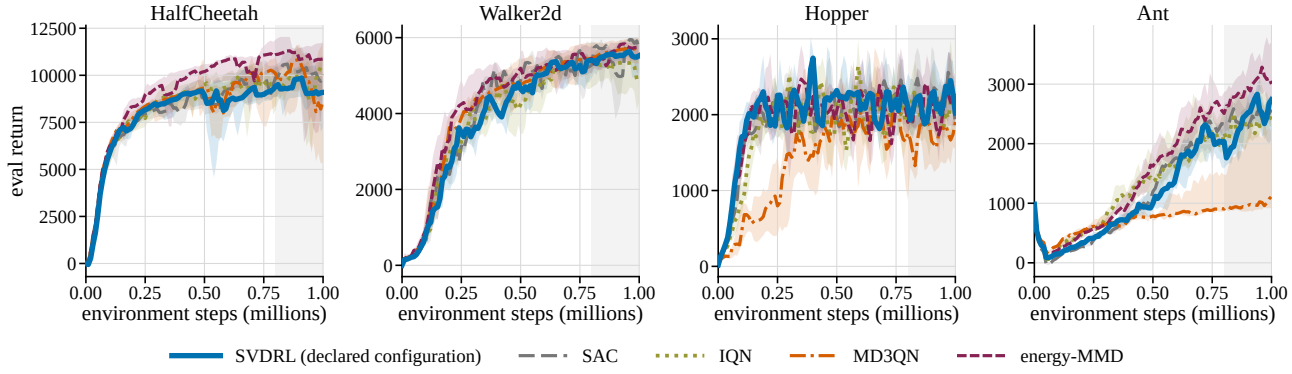

Figure 9: *MuJoCo learning curves.* Evaluation return against environment steps, per run the mean of its last five evaluations, median over seeds with the interquartile band, every head on the same SAC trunk. SVDRL is the declared configuration per environment; the shaded strip marks the steps from 800k to 1M over which Table 3 averages the return. Seeds per head as in Table 3; a run stopped by the divergence guard keeps its last evaluation.

Table 17: *Returns per head.* MuJoCo returns at 1M environment steps: IQM over seeds of each run's mean eval return over 800k–1M (21 evaluations, as in Table 3) and of its last 5 evaluations, with 95 % bootstrap CIs over seeds. The SVDRL configuration per environment is declared (HalfCheetah: the default configuration; Walker2d, Hopper, Ant: target mean, the variant that takes the moment-matching mean from the bootstrap targets; a per-environment choice between two declared variants); the rows marked * are the other SVDRL variant run on that environment (HalfCheetah: log-variance moment head). †Includes runs stopped by the divergence guard, counted with the evaluations they logged: Ant SVDRL (target mean), one seed (16 window evaluations).

Env	Head	800k–1M IQM [95 % CI] (↑)	last-5 IQM [95 % CI] (↑)
HalfCheetah	SVDRL (default)	8720[7565, 9889]	8821[7158, 10035]
	SAC	10014[9446, 10512]	9988[8673, 10742]
	IQN	9219[8338, 9879]	10227[7198, 11087]
	MD3QN	8986[8449, 9877]	8676[5406, 11376]
	energy-MMD	10098[8885, 11186]	10571[8256, 11855]
	SVDRL, log-variance head*	9193[7393, 10571]	8640[4898, 11551]
Walker2d	SVDRL (target mean)	5514[5025, 5780]	5482[4883, 5843]
	SAC	5372[4939, 5744]	5543[4925, 5956]
	IQN	5083[4466, 5496]	4928[4289, 5501]
	MD3QN	5661[5129, 5907]	5695[5235, 5984]
	energy-MMD	5609[5361, 5805]	5713[5434, 5988]
	SVDRL (default)*	3774[3560, 4513]	3422[2249, 4629]
Hopper	SVDRL (target mean)	2187[2047, 2294]	2061[1849, 2334]
	SAC	2187[2083, 2319]	1935[1738, 2483]
	IQN	2019[1841, 2221]	2089[1753, 2416]
	MD3QN	1822[1743, 1889]	1943[1524, 2425]
	energy-MMD	2180[2047, 2255]	2214[1685, 2529]
Ant	SVDRL (target mean)	2429[2059, 2654]	2621[2124, 2952]
	SAC	2426[2214, 2759]	2552[2305, 2995]
	IQN	2411[2095, 3029]	2523[2242, 3074]
	MD3QN	1233[895, 2343]	1316[917, 2545]
	energy-MMD	2928[2628, 3598]	3088[2617, 3910]
	SVDRL (default)*	1052[956, 1165]	1176[887, 1602]

Table 18: *MuJoCo critic accuracy.* Critic laws against frozen-policy Monte Carlo returns at 10^6 steps, all metrics in raw return units relative to the run's own MC draws: shape = w_2 of the centered laws / MC std; std ratio (ideal 1); mean error (critic − MC); full w_2 (CDF estimator, cross-checked by the quantile estimator) / MC std; coverage: the fraction of MC draws inside the critic's central 80 % interval, raw (80 % cov.) and with both means removed (centered), ideal 0.8. IQM over seeds, with the 95 % bootstrap interval of shape and full w_2 below it. Every head's mean is biased (twin-min targets), so full w_2 and raw coverage do not separate the heads. SVDRL (target mean): the variant that takes the moment-matching mean from the bootstrap targets. * The other declared SVDRL variant (appendix row). The sampler row (reverse diffusion from the saved score nets) is shown only for the log-variance head, the only configuration whose nets pass the label-free gate; it is a diagnostic, not a sampler claim. SAC: mean error only (scalar critic).

Head	shape (↓)	std ratio	mean error	full w_2 (↓)	80 % cov.	centered
HalfCheetah						
SVDRL (default)	0.43 [0.28, 0.56]	1.36	-88	4.9 [3.7, 7.0]	0.01	0.92
IQN	1.05 [0.82, 1.41]	2.01	-74	4.6 [3.3, 5.7]	0.05	0.97
MD3QN	0.69 [0.50, 0.92]	1.64	-50	6.8 [4.9, 8.9]	0.00	0.96
energy-MMD	0.52 [0.39, 0.71]	1.46	-76	4.1 [3.3, 4.9]	0.02	0.94
SVDRL, log-variance head*	0.51 [0.44, 0.91]	1.46	-10	4.2 [3.7, 6.0]	0.01	0.94
its sampler*	0.54 [0.43, 0.98]	1.49	-11	4.2 [3.7, 6.2]	0.01	0.95
SAC	–	–	-55	–	–	–
Walker2d						
SVDRL (target mean)	0.92 [0.61, 1.31]	1.34	-68	4.3 [1.9, 9.9]	0.02	0.95
IQN	3.99 [2.32, 6.13]	4.85	-44	6.3 [3.6, 8.8]	0.46	1.00
MD3QN	1.19 [0.90, 2.75]	1.94	-18	3.4 [1.9, 7.9]	0.10	0.99
energy-MMD	2.08 [0.82, 7.51]	2.84	-21	3.9 [1.6, 10.2]	0.51	0.98
SVDRL (default)*	0.70 [0.42, 1.10]	0.67	100	3.7 [1.8, 17.1]	0.02	0.68
SAC	–	–	-14	–	–	–
Hopper						
SVDRL (target mean)	2.25 [0.87, 4.56]	2.95	-8	5.6 [3.5, 18.2]	0.20	0.99
IQN	10.4 [6.7, 20.6]	11.18	2	14.2 [8.9, 26.2]	0.24	1.00
MD3QN	6.10 [1.69, 14.74]	6.69	-21	15.5 [4.4, 33.4]	0.00	0.99
energy-MMD	21.3 [10.7, 44.5]	21.97	-28	26.0 [11.9, 62.5]	0.22	1.00
SAC	–	–	2	–	–	–
Ant						
SVDRL (target mean)	0.59 [0.57, 0.63]	0.74	-27	3.1 [2.7, 3.4]	0.03	0.80
IQN	0.73 [0.42, 1.19]	1.24	-37	1.7 [1.2, 2.2]	0.27	0.85
MD3QN	0.72 [0.56, 2.95]	1.52	-22	3.1 [1.8, 4.3]	0.16	0.93
energy-MMD	0.49 [0.45, 0.56]	0.75	-10	2.0 [1.5, 2.4]	0.08	0.82
SVDRL (default)*	0.57 [0.48, 0.63]	0.51	-23	1.7 [1.1, 2.7]	0.11	0.53
SAC	–	–	-14	–	–	–

Table 19: *Theory checks.* Final checkpoint, mean ± 1.96 s.e. over seeds (raw units unless stated). (a, b) bootstrap source and particle set (Theorem 4); (c) multi-level vs single-level networks (Remark 2), state–action pair s_0; (d) Theorem 3, normalized units except W_1; (e) multi-level risk $\mathcal{E}_{\mathrm{ml}}$ (normalized, summed over evaluated state–action pairs) under DSM, consistency and both; n: seeds with a finite final checkpoint.

(a, b)	bootstrap / particles	sampler $\bar{W}_1$ (↓)	particles $\bar{W}_1$ (↓)	$\varepsilon_{\mathrm{St}}$ (s_0)	sampler mode mass (s_0)
MDP1	particle (Stein particles)	$0.132_{\pm 0.029}$	$0.247_{\pm 0.032}$	$0.214_{\pm 0.011}$	–
MDP1	sampler	$0.454_{\pm 0.165}$	$0.558_{\pm 0.244}$	$0.289_{\pm 0.071}$	–
MDP1	sampler, moment-matched	$0.139_{\pm 0.030}$	$0.209_{\pm 0.042}$	$0.173_{\pm 0.031}$	–
MDP1	sampler-seeded	$0.372_{\pm 0.107}$	$0.472_{\pm 0.156}$	$0.288_{\pm 0.077}$	–
MDP1	sampler-seeded, moment-matched	$0.132_{\pm 0.025}$	$0.208_{\pm 0.036}$	$0.182_{\pm 0.029}$	–
MDP2	particle (Stein particles)	$0.601_{\pm 0.135}$	$1.61_{\pm 0.33}$	$1.40_{\pm 0.31}$	$0.0260_{\pm 0.0166}$
MDP2	sampler	$2.32_{\pm 0.22}$	$2.36_{\pm 0.30}$	$1.42_{\pm 0.37}$	$0.0322_{\pm 0.0233}$
MDP2	sampler, moment-matched	$0.642_{\pm 0.103}$	$1.63_{\pm 0.38}$	$1.33_{\pm 0.43}$	$0.0263_{\pm 0.0145}$
MDP2	sampler-seeded	$2.00_{\pm 0.19}$	$2.14_{\pm 0.19}$	$1.04_{\pm 0.16}$	$0.0300_{\pm 0.0125}$
MDP2	sampler-seeded, moment-matched	$0.668_{\pm 0.125}$	$1.48_{\pm 0.46}$	$1.30_{\pm 0.43}$	$0.0309_{\pm 0.0132}$
MDP3	particle (Stein particles)	$3.34_{\pm 0.40}$	$6.16_{\pm 1.28}$	$4.68_{\pm 1.39}$	$0.0645_{\pm 0.0199}$
MDP3	sampler	$104_{\pm 8}$	$103_{\pm 8}$	$13.0_{\pm 2.3}$	$0.466_{\pm 0.074}$
MDP3	sampler, moment-matched	$3.44_{\pm 0.38}$	$6.36_{\pm 2.07}$	$5.43_{\pm 1.95}$	$0.0616_{\pm 0.0237}$
MDP3	sampler-seeded	$109_{\pm 5}$	$111_{\pm 5}$	$12.7_{\pm 1.7}$	$0.400_{\pm 0.052}$
MDP3	sampler-seeded, moment-matched	$3.35_{\pm 0.33}$	$6.04_{\pm 1.87}$	$4.84_{\pm 1.97}$	$0.0561_{\pm 0.0198}$

(c)	network	mode mass	W_1	multi-level risk (norm.)	own-level $\mathcal{E}_\sigma$ (unnorm.)
MDP2	main (multi-level, DSM + cons.)	$0.0260_{\pm 0.0166}$	$0.601_{\pm 0.135}$	$0.0944_{\pm 0.0229}$	–
MDP2	multi-level, DSM only	$0.0563_{\pm 0.0346}$	$0.909_{\pm 0.275}$	$0.182_{\pm 0.050}$	–
MDP2	single level $\sigma_{\min}$	$0.320_{\pm 0.010}$	$5.96_{\pm 2.30}$	$251_{\pm 155}$	$3.20_{\pm 0.96}$
MDP2	single level $\Delta/2$	$0.181_{\pm 0.022}$	$6.09_{\pm 0.10}$	$2.56_{\pm 0.09}$	$0.00342_{\pm 0.00143}$
MDP3	main (multi-level, DSM + cons.)	$0.0645_{\pm 0.0199}$	$3.34_{\pm 0.40}$	$0.364_{\pm 0.100}$	–
MDP3	multi-level, DSM only	$0.0986_{\pm 0.0440}$	$3.50_{\pm 0.76}$	$0.438_{\pm 0.135}$	–
MDP3	single level $\sigma_{\min}$	$0.461_{\pm 0.059}$	$31.7_{\pm 6.2}$	$236_{\pm 261}$	$18.3_{\pm 4.7}$
MDP3	single level $\Delta/2$	$0.0875_{\pm 0.0358}$	$16.6_{\pm 0.3}$	$30.0_{\pm 9.9}$	$0.0330_{\pm 0.0088}$

(d)	state–action pair	$\mathcal{E}_{\mathrm{ml}}$	$\varepsilon_{\mathrm{disc}}$	KL_0	$\varepsilon_{\mathrm{TV}}$	TV meas.	W_1 bound	W_1 meas.
MDP1	s_0	$0.0285_{\pm 0.0119}$	$0.2490_{\pm 0.0005}$	$0.0175_{\pm 0.0001}$	$0.384_{\pm 0.008}$	$0.0461_{\pm 0.0072}$	$53.9_{\pm 1.1}$	$0.136_{\pm 0.030}$
MDP2	s_0	$0.0944_{\pm 0.0229}$	$0.5505_{\pm 0.0002}$	$0.01469_{\pm 0.00005}$	$0.574_{\pm 0.010}$	$0.0839_{\pm 0.0177}$	$92.8_{\pm 1.6}$	$0.575_{\pm 0.135}$
MDP2	s_1	$0.00449_{\pm 0.00199}$	$0.6954_{\pm 0.0002}$	$1.76 \times 10^{-4}_{\pm 4.8\times 10^{-7}}$	$0.5916_{\pm 0.0009}$	$0.0252_{\pm 0.0098}$	$95.6_{\pm 0.1}$	$0.346_{\pm 0.030}$
MDP2	s_2	$0.0102_{\pm 0.0044}$	$0.6468_{\pm 0.0002}$	$3.15 \times 10^{-4}_{\pm 2.1\times 10^{-6}}$	$0.573_{\pm 0.002}$	$0.0370_{\pm 0.0103}$	$92.7_{\pm 0.3}$	$0.306_{\pm 0.036}$
MDP3	s_0	$0.364_{\pm 0.100}$	$1.328_{\pm 0.002}$	$0.0312_{\pm 0.0002}$	$0.927_{\pm 0.027}$	$0.199_{\pm 0.023}$	$296_{\pm 8}$	$3.10_{\pm 0.34}$
MDP3	s_1	$0.0988_{\pm 0.0264}$	$1.524_{\pm 0.001}$	$1.38 \times 10^{-4}_{\pm 1.3\times 10^{-6}}$	$0.901_{\pm 0.007}$	$0.179_{\pm 0.007}$	$287_{\pm 2}$	$1.76_{\pm 0.04}$
MDP3	s_2	$0.115_{\pm 0.071}$	$1.544_{\pm 0.001}$	$1.02 \times 10^{-4}_{\pm 1.3\times 10^{-6}}$	$0.910_{\pm 0.019}$	$0.169_{\pm 0.010}$	$290_{\pm 6}$	$1.87_{\pm 0.05}$
MDP3	s_3	$0.0733_{\pm 0.0365}$	$1.507_{\pm 0.001}$	$1.46 \times 10^{-4}_{\pm 8.9\times 10^{-7}}$	$0.889_{\pm 0.010}$	$0.200_{\pm 0.006}$	$284_{\pm 3}$	$1.60_{\pm 0.05}$

(e)	training signal	$\mathcal{E}_{\mathrm{ml}}$	sampler $\bar{W}_1$	n
3-cycle	DSM only	$5.08 \times 10^{-5}_{\pm 1.6\times 10^{-5}}$	$0.0420_{\pm 0.0002}$	10/10
3-cycle	consistency only	$6.27 \times 10^{-4}_{\pm 5.2\times 10^{-4}}$	$0.0420_{\pm 0.0002}$	10/10
3-cycle	DSM + consistency	$2.48 \times 10^{-5}_{\pm 1.4\times 10^{-5}}$	$0.0420_{\pm 0.0002}$	10/10
3-cycle	consistency only, grid pairing	$4.51 \times 10^{-4}_{\pm 3.3\times 10^{-4}}$	$0.0420_{\pm 0.0002}$	10/10
3-cycle	DSM + consistency, grid pairing	$1.67 \times 10^{-5}_{\pm 5.0\times 10^{-6}}$	$0.0420_{\pm 0.0002}$	10/10
MDP1	DSM only	$0.0493_{\pm 0.0185}$	$0.172_{\pm 0.031}$	10/10
MDP1	DSM + consistency (deployed)	$0.0285_{\pm 0.0119}$	$0.132_{\pm 0.029}$	10/10
MDP1	consistency only	$5.21 \times 10^{15}_{\pm 1.0\times 10^{16}}$	$79.0_{\pm 21.6}$	2/10
MDP1	consistency only, grid pairing	$5.51 \times 10^{16}_{\pm 8.8\times 10^{16}}$	$74.0_{\pm 19.2}$	5/10
MDP1	consistency only, grid pairing, EMA target	$0.0143_{\pm 0.0037}$	$0.144_{\pm 0.022}$	10/10

Table 20: *Score-space temporal differences.* Level-by-level check of Corollary 3 on the trained networks: per variant, the seeds in which the one-step inequality (16) and its iterated form hold at every level, and the ratio of the iterated bound to the measured normalized error at the highest, a middle and the lowest noise level. Normalized units, grid ratio $\kappa_g = \gamma$.

MDP	variant	one-step holds	iterated holds	bound/error top	middle	bottom
3-cycle	DSM only	10/10	10/10	2.54±0.24	53.7±15.7	109±37
3-cycle	consistency only	10/10	10/10	1.00010±0.00005	6.47±1.74	11.4±3.8
3-cycle	DSM + consistency	10/10	10/10	1.00012±0.00006	32.8±6.0	74.5±19.4
3-cycle	consistency only, grid pairing	10/10	10/10	1.00006±0.00002	6.56±2.13	12.1±5.2
3-cycle	DSM + consistency, grid pairing	10/10	10/10	1.00012±0.00004	35.4±7.9	80.2±18.0
MDP1	DSM only	10/10	10/10	1.31±0.18	3.02±0.83	30.5±11.0
MDP1	DSM + consistency (deployed)	10/10	10/10	1.32±0.15	4.79±2.19	39.9±15.2
MDP1	consistency only	2/2	2/2	1.000000004±0.000000006	124±152	9.94±0.02
MDP1	consistency only, grid pairing	5/5	5/5	1.0000001±0.0000001	62.7±114.8	98.7±80.9
MDP1	consistency only, grid pairing, EMA target	10/10	10/10	3.43±1.74	2.36±0.30	15.8±1.3

Table 21: *Width of the far rare mode.* Static fit of the score network to fixed draws of the root law of the rare-event chain ($p = 0.005$, $D = 2$), without bootstrapping, consistency loss or EMA, three seeds: mass, median offset and standard deviation of the rare component in return units, and the ratio of the standard deviation to the true one.

seed	rare mass	median offset	sd	sd ÷ true sd
0	0.00667	+1.7	7.42	12.9
1	0.00958	+4.7	8.22	14.2
2	0.0110	+8.1	8.48	14.7
truth	0.005	0	0.578	1

Table 22: *The competitors' pixel benchmark.* Pixel maze of Debes and Tuytelaars (2026b), policy evaluation of the uniform random policy: per variant and head, the number of seeds n and the median over seeds of each distance to the Monte Carlo returns at the start state, with its interval (the estimate $\hat{w}_2$ of Debes and Tuytelaars (2026b), $\hat{w}_2$ against 200 Monte Carlo returns, and CDF-based sliced w_2). SVDRL with the log-variance moment head and with its earlier configuration.

variant	head	n	$\hat{W}_2$ (Debes) (↓)	$\hat{W}_2$ (200 MC) (↓)	sliced W_2 (CDF) (↓)
exclusive	MMD-mq	5	**0.309**[0.30, 0.31]	0.121[0.12, 0.13]	0.0385[0.033, 0.040]
	sliced-Cramér	5	0.309[0.30, 0.31]	**0.116**[0.11, 0.12]	**0.0249**[0.022, 0.031]
	SVDRL (log-var. moment head)	5	0.422[0.42, 0.44]	0.367[0.34, 0.39]	0.224[0.19, 0.25]
	sliced-W_2	5	0.512[0.49, 0.54]	0.506[0.50, 0.51]	0.352[0.35, 0.35]
	SVDRL (earlier selected configuration)	5	0.761[0.43, 0.81]	0.702[0.40, 0.71]	0.378[0.26, 0.43]
identical	sliced-Cramér	5	**0.253**[0.25, 0.26]	**0.0960**[0.093, 0.10]	**0.0313**[0.027, 0.040]
	MMD-mq	5	0.263[0.25, 0.27]	0.115[0.11, 0.12]	0.0532[0.050, 0.061]
	SVDRL (log-var. moment head)	5	0.326[0.30, 0.34]	0.272[0.26, 0.28]	0.182[0.17, 0.19]
	SVDRL (earlier selected configuration)	5	0.331[0.30, 0.55]	0.238[0.19, 0.46]	0.139[0.11, 0.26]
	sliced-W_2	5	0.387[0.38, 0.40]	0.377[0.37, 0.38]	0.260[0.26, 0.27]
multireward	MMD-mq	5	**0.271**[0.26, 0.31]	0.134[0.13, 0.14]	0.0322[0.031, 0.038]
	sliced-Cramér	5	0.273[0.26, 0.31]	**0.129**[0.12, 0.14]	**0.0250**[0.024, 0.026]
	SVDRL (log-var. moment head)	5	0.398[0.32, 0.43]	0.383[0.29, 0.40]	0.145[0.12, 0.16]
	SVDRL (earlier selected configuration)	5	0.433[0.42, 0.48]	0.412[0.40, 0.45]	0.144[0.14, 0.16]
	sliced-W_2	5	0.485[0.48, 0.49]	0.487[0.49, 0.49]	0.242[0.24, 0.24]

Table 23: *Multi-objective MuJoCo.* Distributional policy evaluation of fixed policies after $5 \cdot 10^5$ and 10^6 steps: energy distance and sliced w_2^2 to the Monte Carlo law and between the centered laws, and the absolute mean error, each the median over seeds with its 95 % bootstrap interval, a run's value its mean over the last 20 % of the evaluations. SVDRL's initial and selected configurations as particles and sampler; †: the configuration's sampler fails its ground-truth-free gate. IQN models per-dimension marginals, so its joint metrics are empty. On mo-halfcheetah the heads with the centered PopArt accumulator and SVDRL's selected configuration, which uses it, are compared, and the uncentered rows are kept for reference. Bold: inside the interval of the best compared median of its column.

	Law (↓)		Mean (↓)	Centered law (↓)	
Head	energy	sliced w_2^2	abs. error	energy	sliced w_2^2
mo-halfcheetah, $5 \cdot 10^5$ *steps*					
SVDRL selected, particles†	0.1[0.098, 0.14]	0.17[0.16, 0.22]	0.17[0.15, 0.18]	0.065[0.055, 0.087]	0.14[0.12, 0.18]
MD3QN	**0.038**[0.037, 0.039]	**0.062**[0.061, 0.063]	**0.15**[0.15, 0.16]	**0.0073**[0.0066, 0.0073]	**0.027**[0.026, 0.028]
Energy MMD	0.05[0.047, 0.053]	0.09[0.086, 0.096]	0.17[0.16, 0.18]	0.014[0.014, 0.015]	0.048[0.046, 0.05]
IQN	–	–	**0.15**[0.15, 0.16]	–	–
uncentered accumulator, for reference					
SVDRL initial, particles	0.33[0.16, 0.8]	0.76[0.27, 2.2]	0.29[0.21, 0.66]	0.11[0.09, 0.3]	0.37[0.21, 1.5]
SVDRL initial, sampler	13[0.35, 110]	280[1.4, 4 100]	4.1[0.29, 33]	2.1[0.25, 21]	84[1.2, 1 200]
MD3QN	0.038[0.034, 0.04]	0.064[0.058, 0.065]	0.15[0.14, 0.16]	0.0077[0.0067, 0.0091]	0.028[0.027, 0.03]
Energy MMD	0.081[0.053, 1.2]	0.15[0.12, 26]	0.16[0.15, 0.67]	0.033[0.02, 1.2]	0.11[0.084, 25]
IQN	–	–	0.15[0.15, 0.16]	–	–
mo-halfcheetah, 10^6 *steps*					
SVDRL selected, particles†	0.094[0.063, 0.14]	0.22[0.15, 0.28]	**0.2**[0.18, 0.27]	0.05[0.027, 0.069]	0.14[0.097, 0.17]
MD3QN	**0.076**[0.063, 0.093]	**0.17**[0.15, 0.19]	**0.22**[0.19, 0.25]	**0.025**[0.022, 0.027]	**0.093**[0.083, 0.095]
Energy MMD	0.14[0.12, 0.19]	0.29[0.26, 0.34]	0.29[0.28, 0.36]	0.048[0.043, 0.052]	0.15[0.14, 0.16]
IQN	–	–	**0.24**[0.23, 0.28]	–	–
uncentered accumulator, for reference					
SVDRL initial, particles	0.68[0.26, 0.75]	1.6[0.68, 4.4]	0.55[0.33, 0.67]	0.23[0.084, 0.47]	0.91[0.5, 3.9]
SVDRL initial, sampler	15[0.64, 260]	1 300[5.6, 17 000]	6.6[0.43, 85]	11[0.38, 52]	840[4, 3 900]
MD3QN	0.087[0.073, 0.25]	0.18[0.16, 0.42]	0.19[0.19, 0.34]	0.028[0.026, 0.033]	0.098[0.093, 0.11]
Energy MMD	0.61[0.23, 1.6]	21[1.5, 54]	1[0.29, 1.7]	1.2[0.18, 1.8]	16[1.4, 47]
IQN	–	–	0.24[0.19, 0.72]	–	–
mo-hopper, $5 \cdot 10^5$ *steps*					
SVDRL selected, particles	1.1[0.97, 1.1]	1.3[1.1, 1.3]	0.51[0.49, 0.52]	0.23[0.2, 0.23]	0.3[0.28, 0.32]
SVDRL selected, sampler	1.1[0.98, 1.1]	1.2[1.2, 1.3]	0.51[0.5, 0.54]	0.22[0.21, 0.23]	0.3[0.3, 0.33]
SVDRL initial, particles	0.92[0.88, 0.97]	0.97[0.93, 1]	0.45[0.44, 0.46]	0.15[0.14, 0.16]	0.2[0.18, 0.21]
SVDRL initial, sampler	1.4[1, 1.8]	3.6[2.9, 18]	0.73[0.55, 0.85]	0.4[0.26, 0.48]	2.5[1, 17]
MD3QN	0.86[0.68, 1.6]	1.6[0.81, 9.2]	0.56[0.42, 1.2]	0.37[0.21, 1.9]	0.74[0.26, 6]
Energy MMD	**0.74**[0.71, 0.79]	**0.65**[0.6, 0.68]	**0.38**[0.37, 0.41]	**0.07**[0.065, 0.073]	**0.064**[0.063, 0.066]
IQN	–	–	0.48[0.46, 0.52]	–	–
mo-hopper, 10^6 *steps*					
SVDRL selected, particles†	2.9[2.7, 3.1]	8.4[6.5, 8.9]	**1.5**[1.5, 1.6]	0.76[0.36, 0.88]	3[1.7, 3.7]
SVDRL initial, particles	2.8[2.7, 3.1]	**6.8**[6.1, 7.2]	**1.4**[1.4, 1.5]	**0.26**[0.24, 0.44]	**1.3**[1.3, 1.8]
SVDRL initial, sampler	2.7[2.7, 5.2]	12[7.1, 99]	**1.4**[1.4, 2.4]	**0.41**[0.27, 1.9]	7.3[1.8, 63]
MD3QN	6[4.1, 7.4]	110[69, 120]	2.6[2, 3.3]	3[2.7, 3.6]	56[53, 70]
Energy MMD	**2.5**[2.3, 2.6]	8.1[6.8, 11]	**1.4**[1.3, 1.5]	**0.41**[0.36, 0.42]	3.6[1.8, 5.5]
IQN	–	–	**1.5**[1.5, 1.6]	–	–