arXiv:2611.02895v1 [physics.soc-ph] 2 Nov 2026

Zero-shot time-series foundation models against the Khepri v1.5 baselines for 96-hour carbon-intensity forecasts in nine Nordic bidding zones: a preregistered comparison

Eirik Botten Nicolaysen
EcoDeco AS, Norway
eirik@ecodeco.no

2026-11-02

Abstract

We compare three zero-shot time-series foundation models (Timestral-5M, Timestral-209M and Chronos-2, 8192-hour context, no covariates) with the Khepri v1.5 SARIMA and GBM baselines, with their code and hyperparameters unchanged, for 96-hour forecasts of hourly carbon intensity in the Nordic bidding zones NO1–NO5 and SE1–SE4. The design, the analysis code and a minimum detectable effect (MDE) were timestamped before any foundation-model forecast by an OpenTimestamps proof (Bitcoin-confirmed) and the GPG-signed tag prereg-v1; the proofs are archived in Zenodo (DOI 10.5281/zenodo.23881315). The preregistration was also submitted to OSF Registries 2026-11-02 07:49 UTC (osf.io/hncz5; embargo requested, pending OSF approval). On the 2025 test (361 daily origins, contamination not ruled out) all three models have a lower zone-averaged MAPE than both baselines, by 0.952 to 1.347 percentage points (Diebold–Mariano with the Harvey–Leybourne–Newbold correction, all $p < 0.0001$), and each beats SARIMA in 9 of 9 zones. On a post-release holdout of 14 origins from 2026-10-16 the point estimates have the same sign (0.729 to 1.326 percentage points lower), but all lie within the preregistered minimum detectable effect of 2.989 points and are reported as not determined. For calibration (RQ4) the preregistration locked a Diebold–Mariano–HLN test on CRPS; the MDE applies to MAPE only. Holdout CRPS point estimates (n = 14 origins) are 1.1114 (Timestral-5M), 1.0733 (Timestral-209M) and 1.1198 (Chronos-2) against 1.4165 gCO_2-eq/kWh for SARIMA ($p < 0.0001$ for each comparison). The SARIMA quantiles are constructed from the fitted model's Gaussian predictive distribution (deviation D-03), not taken from a native v1.5 forecast, so the difference may partly reflect the SARIMA intervals. The contamination indicator does not exceed its preregistered threshold for any comparison; with 14 holdout origins it has low power (HAC standard errors 0.124 to 0.360 pp). Forecasts, the CI series, the locked analysis and all deviations are released with the code.

1 Introduction

Forecasts of the carbon intensity (CI) of electricity, in grams of CO2-equivalent per kWh, have been proposed as a building block for shifting flexible electricity use to hours with lower emissions (Maji et al. 2022a).

For the nine Nordic bidding zones NO1–NO5 and SE1–SE4, Khepri v1.5 computes an hourly CI series from ENTSO-E generation per production type and evaluates SARIMA and gradient-boosting (GBM) baselines over a 96-hour horizon from daily origins (Khepri v1.5, DOI 10.5281/zenodo.23101127).

On the 2025 test, GBM beat SARIMA in 7 of the 9 zones, by 0.02 to 2.64 percentage points (pp) of mean absolute percentage error (MAPE), and lost in NO4 (+2.92 pp) and NO5 (+0.53 pp).

Time-series foundation models (TSFMs) are pretrained forecasters that can be applied to a new series without training on it ("zero-shot"). Chronos-2 is described by its developers as offering zero-shot support for univariate, multivariate and covariate-informed forecasting (Ansari et al. 2025). Timestral, released by Mistral AI on 2026-10-15 (Mistral AI 2026), comes in four sizes from 5M to 209M parameters with 21 predictive quantiles per time step.

For a CI series, a zero-shot TSFM needs only the history of the series itself. This makes it directly comparable with the Khepri baselines, which are also univariate and fitted per zone.

TSFMs have already been applied to CI: the public CarbonX repository lists zero-shot forecasts with five TSFMs for 214 grids, among them Norway and Sweden as whole countries (CarbonX).

At bidding-zone level the baselines differ between zones: in v1.5, NO5 had by far the lowest MAPE level (SARIMA 0.72 %), and NO4 and NO5 were the two zones where GBM lost to SARIMA.

Evaluating a pretrained model on historical data raises a contamination problem. If the test period, or the series itself, was part of the pretraining corpus, a retrospective test can overstate the accuracy the model would have on data it has not seen. The repository of the live benchmark Impermanent states the same concern for static benchmarks: "[t]raining data may overlap with evaluation data" (Garza et al. 2026).

For the models used here the corpus cannot be checked directly. The Timestral repository discloses neither training data nor a training cut-off. For Chronos-2 the pretraining datasets are named on its model card; in our reading of the card no dataset name refers to ENTSO-E, Norway or Sweden, but a name match cannot rule out that Nordic data entered through aggregated datasets.

We therefore pair two tests. A retrospective test on the Khepri v1.5 2025 test set (361 daily origins per zone) has more statistical power but is labelled "contamination not ruled out". A post-release holdout uses only origins from 2026-10-16, the day after the Timestral release, so that every target hour lies after the release of all three TSFMs; it has 14 origins per zone.

Fourteen origins give little power. The minimum detectable effect (MDE) for the zone-averaged MAPE difference, estimated from the SARIMA − GBM loss difference on the 2025 test, is 2.989 pp on the holdout, against 0.413 pp on the retro test. A holdout difference smaller than this is reported as "not determined", not as equivalence.

The holdout is therefore a contamination control, and the difference between the TSFM advantage on the two tests is reported as a descriptive contamination indicator, not as a test.

The contributions are as follows. We make no claim of priority.

- A preregistered comparison of three zero-shot TSFMs (Timestral-5M, Timestral-209M, Chronos-2) with the Khepri v1.5 SARIMA and GBM baselines and a seasonal-naive reference, for 96-hour CI forecasts in nine Nordic bidding zones. Design, analysis code and MDE were timestamped before any foundation-model forecast by an OpenTimestamps proof (Bitcoin-confirmed) and the GPG-signed tag `prereg-v1`; the proofs are archived in Zenodo (10.5281/zenodo.23881315). The preregistration was also submitted to OSF Registries 2026-11-02 07:49 UTC (osf.io/hncz5; embargo requested, pending OSF approval).

- A design that pairs a retrospective test with a small minimum detectable effect and a post-release holdout, with the minimum detectable effect and the decision rule fixed before any TSFM forecast was made.

- Hourly predictions of every arm, the derived CI series and the code, stored with sha256 sums and OpenTimestamps stamps, and a deviation log in which every deviation is dated.

- A prospective log that issues all arms daily from the registration onward, to be analysed later with the same locked analysis code.

The rest of the paper is organised as follows. Section 2 reviews related work, Section 3 describes the data, arms, tests and deviations, Section 4 reports the results and Section 5 describes the software and its availability.

2 Related work

2.1 Carbon-intensity forecasting

DACF (Maji et al. 2022a) forecasts the hourly average CI of a grid for the next 24 hours. It predicts the production of each generation source with a separate neural network that uses historical production, calendar features and day-ahead weather forecasts, and combines the source forecasts with direct (operational) emission rates.

It is evaluated on six regions, four in the US and two in Europe, Sweden (SE) and Germany (DE), with ENTSO-E data for the European regions. It uses data from 2019–2021, retrains its models every six months and reports predictions for the last six months of 2021 as the test, with MAPE as the metric.

DACF is compared with the authors' own reimplementation of the method of Leerbeck et al. (2020), since that code was not public.

According to its public repository, CarbonCast (Maji et al. 2022b) extends DACF to forecasts of up to 96 hours. The repository covers 13 regions, among them Sweden (SE) as the only Nordic region, and its saved models are trained on 2020 and validated on January–June 2021; the repository states that results similar to the paper are obtained when testing on July–December 2021. It also points to a real-time v3.0 branch in testing, which we did not read.

The repository of EnsembleCI (Yan et al. 2025) implements CI forecasting with gradient-boosting, neural and ensemble learners (LightGBM, a FastAI network, CatBoost, XGBoost and an AutoGluon weighted ensemble) and lists 11 regions, with Sweden (SE) as the only Nordic one.

In the region lists and released data of these three repositories, Sweden appears only at country level and Norway does not appear; no series for the bidding zones NO1–NO5 or SE1–SE4 is included.

Leerbeck et al. (2020) forecast average and marginal CO2 emission intensity for the Danish bidding zone DK2 up to 24 hours ahead with 95 % prediction intervals, combining linear models with exogenous inputs (ARIMAX) fitted to data supplied by Tmrow.

The EnsembleCI paper (Yan et al. 2025) evaluates a stacked ensemble of gradient-boosting and neural learners on 11 grids, six in the US and five in the EU with Sweden as a country, up to four days ahead with MAPE, testing the EU grids on the second half of 2021, and reports lower MAPE than CarbonCast on all four days.

Clauß et al. (2019) compute hourly average CO2-equivalent intensities for Scandinavian bidding zones with an input–output method that accounts for trade between zones, from ENTSO-E generation data (Swedish zones from the Swedish TSO), and use them as a control signal for the demand response of residential heating; the intensities are computed for 2015 and are not forecast.

GEFCom2014 had four probabilistic forecasting tracks, on load, price, wind and solar power (Hong et al. 2016).

2.2 Foundation models for carbon intensity

The CarbonX paper (Maji et al. 2025) evaluates five TSFMs (MOMENT, Chronos, TimesFM, Sundial and Time-MoE) zero-shot for CI forecasts up to 96 hours, extended to 21 days. Its benchmark

has 13 grids; the European ones are Sweden (SE), Germany, Poland, Spain and the Netherlands, evaluated on the last six months of 2021, and a global evaluation uses Electricity Maps data for 214 grids (2021–2024, 70/30 split) whose names are not given in the paper. The evaluations are on historical data; we found no discussion in the paper of overlap between the models' pretraining data and the evaluation data, and no Norwegian grid or bidding zone is named.

The CarbonX repository provides zero-shot CI forecasts with five TSFMs (Chronos, MOMENT, Sundial, Time-MoE and TimesFM) for 214 grids at a horizon of up to 96 hours, with an extended mode of up to 21 days.

Norway and Sweden are included as whole countries; the released data for both are hourly series from 2021-01-01 to 2024-12-31, and no bidding-zone series is released.

We found no statement in the repository about prospective evaluation, evaluation on data after the models' release, or overlap between the models' pretraining data and the evaluation data.

Whether the repository corresponds exactly to the evaluation reported in the CarbonX paper could not be verified.

The three TSFMs evaluated here come from two families: Timestral (Mistral AI) and Chronos-2 (Ansari et al. 2025). The Chronos-2 technical report describes a pretrained model for univariate, multivariate and covariate-informed forecasting in a zero-shot manner; its context length is extended to 8192 in a second training stage, multivariate and covariate-informed training relies entirely on synthetic data, and none of the real univariate pretraining datasets it lists (its Table 6) is an ENTSO-E or CI series; the energy-domain datasets listed are Electricity, Solar, Wind Farms, London Smart Meters and Buildings 900K. The report gives no training cut-off or release date.

The Timestral code is released under Apache-2.0 and the weights separately under a "Modified MIT License"; its repository discloses neither training data nor a training cut-off. Chronos-2 became public in October 2025: its technical report appeared on 2025-10-17, its repository announces the release on 2025-10-20, and the Hugging Face model repository amazon/chronos-2 dates from 2025-10-30 (D-12). The quickstart notebook in its repository gives a maximum context of 8192 steps.

2.3 Contamination and live evaluation

According to its repository, Impermanent (Garza et al. 2026) is a live benchmark for temporal generalisation in which models must forecast before the future is observed and are scored as data arrive. The repository motivates this with the risk that, for foundation models, "[t]raining data may overlap with evaluation data" in static benchmarks. The signals it lists are activity series derived from GitHub Archive (stars, pushes, pull requests, issues).

Our post-release holdout is not a live evaluation: its forecasts were produced after the truth had been observed. It guarantees only that the target hours postdate the public release of Timestral (2026-10-15, the date of its GitHub release commit) and of Chronos-2 (October 2025). In the prospective log (Section 3.10) a scheduled job issues the forecast for each 00:00 UTC origin a few hours later, while most of the 96 target hours are still unobserved; the context ends at the origin.

In the Impermanent paper, every model must produce point and probabilistic forecasts for the next horizon at each cutoff before any ground truth is available, and forecasts are scored after the outcomes arrive. The paper evaluates four TSFMs, among them Chronos-2, on GitHub activity series from GH Archive and contains no energy or CI series.

2.4 Forecast comparison and evaluation

Real-time probabilistic forecasting exercises outside energy evaluate forecasts submitted as quantiles with the weighted interval score and the empirical coverage of prediction intervals, for example the

US COVID-19 Forecast Hub, whose forecasts were submitted as 23 quantiles from 0.01 to 0.99 (Cramer et al. 2022).

We compare forecast accuracy with the Diebold–Mariano test (Diebold & Mariano 1995) with the small-sample correction of Harvey, Leybourne & Newbold (1997), and evaluate predictive quantiles with the CRPS (Gneiting & Raftery 2007); details are in Section 3.

3 Methods

3.1 Study design and preregistration

The design, the analysis code and deviations D-01 to D-08 were fixed in the tagged preregistration before any time-series foundation model (TSFM) forecast was produced on real carbon-intensity (CI) data; D-09 to D-12 came later (Section 3.9).

The design, the analysis code and the MDE were timestamped before any foundation-model forecast by an OpenTimestamps proof (Bitcoin-confirmed) of the tarball of the repository at the GPG-signed git tag `prereg-v1` (commit 74539f2af7bb6d203d8c8452e01d96a222156078) and by the signed tag itself; the proofs are archived in Zenodo (10.5281/zenodo.23881315).

The preregistration was also submitted to OSF Registries 2026-11-02 07:49 UTC (osf.io/hncz5; embargo requested, pending OSF approval).

The locked analysis is `analysis.py` with `khepri_tsfm/stats.py` at that tag.

Deviations are listed in Section 3.9 and in full in the repository (`truth/DEVIATIONS_KHEPRI-TSFM_v1_2026-11-02.md`).

The study has two test windows with different roles. A *retrospective* test (retro) on the Khepri v1.5 2025 test set carries the statistical power but is labelled "contamination not ruled out": the training cut-off of Timestral is not disclosed in its repository, and most of the 2025 test period precedes the public release of Chronos-2 in October 2025 (D-12). A *post-release holdout* uses only forecast origins after the public release of Timestral and serves as a contamination control with low power.

3.2 Data and carbon intensity

We use ENTSO-E Transparency Platform data on actual generation per production type (A75) for the nine Nordic bidding zones NO1–NO5 and SE1–SE4.

For 2019–2025 we use the raw files of Khepri v1.5 unchanged. Data for 2026-01-01 to 2026-11-02 were fetched for this study on 2026-11-02 between 07:12 and 07:28 UTC with entsoe-py 0.8.0; the last hour complete for all nine zones is 2026-11-02 05:00 UTC.

Raw ENTSO-E data are not redistributed; their sha256 sums are stamped with OpenTimestamps.

Hourly CI (gCO2-eq/kWh) is computed with Khepri v1.5 (DOI 10.5281/zenodo.23101127, git tag v1.5) without study-specific code: AR5 life-cycle emission factors, the ADR-0002 materiality threshold, pumped-storage hours set to missing, Waste/Other excluded, hourly mean, and the v1.5 gap filling `interpolate(limit=6).ffill().bfill()`.

The same filled series is used as model input and as truth, exactly as in v1.5.

The derived CI series is in the repository (`data/CI_KHEPRI-TSFM_v1_2026-11-02.parquet`, sha256 ab308ae9f49fb9a4229d3ef089bb8b005fd149101a1e9596c9295c3e5cfa95f0).

3.3 Forecast origins, horizon and splits

Forecast origins are daily at 00:00 UTC. Each forecast covers the 96 hours o+1 h … o+96 h, and an origin is used only if the full 96-hour truth is available.

All arms share the same context end: the origin hour, inclusive.

Retro. The Khepri v1.5 2025 test exactly as published: history from 2021 (NO) and 2022 (SE) (v1.5 data years 2021–2024 and 2022–2024), `train_end` 2023-12-31 23:00, test period 2025-01-01 to 2025-12-31, giving 361 origins per zone (2025-01-01 to 2025-12-27).

Holdout. Origins from 2026-10-16, the day after the Timestral release on 2026-10-15, with o+96 h $\leq$ 2026-11-02 05:00 UTC. This gives 14 origins per zone (2026-10-16 to 2026-10-29).

For SARIMA and GBM the v1.5 training rule is applied with test start 2026-10-16, changing parameters only (Deviation D-06): `train_end` 2025-10-15 23:00, history from 2023 (NO) and 2024 (SE) (data years 2023–2025 and 2024–2025), and GBM training origins 2024-10-12 to 2025-10-11.

Because v1.5 reads whole calendar years, the SARIMA training span becomes 2.79 years (NO) and 1.79 years (SE), against 3.0 and 2.0 years in v1.5; preserving the span exactly would have required a change to the v1.5 forecasting code, which we did not make.

The window 2026-01-01 to 2026-06-29, already used as a final holdout by Khepri v1.4, is not used.

3.4 Forecasting arms

All arms are univariate; no covariates are used.

Arm	Specification
Seasonal naive (week)	f(o+h) = y(o+h−168 h), h = 1…96. A reference arm added for this study (D-04); it enters no test.
SARIMA	Khepri v1.5, unchanged: SARIMAX (2,0,1)(1,1,0,24), `enforce_stationarity=False`, `enforce_invertibility=False`, `fit(disp=False, maxiter=50)` once on the training data; per origin `apply()` on a 45-day window ending at the origin and a 96-step forecast; diurnal persistence substituted and flagged if a forecast fails.
GBM	Khepri v1.5, unchanged: LightGBM with 200 estimators, 31 leaves, learning rate 0.05, seeds 0, 1 and 2 averaged; features day of week, month, CI at the origin, 24 lags and horizon; trained once on the last 365 daily origins whose 96-h targets lie at or before `train_end`.

Arm	Specification
Timestral-5M	`mistralai/Timestral-5M` at revision bdd20cfc15a3fea0e4cb5ccef37f8d3ca7aa335f, package `timestral` 1.0.0 (github.com/mistralai/timestral at a534bdbd244d64f3d9d743f46c7bf6fc330ed4ff), maximum context 8192, 96-step forecast with library defaults.
Timestral-209M	`mistralai/Timestral-209M` at revision 9874d252ab0aae16f7f7c999f98f98d6ad7d8365, as above.
Chronos-2	`amazon/chronos-2` at revision 29ec3766d36d6f73f0696f85560a422f50e8498c, `chronos-forecasting` 2.3.2, `predict_quantiles` with prediction length 96 and context length 8192.

The TSFM context is the last 8192 hourly values of the same CI series the baselines use, and the TSFM point forecast is the 0.5 quantile.

All three TSFMs natively provide the same 21 quantile levels 0.01, 0.05, 0.10, ..., 0.90, 0.95, 0.99, which form the common quantile set. SARIMA provides them from its Gaussian predictive distribution (`get_forecast`, predicted mean + $z_q \cdot$ standard error), with the point forecast unchanged (D-03). GBM and the seasonal-naive arm give no quantiles.

Khepri v1.5 does not store hourly predictions, so a driver repeats the v1.5 forecasting loop line by line, calls the same v1.5 functions and keeps the predictions; the v1.5 forecasting module itself is unchanged (D-01). Before registration this driver was validated against stored v1.5 runs: for SE1 the SARIMA per-(origin, day) MAPE differs by at most 0.00036 pp and GBM by at most 3.6e-15; for NO1 GBM differs by at most 4.4e-16.

The NO1 SARIMA fit is not bit-reproducible (D-02): in the pinned environment it converges after 1 L-BFGS iteration, against 22 in the 2026-10-01 v1.5 run; in the validation before registration this gave a NO1 retro mean SARIMA MAPE of 2.4596 against 2.4656 in the stored v1.5 rows. The reproduction table (Section 3.8) reports the final run. Data, code, library versions and BLAS thread count were ruled out as causes; the cause was not found. We report this rather than adjust the fit.

The environment is pinned by `uv.lock` (Python 3.12.13). Inference ran on a CPU (Apple M1, 16 GB) with `torch.set_num_threads(8)` and the Hugging Face hub offline after download.

3.5 Loss

The loss per arm, split, zone and origin is the Khepri v1.5 MAPE over the 96 target hours of the point forecast; hours with |actual| $\leq$ 1e-6 are excluded.

The test series is the zone-averaged loss per origin, the mean over the nine zones.

For probabilistic forecasts (RQ4) we use the quantile-score approximation of the continuous ranked probability score (CRPS) over the 21 levels, CRPS $\approx \frac{2}{21} \sum_k \text{pinball}_{\tau_k}(y, q_k)$, averaged over the 96 hours, in gCO2-eq/kWh, and the empirical coverage of the central 80 % interval [q0.10, q0.90] and 90 % interval [q0.05, q0.95] (Gneiting & Raftery 2007).

Origins are never back-filled. If any compared arm misses any zone at an origin, that origin drops out of the comparison, and the numbers of origins per arm are reported.

3.6 Research questions, hypotheses and tests

The research questions are: RQ1, whether the TSFMs have lower 96-h MAPE than the Khepri v1.5 SARIMA and GBM baselines on the post-release holdout; RQ2, the same question on the retro test; RQ3, how the TSFMs do in NO4 and NO5, the two zones where GBM lost to SARIMA in Khepri v1.5 (+2.92 pp and +0.53 pp), reported descriptively; RQ4, whether TSFM predictive quantiles are better than SARIMA's on the holdout (CRPS and coverage); and RQ5, whether the TSFM advantage is larger on the retro test than on the holdout.

Hypotheses. H1: on the retro test a TSFM beats SARIMA (lower mean 96-h MAPE) in at least 5 of the 9 zones, evaluated separately for each TSFM. H2: on the retro test the TSFM advantage over SARIMA, in percentage points (pp) of MAPE, is smallest in NO5. Because NO5 has by far the lowest MAPE level (SARIMA 0.72 % in v1.5), the pp measure favours H2 by scale alone, so the relative advantage (pp divided by SARIMA MAPE) is reported as a secondary measure. The null of no difference is reported as a result in its own right.

Diebold–Mariano tests. For each TSFM against each baseline we form $d_t = \text{loss(TSFM)}_t - \text{loss(baseline)}_t$ on the zone-averaged per-origin losses and apply the Diebold–Mariano test (Diebold & Mariano 1995) with the Harvey–Leybourne–Newbold small-sample correction (Harvey et al. 1997), h = 4, because 96-hour horizons at daily origins overlap. The long-run variance uses a rectangular kernel over lags 0…3, with Bartlett weights over the same lags if it is not positive. The statistic is referred to t(n−1), two-sided. There are 6 tests per split (3 TSFMs × SARIMA and GBM), each at $\alpha = 0.05$ unadjusted, as designed; Holm-adjusted p-values across the 6 tests per split are reported alongside as a secondary safeguard.

RQ1 uses the holdout tests and RQ2 the retro tests. RQ4 applies the same test to the zone-averaged per-origin CRPS difference TSFM − SARIMA on the holdout.

Holdout decision rule. If |mean d| is smaller than the minimum detectable effect (MDE, Section 3.7) the holdout result is "not determined", regardless of p, and never read as equivalence. If |mean d| ≥ MDE and $p < 0.05$ the result is "TSFM better" ($d < 0$) or "TSFM worse" ($d > 0$); otherwise "not significant".

Contamination indicator (RQ5). With the advantage $A_t = \text{loss(baseline)}_t - \text{loss(TSFM)}_t$ (zone-averaged), the indicator is $\overline{A}_{\text{retro}} - \overline{A}_{\text{holdout}}$, with standard error $\sqrt{\text{SE}^2_{\text{retro}} + \text{SE}^2_{\text{holdout}}}$ from HAC (rectangular, lags 0…3) standard errors of each mean. A value with z = indicator/SE > 1.96 is reported as "possible contamination". The indicator also reflects any genuine change between 2025 and October 2026 and is not a test of contamination.

Per-zone results are descriptive tables of mean loss per arm, zone and split. Analyses not listed in the preregistration are labelled exploratory.

3.7 Statistical power

The MDE was computed before any TSFM run, from the stored Khepri v1.5 per-(origin, day) metric rows of the 2025 test (D-05), because no 2025 per-origin predictions were stored. The series is the zone-averaged SARIMA − GBM loss difference over 361 origins; its long-run variance with a rectangular kernel and h = 4 is $\sigma^2_{\text{LR}} = 7.6455\ \text{pp}^2$.

With $\text{MDE} = (t_{0.975,\,n-1} + t_{0.80,\,n-1})/\text{HLN}(n,h) \cdot \sqrt{\sigma^2_{\text{LR}}/n}$, where $\text{HLN}(n,h) = \sqrt{(n+1-2h+h(h-1)/n)/n}$, $\alpha = 0.05$ two-sided and power 0.8, the MDE is 2.989 pp for the n = 14 holdout origins and 0.413 pp for the n = 361 retro origins.

The holdout is therefore a contamination control with low power; the retro test carries the power.

3.8 Data traps and sensitivity analysis

Data traps are handled exactly as in Khepri v1.5, without new methods.

In the holdout window, however, the share of hours where the v1.5 CI is missing before gap filling (pumped-storage hours) is large: 54.5 % in NO2, 41.6 % in NO3 and 16.7 % in NO5, with longest gaps of 111, 45 and 58 hours; in 2025 the shares are 9.5 %, 4.7 % and 4.8 %. NO1, NO4 and SE1–SE4 have no filled hours in either window (D-07).

Since filled hours enter the truth, the preregistration adds sensitivity analysis S1: the same Diebold–Mariano tests on the zone average over the six zones without filled hours (NO1, NO4, SE1, SE2, SE3, SE4), reported descriptively next to the primary tests. The zone set was fixed from the CI data before any TSFM forecast.

A descriptive reproduction table compares our retro mean MAPE for SARIMA and GBM with the published v1.5 tables per zone (D-08); it enters no test.

3.9 Deviations

All deviations from the design and the preregistration are logged with dates in the deviation log; the table below summarises them. D-01 to D-08 were made before registration and are part of it; D-09 and D-10 were made after the `prereg-v1` tag and concern only the prospective log, not the arms or the analysis; D-11 (packaging) and D-12 (wording) change no analysis and no number.

ID	Deviation	Effect
D-01	Prediction-keeping driver around the unchanged v1.5 forecasting code	Validated against stored v1.5 runs (Section 3.4)
D-02	NO1 SARIMA fit not bit-reproducible; cause not found	Pre-registration validation: NO1 retro SARIMA mean MAPE 2.4596 vs. 2.4656
D-03	SARIMA quantiles from the Gaussian predictive distribution	Point forecast unchanged
D-04	Seasonal naive (week) added as a new reference arm	Enters no test
D-05	MDE from stored v1.5 per-origin metric rows	No 2025 predictions were stored
D-06	Holdout training rule mapped by parameters only	SARIMA training span 2.79 y (NO), 1.79 y (SE) vs. 3.0 and 2.0 y
D-07	Large share of filled pumped-storage hours in the holdout	Sensitivity S1 on six zones without filled hours
D-08	Reproduction table against published v1.5 results	Descriptive
D-09	Bug fix in the prospective-log issuer (chronological splicing of new A75 data)	Prospective log only; analysis and arms unchanged

ID	Deviation	Effect
D-10	Guard in the prospective-log issuer that skips a day's issue while another heavy job runs; such days are logged as lost, never back-filled	Prospective log only; analysis and arms unchanged
D-11	Khepri v1.5 included byte-identically as `khepri_v15/` because its tag has no package definition; only import lines changed	Packaging only; RESULTATER byte-identical in a clean environment
D-12	Preregistration wording on training cut-offs corrected: Chronos-2 became public in October 2025	Wording only; analyses unchanged

3.10 Prospective log

From registration onward, a daily job issues all arms for the 00:00 UTC origin with the holdout training rule and writes each issue with its sha256 and an OpenTimestamps stamp. Lost days are logged and never back-filled. The log will be analysed at a later date, not yet fixed, with the same locked `analysis.py`, after joining the truth from a later A75 pull through the same Khepri v1.5 CI code. It is not part of the results of this paper.

3.11 Deferred to a later version

A 32 768-step context, Timestral-28M and Timestral-123M, and covariates are deferred to a later version and are not part of this study.

3.12 Licences

The code is licensed under Apache-2.0 and the derived data under CC-BY-4.0. The Timestral weights are under Mistral AI's "Modified MIT License", which includes a revenue restriction; Chronos-2 is under Apache-2.0.

4 Results

Unless stated otherwise, the numbers in this section come from the locked `analysis.py` (tag `prereg-v1`; run with the import change of deviation D-11) applied to the six prediction files; the output is `results/RESULTATER_KHEPRI-TSFM_v1_2026-11-02.md`. Every arm produced forecasts for all nine zones at all 361 retro origins and all 14 holdout origins; no origin was lost and SARIMA never fell back to diurnal persistence. Loss is the Khepri v1.5 MAPE over the 96 target hours; the test series is the zone-averaged loss per origin; a negative difference means that the TSFM has the lower loss.

Decision rule for the holdout. The minimum detectable effect for the holdout, computed before any TSFM run, is MDE = 2.989 pp (Section 3). A holdout difference whose absolute value is below

Table 1: Retro test: zone-averaged 96-h MAPE (%), difference TSFM − baseline (pp) and the Diebold–Mariano statistic with the Harvey–Leybourne–Newbold correction ($h = 4$, $t(n-1)$, $n = 361$). All $p < 0.0001$, also after Holm adjustment over the six tests.

TSFM	baseline	MAPE TSFM	MAPE baseline	diff. (pp)	HLN stat.
Timestral-5M	SARIMA	6.253	7.561	−1.308	−13.30
Timestral-5M	GBM	6.253	7.347	−1.094	−9.86
Timestral-209M	SARIMA	6.213	7.561	−1.347	−12.77
Timestral-209M	GBM	6.213	7.347	−1.134	−9.97
Chronos-2	SARIMA	6.395	7.561	−1.166	−12.65
Chronos-2	GBM	6.395	7.347	−0.952	−7.79

2.989 pp is reported as *not determined*, never as equality, whatever its p-value.

4.1 Reproduction of the baselines

Our rerun of SARIMA and GBM reproduces the published Khepri v1.5 retro means within ±0.005 pp in every zone except NO1 SARIMA (2.4596 against the published, rounded 2.47, a difference of −0.0104 pp; against the unrounded stored v1.5 rows, 2.4656, the difference is −0.006 pp; the fit converges after one L-BFGS iteration in the current environment, deviation D-02). These reproduction figures come from `results/REPRO-V15_KHEPRI-TSFM_v1_2026-11-02.md`; the NO5 note below comes from `results/TILLEGG_KHEPRI-TSFM_v1_2026-11-02.md`.

4.2 RQ2: retro test (2025, 361 origins)

On the 2025 test all three TSFMs have a lower zone-averaged loss than both SARIMA and GBM (Table 1). The differences range from −0.952 pp (Chronos-2 against GBM) to −1.347 pp (Timestral-209M against SARIMA). The retro test is labelled *contamination not ruled out*: Timestral discloses neither its training data nor a training cut-off, and most 2025 origins precede the public release of Chronos-2 in October 2025 (deviation D-12).

H1. Each TSFM has a lower retro mean loss than SARIMA in 9 of 9 zones (Table 2); H1 (at least 5 of 9) is supported for all three arms.

H2. Measured in percentage points, the smallest advantage over SARIMA is in NO5 for all three TSFMs, so H2 is supported on its primary measure. As stated in the preregistration, this measure favours H2 by scale alone: NO5 has the lowest MAPE level (SARIMA 0.722 %). On the secondary, relative measure the smallest advantage is in SE1 for all three arms; for Timestral-209M the relative advantage in NO5 is 17.8589 %, against 13.8126 % in SE1.

4.3 RQ3: NO4 and NO5

In NO4, where GBM lost to SARIMA in Khepri v1.5, the TSFMs have the lowest retro loss (Timestral-209M 6.133 %, Timestral-5M 6.171 %, Chronos-2 6.339 %, against SARIMA 7.242 % and GBM 10.158 %). In NO5 the TSFMs are also below SARIMA (0.593–0.629 % against 0.722 %). On the holdout, NO4 shows the same order (TSFM 6.984–7.314 %, SARIMA 7.480 %, GBM 9.288 %). The NO5 holdout carries no information about ranking: the Khepri v1.5 CI series for NO5 is constant at 24.0 gCO_2-eq/kWh throughout the holdout window, so every arm except GBM (0.047 %)

Table 2: Mean 96-h MAPE (%) per zone, retro test (descriptive).

zone	snaive (week)	SARIMA	GBM	Timestral-5M	Timestral-209M	Chronos-2
NO1	2.969	2.460	2.307	2.108	2.120	2.131
NO2	4.194	3.964	3.423	3.151	3.111	3.182
NO3	8.139	6.874	6.583	5.868	5.773	5.910
NO4	10.104	7.242	10.158	6.171	6.133	6.339
NO5	1.029	0.722	1.256	0.605	0.593	0.629
SE1	14.322	11.638	10.758	10.096	10.031	10.266
SE2	13.293	10.416	9.568	8.920	8.965	9.028
SE3	5.445	4.543	4.526	3.860	3.753	3.896
SE4	22.605	20.187	17.545	15.495	15.441	16.174

Table 3: Holdout (origins 2026-10-16 to 2026-10-29): zone-averaged 96-h MAPE (%), difference TSFM − baseline (pp), HLN statistic ($t(n-1)$, $n = 14$), unadjusted and Holm-adjusted p, and the preregistered verdict (MDE = 2.989 pp).

TSFM	baseline	TSFM	base	diff.	stat.	p	p_{Holm}	verdict
Timestral-5M	SARIMA	8.415	9.332	−0.917	−1.98	0.0693	0.1385	not determined
Timestral-5M	GBM	8.415	9.223	−0.808	−5.00	0.0002	0.0012	not determined
Timestral-209M	SARIMA	8.005	9.332	−1.326	−3.79	0.0023	0.0068	not determined
Timestral-209M	GBM	8.005	9.223	−1.218	−17.97	<0.0001	<0.0001	not determined
Chronos-2	SARIMA	8.494	9.332	−0.838	−1.89	0.0810	0.1385	not determined
Chronos-2	GBM	8.494	9.223	−0.729	−4.95	0.0003	0.0012	not determined

has a MAPE of 0.000 %. This is a property of the data in that window, not a result about the models.

4.4 RQ1: post-release holdout (14 origins)

On the holdout every TSFM again has the lower point estimate, with differences from −0.729 to −1.326 pp (Table 3). All six differences lie within ±2.989 pp, so by the preregistered rule all six comparisons are *not determined*, although four of them have $p < 0.05$. With 14 origins the holdout cannot resolve differences of the size seen on the retro test; it was designed as a contamination control, not as the test that carries the power.

Sensitivity S1. Restricted to the six zones without filled pumped-storage hours (NO1, NO4, SE1–SE4), the retro differences are larger (−1.171 to −1.674 pp, all $p < 0.0001$) and the holdout differences lie between −0.981 and −1.662 pp. NO5 is not in the S1 set; in the primary zone average it enters as the preregistration requires, but on the holdout it contributes a zero loss for every arm except GBM and so carries no information about ranking.

4.5 RQ4: calibration (holdout)

RQ4 is the preregistered calibration comparison: a DM-HLN test on the zone-averaged per-origin CRPS difference TSFM − SARIMA was locked in the preregistration, and the MDE applies to MAPE only. All three TSFMs have a lower holdout CRPS than SARIMA (differences −0.3052,

Table 4: Holdout calibration: CRPS (quantile-score approximation over 21 levels, gCO_2-eq/kWh) and empirical coverage of the central 80 % and 90 % intervals.

arm	CRPS	coverage 80 %	coverage 90 %
SARIMA	1.4165	0.835	0.911
Timestral-5M	1.1114	0.838	0.915
Timestral-209M	1.0733	0.838	0.921
Chronos-2	1.1198	0.806	0.905

-0.3432 and -0.2967 gCO_2-eq/kWh for Timestral-5M, Timestral-209M and Chronos-2; HLN statistics -11.75, -11.23 and -8.52; all $p < 0.0001$). The SARIMA quantiles are constructed from the fitted model's Gaussian predictive distribution (deviation D-03), not taken from a native v1.5 probabilistic forecast, so the difference may partly reflect the SARIMA intervals rather than the TSFMs alone. Empirical coverage lies between 0.806 and 0.838 for the 80 % interval and between 0.905 and 0.921 for the 90 % interval across the four arms (Table 4); these are 14 origins and are not tested.

4.6 RQ5: contamination indicator

The indicator is the TSFM advantage on the retro test minus the advantage on the holdout. It ranges from -0.084 pp (Timestral-209M against GBM) to 0.391 pp (Timestral-5M against SARIMA), with z between -0.68 and 1.75. No comparison exceeds the preregistered threshold $z > 1.96$, so none is reported as possible contamination. The indicator is descriptive: it also absorbs any genuine difference between 2025 and October 2026, and with 14 holdout origins its standard errors are wide (0.124–0.360 pp).

5 Software and availability

5.1 Summary

Khepri-TSFM is a small, preregistered evaluation package: it fetches ENTSO-E A75 data, derives hourly carbon intensity with Khepri v1.5, runs three zero-shot TSFMs and the Khepri v1.5 baselines on a fixed set of 96-hour forecast origins, and applies one locked analysis script. The repository holds the CI series every arm saw, all forecasts with 21 quantiles, the analysis output and the ledgers that tie every number in this paper to a file and line.

5.2 Statement of need

Carbon-aware scheduling needs day-ahead to four-day CI forecasts per bidding zone, and TSFMs are now offered as drop-in forecasters. A comparison against an already published baseline is only informative if the baseline is not retuned, the forecast origins and context are shared across arms, and the analysis is fixed before the TSFM forecasts exist. Khepri-TSFM provides such a comparison for the nine Nordic zones, with a post-release holdout as a contamination control.

5.3 State of the field

TSFM-based CI forecasting tools and live TSFM benchmarks exist (Section 2). This study is kept as its own repository rather than contributed to a general toolkit or benchmark because its object is a frozen, versioned comparison: the baselines, splits, pinned model revisions and analysis must

stay bound to the preregistration and its timestamps, which a contribution into a moving code base would not preserve.

5.4 Software design

Four trade-offs shape the package. (i) The Khepri v1.5 SARIMA and GBM code and hyperparameters are used unchanged; a driver repeats the v1.5 evaluation loop only to keep the hourly forecasts (deviation D-01). (ii) Khepri v1.5 is vendored byte-identically as `khepri_v15`, because its tag has no package definition (D-11). (iii) All TSFMs get the same 8192-hour context, the maximum of Chronos-2; longer contexts are deferred to v2. (iv) Everything runs on a CPU (Apple M1, 16 GB): with a batch of nine zones one forecast call takes a median of 0.143 s (Timestral-5M), 1.800 s (Timestral-209M) and 0.922 s (Chronos-2), and a SARIMA fit takes 17.8 s (`results/LATENS_KHEPRI-TSFM_v1_2026-11-02.md`).

5.5 Research impact statement

The package reproduces the published Khepri v1.5 baselines and extends them with probabilistic forecasts; the CI series, all forecasts and the analysis can be re-run from the repository without ENTSO-E access (`repro_from_ci.py`); a daily prospective log continues the comparison beyond the holdout. We make no claim beyond the nine zones and the periods tested.

5.6 AI usage disclosure

Implementation was assisted by Claude Code (Anthropic) as a tool under the author's direction; roles, tools and versions are listed in `AIUSE_KHEPRI-TSFM_v1_2026-11-02.md`. The author is responsible for all choices and numbers.

Data and code availability

Code, the hourly CI series, all forecasts and the analysis output: Zenodo, DOI 10.5281/zenodo.23881315 (repository github.com/avalyset/khepri-tsfm). Khepri v1.5.0: DOI 10.5281/zenodo.23101127.

AI use

Implementation was assisted by Claude Code (Anthropic) as a tool under the author's direction (`AIUSE_KHEPRI-TSFM_v1_2026-11-02.md` in the repository). The author is responsible for all choices and numbers.

References

Ansari, A. F., Shchur, O., Küken, J., et al. (2025). Chronos-2: From univariate to universal forecasting. arXiv:2510.15821.

Botten Nicolaysen, E. (2026). Khepri v1.5.0. Zenodo. doi:10.5281/zenodo.23101127.

Clauß, J., Stinner, S., Solli, C., Lindberg, K. B., Madsen, H., & Georges, L. (2019). Evaluation method for the hourly average CO2eq. intensity of the electricity mix and its application to the demand response of residential heating. *Energies*, 12(7), 1345. doi:10.3390/en12071345.

Cramer, E. Y., Ray, E. L., Lopez, V. K., et al. (2022). Evaluation of individual and ensemble probabilistic forecasts of COVID-19 mortality in the United States. *Proceedings of the National Academy of Sciences*, 119(15), e2113561119. doi:10.1073/pnas.2113561119.

Diebold, F. X., & Mariano, R. S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263. doi:10.1080/07350015.1995.10524599.

Garza, A., Rosillo, R., Mendoza-Smith, R., Salinas, D., Williams, A. R., Ashok, A., Goswami, M., & Juárez, J. M. (2026). Impermanent: A live benchmark for temporal generalization in time series forecasting. arXiv:2603.08707.

Gneiting, T., & Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477), 359–378. doi:10.1198/016214506000001437.

Harvey, D., Leybourne, S., & Newbold, P. (1997). Testing the equality of prediction mean squared errors. *International Journal of Forecasting*, 13(2), 281–291. doi:10.1016/S0169-2070(96)00719-4.

Hong, T., Pinson, P., Fan, S., Zareipour, H., Troccoli, A., & Hyndman, R. J. (2016). Probabilistic energy forecasting: Global Energy Forecasting Competition 2014 and beyond. *International Journal of Forecasting*, 32(3), 896–913. doi:10.1016/j.ijforecast.2016.02.001.

Leerbeck, K., Bacher, P., Junker, R. G., Goranović, G., Corradi, O., Ebrahimy, R., Tveit, A., & Madsen, H. (2020). Short-term forecasting of CO2 emission intensity in power grids by machine learning. *Applied Energy*, 277, 115527. doi:10.1016/j.apenergy.2020.115527.

Maji, D., Sitaraman, R. K., & Shenoy, P. (2022a). DACF: Day-ahead carbon intensity forecasting of power grids using machine learning. In *Proceedings of the Thirteenth ACM International Conference on Future Energy Systems (e-Energy '22)*, 188–192. doi:10.1145/3538637.3538849.

Maji, D., Shenoy, P., & Sitaraman, R. K. (2022b). CarbonCast: Multi-day forecasting of grid carbon intensity. In *Proceedings of the 9th ACM International Conference on Systems for Energy-Efficient Buildings, Cities, and Transportation (BuildSys '22)*, 198–207. doi:10.1145/3563357.3564079.

Maji, D., Yang, K., Shenoy, P., Sitaraman, R. K., & Srivastava, M. (2025). CarbonX: An open-source tool for computational decarbonization using time series foundation models. arXiv:2510.01521.

Mistral AI (2026). Timestral. github.com/mistralai/timestral, commit a534bdbd244d64f3d9d743f46c7bf6fc330ed4ff.

Yan, L., Wang, L., Liu, S., & Ding, Y. (2025). EnsembleCI: Ensemble learning for carbon intensity forecasting. In *Proceedings of the 16th ACM International Conference on Future and Sustainable Energy Systems (e-Energy '25)*, 208–212. doi:10.1145/3679240.3734630.